EU AI ACT · ARTICLE 50

THE AI DISCLOSURE PLAYBOOK

A Practical Guide to **Article 50** of the EU AI Act for Marketers, Creators and Small Businesses

ZIRLANDIA MILKOVIC

EU AI ACT · ARTICLE 50

THE AI DISCLOSURE PLAYBOOK

A Practical Guide to Article 50 of the EU AI Act for Marketers, Creators and Small Businesses

ZIRLANDIA MILKOVIC

The AI Disclosure Playbook
A Practical Guide to Article 50 of the EU AI Act for Marketers, Creators and Small Businesses

© 2026 Zirlandia Milkovic. All rights reserved. No part of this book may be reproduced without written permission, except for brief quotations in reviews.

First edition, 2026. Research cut-off: 11 November 2026. British English. Companion resource: GDPRGard.eu.

This book is practical guidance, not legal, security or data-protection advice. All examples, businesses and people in worked scenarios are fictional. Cover and illustrations created by the author with design and AI image tools.

Contents

START HERE

Educational Disclaimer 4
Introduction: Less Explanation, More Instructions 4
Who This Book Is For 5
How to Use This Book 6
Where Do You Stand? Ten Quick Questions 7

CHAPTERS

1. What Article 50 Requires and What It Does Not 9
2. The Disclosure Test: Does This Content Need a Label? 14
3. AI-Generated Images, Audio and Video 19
4. Deepfake Transparency 24
5. AI-Generated Text and Public-Interest Publications 30
6. Human Review and Editorial Responsibility 35
7. Customer-Facing AI Transparency 39
8. Practical Disclosure Workflows 44
9. The Label and Wording Library 49
10. Your 30-Day Rollout and Three Case Files 53

CLOSING

Conclusion: Decide, Label, Own, Record 57

APPENDICES

Appendix A: Content Classification Sheet 58
Appendix B: Label Quick Reference Card 59
Appendix C: AI Use Register 60
Appendix D: Human Review Sign-Off Form 61
Appendix E: Client AI-Use Clause Checklist 62
Appendix F: One-Page AI Disclosure Policy Template 63

BACK MATTER

Glossary 64
A Note on GDPRGard.eu 65
About the Author 66
Notes and Sources 67

Educational Disclaimer

This book is practical, educational guidance for marketing agencies, online retailers, content creators and small brands. It is not legal advice, and it does not create a professional relationship between the author and the reader.

The EU AI Act, consumer law, advertising rules and data-protection law differ by country, apply differently to different businesses, and are still being clarified by guidelines, codes of practice and regulators. Check the current position with a qualified adviser or your regulator before relying on anything here for a decision that matters.

Following the steps in this book may reduce your risk. It cannot guarantee that you will meet every legal duty. Checklists, wording and templates are starting points to adapt, not finished legal documents.

Every business, person and product used in a worked example is **fictional** and is labelled as such. Where the book describes an official document, it says which one and whether it was read in full. Where it gives a recommendation rather than a legal requirement, it says so.

Introduction: Less Explanation, More Instructions

There is no shortage of articles explaining the EU AI Act. Most of them are long, and most of them stop where your problem starts.

Your problem is not "what is a deployer?" Your problem is Tuesday. A designer has just exported six product photos from an image generator. A freelancer has sent a voiceover that sounds suspiciously like a real actor. The chatbot vendor has asked whether you want to "enable the AI badge". A client wants a newsletter about a new health regulation out by five. Someone in the team asks, *"Do we have to say this is AI?"* and nobody knows who decides.

Article 50 of the EU AI Act, the transparency rule, has applied since **2 August 2026**. [1][2] The European Commission has since published guidelines on how to read it, and a final Code of Practice on marking and labelling AI-generated content. [3][4] Those documents are useful. They are also written for lawyers, regulators and the companies that build AI models.

This book is written for the people who have to *ship*. It does three things:

1. It reduces the rule to **one decision test** you can apply to any piece of content in under a minute.
2. It gives you **the wording, the placement and the paperwork** for each kind of content: images, audio, video, text, chatbots and voice agents.
3. It sets out a **workflow** with named roles, so that "do we have to say this is AI?" always has an owner and a recorded answer.

> **The rule.** Decide before you publish, label where people will see it, name a human who answers for it, and write down what you did.

That is the whole playbook in one sentence. The rest of the book makes each part of it concrete.

Who This Book Is For

This book is for you if you run, manage or advise a small or mid-sized organisation in Europe or serving European audiences, and any of the following is true:

- You work at, or run, a **marketing or creative agency** that produces copy, images, audio or video for clients, with or without AI tools.
- You run an **online shop** that uses AI for product images, descriptions, chat support, emails or reviews.
- You are a **content creator, newsletter writer, podcaster, course author or consultant** who uses AI to draft, illustrate, translate, voice or edit.
- You manage a **small brand** with a social-media calendar, a website and a customer inbox.

- A client, platform, insurer or regulator has started asking, "Did you use AI for this?"

It assumes no technical background. Terms are explained when they first appear and collected in the glossary.

It is **not** written for organisations that build and sell AI models or AI systems to others. They carry heavier duties, including the technical marking of outputs. If that is you, treat this book as orientation and take specialist advice.

How to Use This Book

You can read the book in an evening, or keep it open as a manual.

Part one, Chapters 1 and 2, gives you the map: what Article 50 requires and the decision test. **Part two, Chapters 3 to 7,** goes content type by content type: images, audio and video; deepfakes; AI-written text; human review; and customer-facing AI. **Part three, Chapters 8 to 10,** is the machinery: workflows, a wording library, agencies and clients, and a 30-day rollout with case files.

Look for these markers:

- **The rule.** A principle worth remembering, set in a dark panel.
- **Fictional.** A worked example using invented people and businesses.
- **Say this.** Wording you can adapt.
- **Our advice.** A recommendation that goes beyond what the rule strictly requires. These are labelled so that you can tell what is the law and what is good practice.
- Check boxes (☐). A step you can tick off. Where a line is blank, write on it.

The appendices hold the templates: a content classification sheet, a label quick-reference, an AI-use register, a sign-off form, a client AI clause checklist and a one-page disclosure policy. Copy them into a spreadsheet or print them.

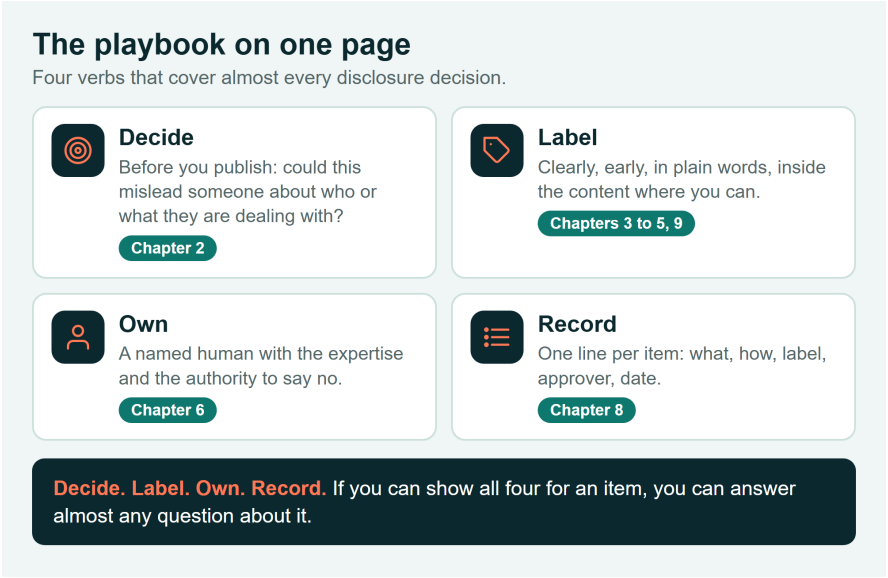

The playbook on one page: decide, label, own, record

Where Do You Stand? Ten Quick Questions

Answer these before you read further. Each "yes" points to a chapter.

1. ☐ Do we publish AI-generated or AI-edited images, video or audio that could be mistaken for the real thing? (*Chapters 3 and 4*)
2. ☐ Do any of them show a real person, product, place or event doing or being something that did not happen or does not exist as shown? (*Chapter 4*)
3. ☐ Do we publish AI-written or AI-assisted text about news, health, law, politics, money or other public-interest topics? (*Chapter 5*)
4. ☐ Does a named human review that text and answer for it before it goes out? (*Chapter 6*)
5. ☐ Do customers chat or speak with an AI system in our name? (*Chapter 7*)
6. ☐ Is it obvious, from the first message or the first second, that it is a machine? (*Chapter 7*)

7. ☐ Does anyone decide, per item, whether something needs a label, and is that decision written down? (*Chapters 2 and 8*)
8. ☐ Do we produce work for clients, or take work from freelancers, without a written statement of where AI was used? (*Chapter 9*)
9. ☐ Could we show a regulator, client or platform, today, what we disclosed and when? (*Chapter 8*)
10. ☐ Does everyone who publishes know the three or four labels we use? (*Chapter 9*)

Questions 4, 6, 7 and 9 are the ones where a "no" is the problem. If you answered "no" to any of them, start with the 30-day plan in Chapter 10, then read the chapters your other answers point to.

CHAPTER 1

What Article 50 Requires and What It Does Not

Article 50 is short. Its effect on a small business is shorter still, once you strip out the parts that are aimed at someone else. This chapter sorts out which parts are aimed at you.

Four duties, two kinds of organisation

Article 50 creates four transparency duties. Each one lands on either the **provider** or the **deployer** of an AI system. [1]

- A **provider** develops an AI system, or has one developed, and puts it on the market under its own name. The company behind a chatbot platform or an image generator is a provider.
- A **deployer** uses an AI system under its own authority in a professional capacity. If you use that chatbot on your website, or that image generator for client work, you are a deployer.

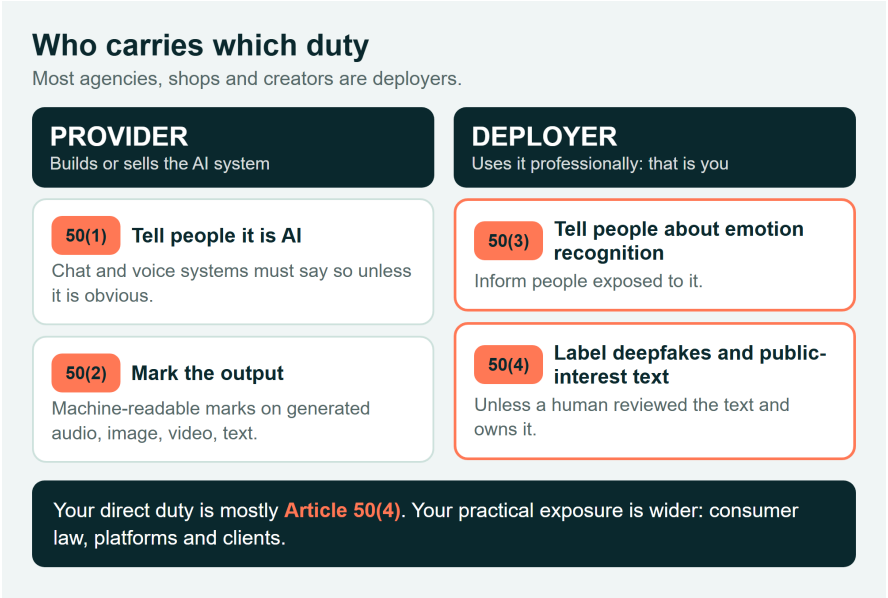

Who carries which duty: provider or deployer

Duty	Carried by	What it asks for
Art. 50(1)	Provider	Design AI systems that interact directly with people so that people are told they are dealing with AI, unless that is obvious
Art. 50(2)	Provider	Mark AI-generated audio, image, video and text in a machine-readable way so it can be detected as artificial
Art. 50(3)	Deployer	Tell people who are exposed to an emotion-recognition or biometric-categorisation system that it is operating
Art. 50(4)	Deployer	Disclose deepfakes, and disclose AI-generated text published to inform the public on matters of public interest, unless a human reviewed it and someone holds editorial responsibility

Two things follow for a typical agency, shop or creator.

First, most of your direct legal duty is in Article 50(4). That means deepfakes and public-interest text. Those are the subject of Chapters 4 to 6. The chatbot duty (50(1)) and the machine-readable marking duty (50(2)) sit with the vendor, although you must not defeat them (Chapter 7 and Chapter 3).

Second, you can become a provider without noticing. If you build your own AI system, or put your name and branding on someone else's tool and sell it on, the provider duties may be yours too. An agency that packages a

white-labelled chatbot for clients is the standard example. If that is you, read Chapter 9 and take specialist advice.

What the rule is not

It is just as useful to know what Article 50 does not do.

- It does **not** say that every piece of AI-assisted content needs a label. A grammar fix, an auto-cropped photo or a translated product description is not a deepfake and is not public-interest text.
- It does **not** ban any AI use. It is a transparency rule, not a prohibition.
- It does **not** replace older law. Consumer protection rules on misleading practices, the ban on fake reviews, advertising standards and GDPR transparency duties still apply and often bite harder. [9][10]
- It is **not** a duty to watermark your own content. The technical marking duty is the provider's. Your duty is not to strip it out, and to add a visible label where Article 50(4) requires one.

The dates that matter

Article 50 has applied since **2 August 2026**. One narrow transition exists: providers of generative AI systems that were already on the EU market before that date have until **2 December 2026** to comply with the *marking* duty in Article 50(2). Content created before 2 August 2026 does not have to be labelled retroactively, although the Commission encourages it. [2]

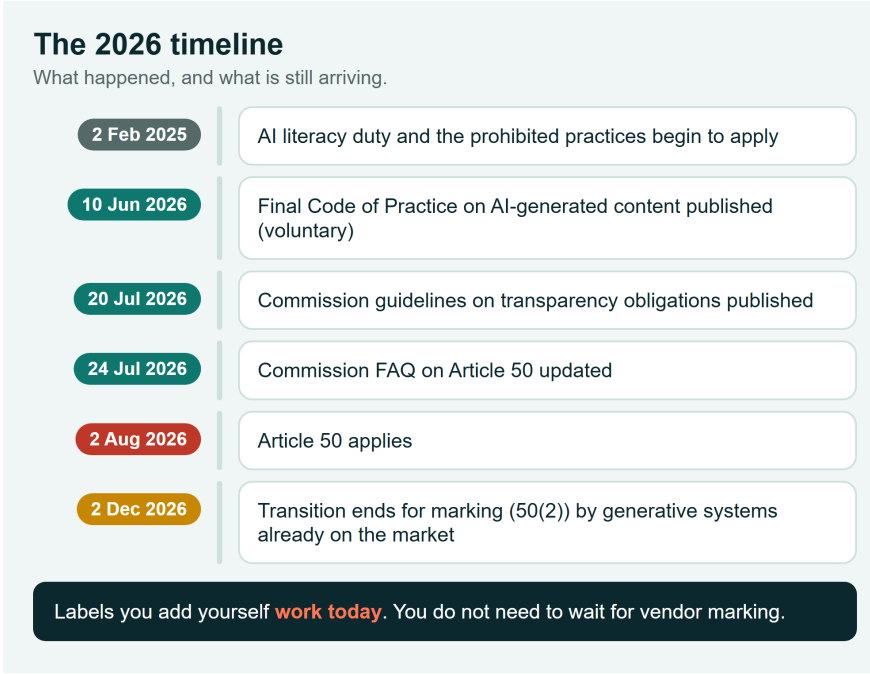

The 2026 timeline for Article 50

For you as a deployer, the practical reading is simple: the labelling and disclosure duties are live now, and the vendor-side marking is still arriving. Check the current dates against the Official Journal, because amending legislation has been discussed. [8] Do not wait for the vendors. A label you add yourself is the thing a customer, platform or regulator can actually see.

Guidelines and the Code of Practice

Two official documents sit beside the Act.

The Commission's **guidelines on transparency obligations** were published on 20 July 2026. They explain how to read terms such as "deepfake", "obvious", "human review" and "editorial responsibility". [3] A set of Commission FAQs summarises them in plainer language. [2]

The **Code of Practice on transparency of AI-generated content** reached its final version on 10 June 2026. It is **voluntary**. It has two sections: one for providers (marking and detection) and one for deployers (labelling deepfakes and public-interest text). The Commission and the AI Board

have assessed it as an adequate way to demonstrate compliance. If you follow it, you can rely on its measures as evidence. If you do not sign, you need to show by other means that what you do is equally adequate. [4][5]

> **Our advice.** You do not need to sign the Code to follow its good ideas. Its deployer section is a sensible reference for label design and placement. Use it as a checklist, and sign only if a client or procurement process asks you to.

What happens if you get it wrong

Breaches of Article 50 can attract administrative fines of up to **EUR 15 million or 3% of worldwide annual turnover**. [2][7] The Commission's FAQ says proportionality can be taken into account for small and mid-sized firms, and the penalty article of the Act contains a rule for SMEs on which figure applies. Check that rule with an adviser for your size and country. Enforcement is mainly through national market-surveillance authorities, so what you actually face is a national regulator, and the way they prioritise will differ by country. [2]

In practice, the first consequences a small business meets are rarely fines. They are a platform removing a post, a client terminating a contract, a customer complaint under consumer law, or a journalist asking questions. The playbook is designed to make all four easy to answer.

> **The rule.** Your direct legal duty is mostly about deepfakes and public-interest text. Your practical exposure is wider. Build for both.

CHAPTER 2

The Disclosure Test: Does This Content Need a Label?

You cannot memorise four duties, three exceptions and a grey zone for every post. You need a test that takes a minute and gives the same answer whoever applies it.

Six questions, in order

Ask these about each piece of content, in this order. Stop at the first "yes" that applies and follow the instruction.

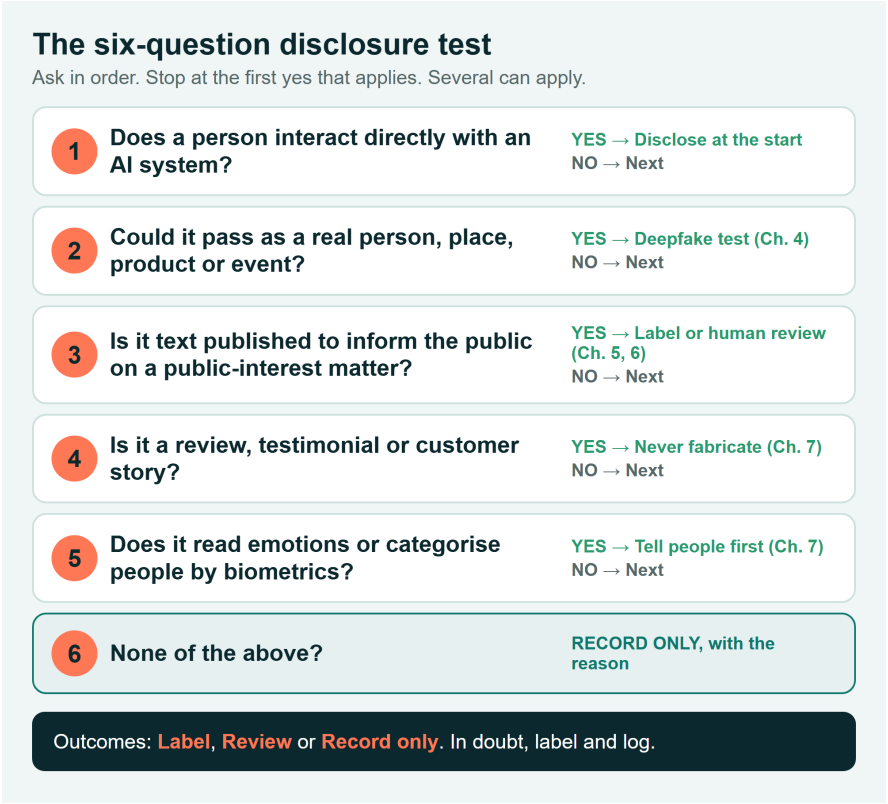

The six-question disclosure test

1. Does a person interact directly with an AI system? A chat window, a voice agent, an AI receptionist. If yes, go to Chapter 7: tell them at the start, unless it is obvious.

2. Is it an image, audio or video that resembles real people, objects, places or events and could pass as authentic? If yes, it may be a deepfake. Go to Chapter 4 and apply the three-part test.

3. Is it text published to inform the public on a matter of public interest? News-style articles, explainers on law, health, money, politics, safety. If yes, go to Chapter 5, and then Chapter 6: label it, or have a human review it and take editorial responsibility.

4. Is it a review, testimonial, endorsement or "customer story"? If AI wrote or invented any part of it, the problem is not labelling. It is consumer law. Do not publish fabricated reviews. Go to Chapter 7.

5. Does the system detect emotions or sort people into categories from biometric data? If yes, people must be told. Go to Chapter 7.

6. None of the above? Ordinary marketing copy, product descriptions, schedules, emails, internal documents, stylised illustrations. No Article 50 label is required. Apply your own honesty standard and record the decision.

Why the order matters

The order is deliberate. The first questions catch the situations where *a person might be misled about who or what they are dealing with*. The last catches everything else. A product description drafted by AI and checked by a human is in question six. The same text, if it were a news-style explainer on a new tax rule, would be in question three.

The same content can fall into more than one question. A synthetic presenter reading an AI-written script about a public-health change is a deepfake candidate *and* public-interest text. Apply every question that fits, and use the most protective answer.

The three outcomes

Every decision lands in one of three outcomes.

Outcome	Meaning	What you do
Label	A disclosure duty applies, or applying one is clearly wise	Add the label from the library (Chapter 9), in the right place
Review	The exemption depends on a human reviewing and owning it	Run the sign-off (Chapter 6) and record the name
Record only	No label needed	Note the decision and the reason in the log

> **Our advice.** The third outcome matters more than it looks. "No label needed, and here is why" is what you will be asked to show. A one-line entry costs thirty seconds.

A creator's desk where a quick decision gets recorded

Edge cases that come up every week

Stylised illustration. A cartoon mascot, a flat-colour icon or an obviously artistic poster does not look like a real person or event and cannot falsely appear authentic. It is unlikely to be a deepfake. Label it if you wish; you probably need not.

AI-edited real photo. Removing a bin from a street photo or adjusting light is routine editing. Putting a person in a place they never were, or changing what someone is doing or saying, is manipulation that may fall within the deepfake definition. [3]

Translation, summaries and rewrites. A translation of your own human-written text is ordinary assistance. A summary of a public-interest report published to inform the public may be public-interest text. Check question three.

Social posts. A short, informal post is unlikely to be the kind of public-interest publication Article 50(4) targets. A post that presents a synthetic video of a real person is a deepfake question regardless of its length.

Voice clones. A cloned voice of a real person delivering words they did not say is a deepfake candidate. A cloned voice of a person who agreed, speaking words they approved, is a consent and contract matter, and a label is still wise.

When you are not sure

If you cannot answer a question with confidence, treat the content as *in scope*, label it and record why. Over-disclosure in a single doubtful case costs little. A pattern of silence in doubtful cases is what regulators and platforms notice.

> **The rule.** When in doubt, label and log. A rare unnecessary label is a small cost. A missing one is not.

CHAPTER 3

AI-Generated Images, Audio and Video

Images, audio and video are where AI content is most convincing and where the rules bite hardest. Two different duties apply, and mixing them up is the most common mistake.

Marking is not labelling

Marking is technical. It means embedding signals, such as watermarks and metadata, so that software can detect that content is AI-generated. It is the provider's duty under Article 50(2). [1]

Labelling is visible to humans. It means telling the audience in words or an icon. For deepfakes, it is the deployer's duty under Article 50(4). [1]

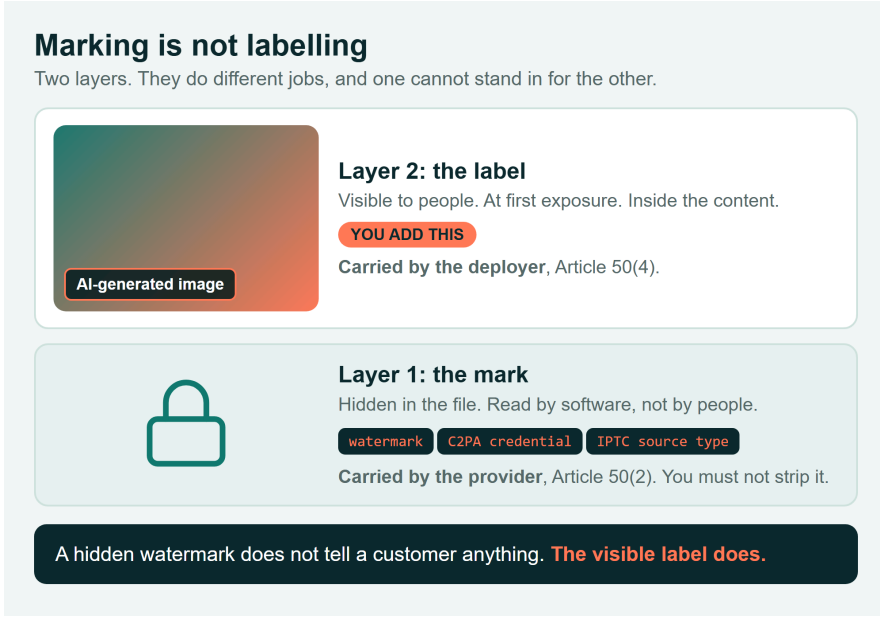

The Commission's FAQ is explicit that a deployer cannot rely on the provider's machine-readable mark to meet the labelling duty. The label must be perceivable by people, without special tools, at the latest at first exposure. [2]

For you, this produces three working rules:

1. **Never strip marks or metadata** from generated files before publishing. Cropping and re-exporting can remove them without anyone intending it. Check your export settings.
2. **Add a visible label yourself** whenever the content is a deepfake or public-interest text.
3. **Keep the original file** with its metadata, in your archive, with the log entry.

Product images

Product imagery is the biggest grey zone for shops and agencies. It is worth being precise about it.

A flat lay: real product, generated background, or neither?

Image type	Typical risk	What to do
Real product photographed, AI used for retouching or background clean-up	Low	Record only

Image type	Typical risk	What to do
Real product, AI-generated scene around it, clearly styled	Low to medium	Label if the scene could mislead about the product's context or size
AI-generated image of the product itself	**High** (consumer law)	Label, and make sure it matches the real product. Never show features it does not have
AI-generated "customer" or "model" shown as a real person wearing or using the product	Medium to high	Label as AI-generated. Never present as a real customer or real result
AI-generated "before and after" or result	**High**	Do not use unless it is a clearly labelled illustration

> **Our advice.** The AI Act question is whether the image is a deepfake. The consumer-law question is whether it misleads a buyer about the product. The second is the one most small retailers fail first. Whatever the label says, the product in the photo must be the product in the box.

Audio

Audio deserves attention because it is increasingly used for ads, podcasts, explainer videos and phone lines.

Synthetic voice reading your own script, in a clearly artificial voice, for a product explainer, is rarely a deepfake. A short spoken line at the start or in the show notes ("Voiceover generated with AI") costs nothing and is good practice.

A cloned voice of a real person is the risky case. If it could be taken for the real person speaking words they did not say, treat it as a deepfake: label it audibly, at the start, not only in the caption. Get written consent from the person, and write down what words the clone may be used for.

Audio-only formats need an audible label. A text caption nobody reads is not enough for a podcast or a radio spot. The Commission's icon guidance covers image, audio and video deepfakes, and asks for accessible, plain-language labels. [6]

Video

Video combines everything. A single video may contain synthetic visuals, a synthetic voice, AI-written subtitles and a real presenter.

- **Fully synthetic video** that looks like filmed reality: label at the start, and keep the label visible or repeat it.
- **Real footage with AI-altered elements**, such as a changed background, a face swapped in, or a voice replaced: label the altered element, and the video as a whole if the alteration matters to what viewers believe.
- **AI avatars or "digital presenters"**: tell viewers that the presenter is synthetic. A viewer who thinks they are watching a real expert is being misled even if the script is accurate.

Video and audio need a label at the start, not just in the caption

Platform labels are not enough

Major platforms offer their own "AI-generated" toggles and detect some marked content. Use them, but do not treat them as your compliance. Platform rules change often, are applied differently by format, and may label content in a way the viewer cannot see in every context. [12] A label embedded in the file or visible in the frame travels with the content when it is shared, downloaded or embedded elsewhere. A platform setting does not.

Checklist for visual and audio content

- ☐ Do we know which tool produced each asset, and is the original kept?
- ☐ Have we checked that export did not remove embedded marks or metadata?
- ☐ Could a reasonable person take it for a real person, place, product or event?
- ☐ If yes, is there a visible or audible label at first exposure?
- ☐ Does the product, result or person shown match reality?
- ☐ For voice clones: is there written consent and a record of what it covers?
- ☐ Is the decision logged with a name and a date?

CHAPTER 4

Deepfake Transparency

"Deepfake" sounds like something that happens to politicians and film stars. The legal definition is wider and plainer, and it reaches ordinary commercial content.

The definition, in working form

The AI Act defines a deepfake as AI-generated or manipulated **image, audio or video content** that **resembles existing persons, objects, places, entities or events** and **would falsely appear to a person to be authentic or truthful**. [1] In practice the definition breaks into three tests, all of which must be met. The Commission's guidelines stress that "existing" includes things that plausibly could have existed, and that context and message decide whether something would falsely appear authentic. [2][3]

Three questions that make content a deepfake

All three must be present. Remove any one and the label duty falls away.

1. Resembles something real
A person, product, place, company or event, including things that plausibly could have existed.

+

2. AI-generated or altered
Produced or materially changed by an AI system.

+

3. Looks authentic
A reasonable viewer would take it as real or truthful, judged by context and message.

↓

ALL THREE = DEEPFAKE

Label it at first exposure. Fiction that is **evidently** artistic or satirical needs only a light-touch disclosure.

1. **Resemblance.** Does it look or sound like something real: a person, a product, a street, a company, an event? That includes things that plausibly could have existed, not only things that did.
2. **Generation or manipulation.** Was it produced or materially altered by an AI system?
3. **False appearance of authenticity.** Would a viewer or listener take it as real or truthful? The guidelines point to resemblance, the message, the deployment context and what the audience expects.

All three must be present. A fantasy dragon fails the first. A photo you took yourself and cropped fails the second. A clearly cartoon rendering of a real celebrity may fail the third.

Everyday examples

Content	Deepfake?	Why
Realistic synthetic "customer" praising a real service, presented as a real customer	**Yes, and misleading**	Resembles a person, AI-generated, appears authentic
Synthetic spokesperson video, labelled in the first frame	Yes, **disclosed**	Resembles a person, appears authentic unless told
AI image of a real street with a fictional shop on it, for an ad	Likely	Resembles an existing place, appears authentic
A pastel illustration of a smiling mascot	Unlikely	Does not appear authentic
A real photo, with the sky replaced by AI for a better mood	Depends	Judge by whether the change misleads about the place or event
A cloned voice of a founder reading a new announcement they never recorded	**Yes**	Resembles a real person, words they never said

> **Our advice.** The guidelines say routine production effects used in filmmaking are unlikely to count. Do not rely on that to wave through anything polished. Ask the blunt question: *would the average viewer be surprised to learn this was* AI? If yes, label it.

The artistic, creative and satirical exception

Where deepfake content forms part of an **evidently artistic, creative, satirical or fictional work**, the duty is lighter. You still disclose, but only in a way that does not hamper the display or enjoyment of the work. [2][6]

That is useful for a brand film, a comic sketch or a fictional ad. It is not a loophole. "Evidently" is the key word. If viewers cannot tell at a glance that it is fiction, the lighter regime is not available. A realistic fake testimonial is not an "artistic work".

How to label a deepfake

The Commission's guidance and the Code of Practice share four expectations for the label. [2][6]

1. **At first exposure.** The label must be perceivable at the latest when the viewer or listener first meets the content.
2. **Clear and distinguishable.** Plain language, readable size, not hidden in a footer.
3. **Not covered or removable by ordinary use.** Not buried under an overlay or a cropped edge. Stays attached when the content is reshared or downloaded.
4. **Accessible.** Alt text for images, spoken or on-screen equivalents for audio and video, and plain-language wording. Where an abbreviation is used, "AI" is the one the guidance accepts.

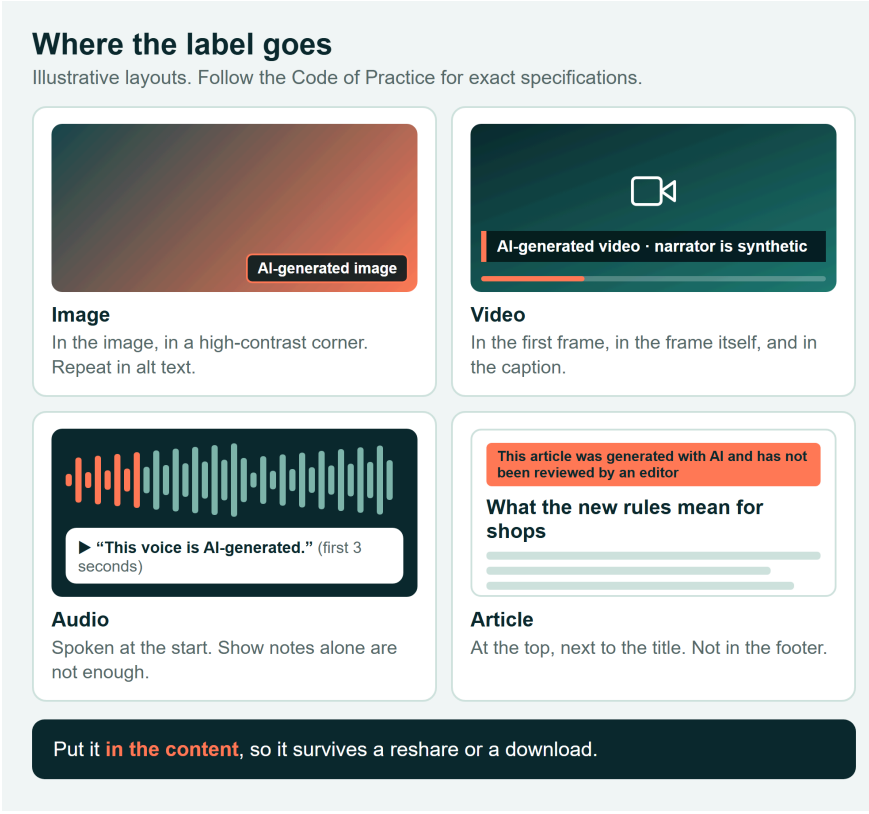

Where the label goes: image, video, audio and article

The EU icons

The Code of Practice offers an optional **EU icon** to signal AI involvement, with three designs: a basic icon for AI involvement, a "fully AI-generated" version for content with no human-created elements, and a "partially AI-modified" version where human content was altered by AI. Each comes in black, white and half-transparent variants. [6]

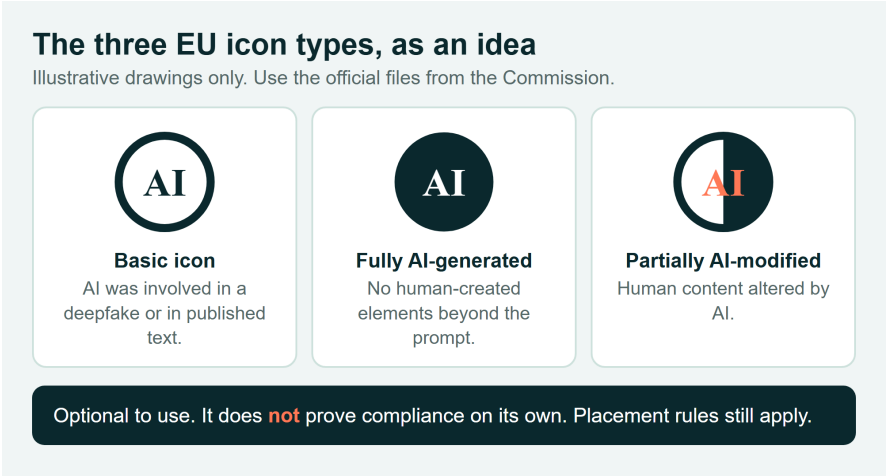

How the three icon types map to your content (illustrative; download the official artwork from the Commission)

Three practical points:

- **Using the icon is optional,** and it does not by itself prove compliance.
- **You must still follow the placement and visibility rules.** An icon in the wrong place is not a label.
- **Use the official files** from the Commission's page, not a lookalike. The artwork shown in this book is an illustration of the idea only.

Consent, likeness and the rest of the law

A label does not cure everything. A realistic synthetic version of a **real person** raises separate issues:

- **Consent and image rights** under national law.
- **Data protection.** A real person's face or voice is personal data. [10]
- **Defamation or impersonation,** if the content suggests the person said or did something they did not.
- **Contract.** Agency and creator agreements should say who may use a person's likeness, for what and for how long.

> **The rule.** Label the deepfake, and get permission for the person. The label tells the audience what it is. It does not give you the right to make it.

Checklist: deepfake transparency

- ☐ Resemblance, generation and authenticity, all three checked and recorded
- ☐ If a real person's likeness or voice is used: written consent on file
- ☐ Label visible or audible at first exposure
- ☐ Label stays with the file when shared or downloaded
- ☐ Alt text or audio equivalent provided
- ☐ If relying on the artistic exception: is the fiction evident at a glance?
- ☐ Original file with metadata archived

CHAPTER 5

AI-Generated Text and Public-Interest Publications

Most of the text your business produces needs no AI label. The exception is narrow, specific and easy to miss, and it is where agencies and creators are most likely to slip.

The text duty in one sentence

Article 50(4) requires deployers to disclose that text was artificially generated or manipulated when it is **published with the purpose of informing the public on matters of public interest**. The duty does not apply where the text has undergone **human review or editorial control** and a **natural or legal person holds editorial responsibility** for the publication. [1][2]

Read that as a three-part gate. The label is required only if **all** of these are true:

1. The text is **AI-generated or manipulated**.
2. It is **published**, to inform the public on a **matter of public interest**.
3. **No human review and editorial responsibility** exist.

Remove any one of the three and the duty falls away. That is the lever you control: the third.

The three-part gate for AI-written text

What counts as a matter of public interest

The Commission's FAQ gives examples including politics, public health and economic developments. [2] Other topics sit in the same family: law and regulation, safety, the environment, elections, justice, consumer rights, public services.

Work through the usual content types.

Content	Public-interest text?	Notes
Product description for a linen shirt	No	Commercial copy, not informing the public on public affairs
Blog: "What the new packaging rules mean for shops"	**Likely yes**	Informs the public on regulation
Newsletter on a new health guideline	**Yes**	Public health
Social post about your sale	No	Promotional
Local news summary	**Yes**	News
"Five ways to save on energy bills" with policy claims	Depends	Look at the claims: policy facts push it towards yes
Internal memo, private email	No	Not published to the public

> **Our advice.** If your content strategy includes explainer articles, regulatory updates, health or money advice, you are already in the zone. Treat all such content as in scope and run it through Chapter 6, whether or not you are sure it is "news".

What does not count as review

The guidelines draw the line sharply. **Human review** means deliberate, substantive examination by people with relevant expertise. **Editorial control** means an entity able to approve, alter or reject the substance, including fact-checking. **Spell-checking and grammar fixes do not count. Editorial responsibility** means ultimate legal responsibility for the publication. [2]

That produces a ladder.

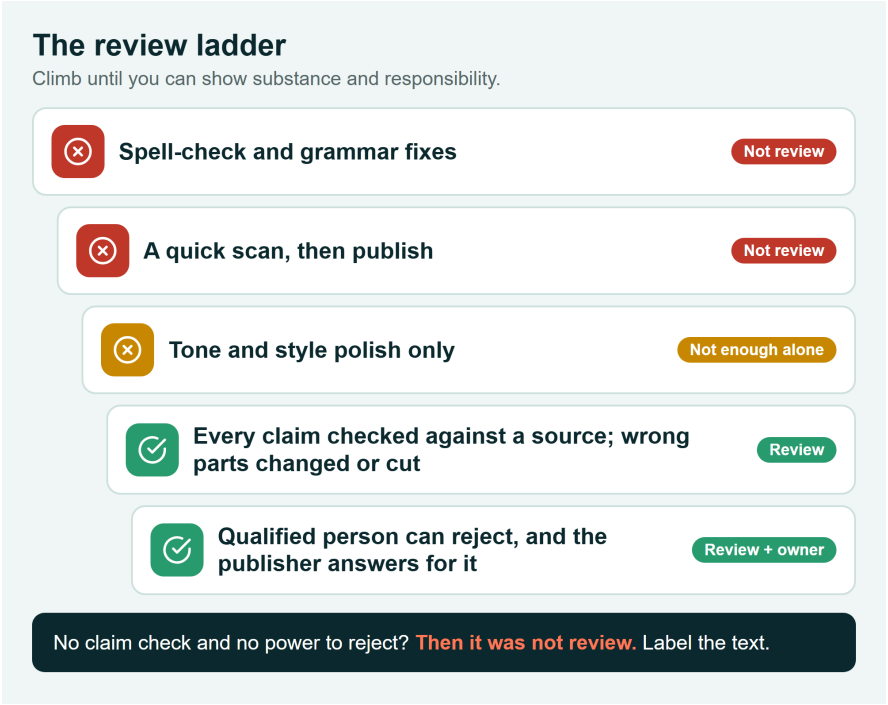

The review ladder: what counts and what does not

Rung	Is it review?
Spell-check and grammar corrections	No

Rung	Is it review?
Scanning it quickly and publishing	No
Reading for tone and fixing style	Not on its own
Checking every factual claim against a source, changing or removing what fails	**Yes**
A qualified person approves, can reject, and the publisher answers for it	**Yes, with responsibility**

When to label instead

Review is the preferred route for most businesses, but sometimes labelling is the better call:

- You publish at volume and cannot review each item to a substantive standard.
- The topic is sensitive and you lack in-house expertise.
- You want to be transparent as a matter of brand trust.
- The text is a **summary or draft** of third-party material and you cannot verify it.

In those cases, label. The Code's text-labelling approach is a plain statement visible at the start of the piece.

> **Say this.** "*This article was generated with AI and has not been reviewed by an editor.*"
>
> **Or, where reviewed but you still want to be open:** "*Drafted with AI, fact-checked and edited by [name], [role].*"

Be careful with the second version

The reviewed version is not required by law once real editorial responsibility exists. It is, however, honest and often helpful. Two cautions apply.

- **Never claim review that did not happen.** A false "reviewed by" is worse than no label. It is a misleading statement and it destroys the exemption.

- **Do not hide behind "AI-assisted".** If AI wrote the first draft and a human only polished the style, say what happened accurately.

Checklist: text

- ☐ Is it published to inform the public on a matter of public interest?
- ☐ Is any part AI-generated or materially AI-altered?
- ☐ If yes, has a named person reviewed it substantively, and does the publisher hold editorial responsibility?
- ☐ If not, is the label in place at the top?
- ☐ Is the claim "reviewed" or "fact-checked" true?
- ☐ Sources recorded for every factual claim

CHAPTER 6

Human Review and Editorial Responsibility

The human review exemption is powerful and fragile. It works only if you can prove that a person did what the guidelines describe. This chapter shows how to build a review system that survives a challenge.

The three things you must be able to show

If a regulator, client or platform asks why a public-interest article carried no AI label, you need three things.

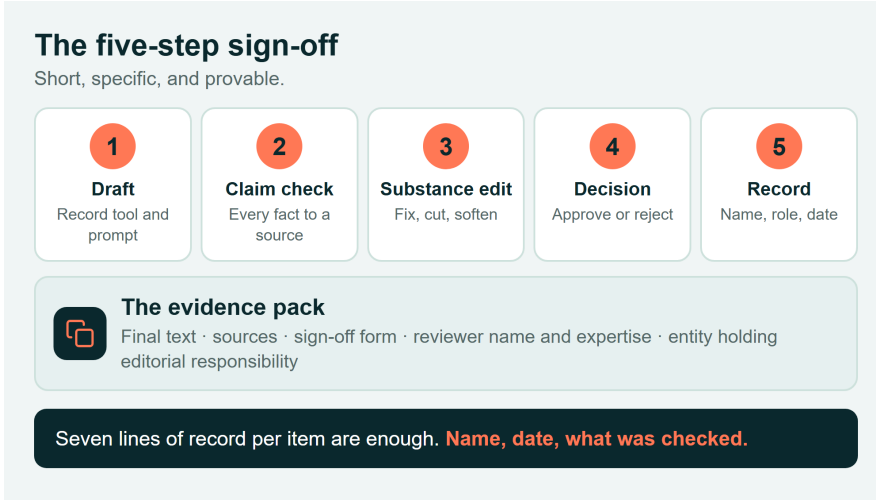

The sign-off path: draft, check, decide, publish, record

1. **Who** reviewed it. A named person, with the expertise to judge the content.
2. **What** they did. Checked facts against sources, edited the substance, or rejected parts.

3. **Who answers for it.** The legal entity or named person that holds editorial responsibility.

> **The rule.** An exemption you cannot evidence is not an exemption. Write the name, the date and what was checked.

The five-step review

A review that meets the standard in the guidelines is not long. It is specific.

Step 1. Draft. Whoever produces the text, human or machine, records which tool was used and what the prompt asked.

Step 2. Claim check. The reviewer lists every factual claim: numbers, dates, names, legal and medical statements, quotes. Each is checked against a primary or reputable source. Unverifiable claims are cut or softened.

Step 3. Substance edit. The reviewer changes what is wrong, thin or misleading. Cosmetic edits do not count.

Step 4. Decision. The reviewer approves, requests changes or rejects. They have the real power to say no.

Step 5. Record. A sign-off note with the reviewer's name, role, date, sources checked and any changes made. Store it with the published version.

Who qualifies as the reviewer

"Relevant expertise" is the phrase to take seriously. A marketing assistant is not an adequate reviewer of an article on a new employment law. A solicitor, accountant, clinician or experienced subject editor may be. For a small team, that often means one of three options:

- **The owner or an experienced editor,** if the topic is within their competence.
- **An external expert,** paid per review, for specialist topics.
- **No publication,** if no qualified reviewer is available. This is a legitimate outcome.

Fictional. Kestrel & Lane, a six-person agency (invented), publishes a weekly newsletter on workplace law for HR managers. Drafts come from an AI tool. The managing editor, Dara, reads each draft, checks every legal claim against the source legislation, rewrites two paragraphs most weeks and signs the sign-off form. A freelance employment solicitor reviews anything involving a new regulation. Kestrel & Lane does not label the newsletter as AI-generated, and its sign-off forms show why.

Responsibility has a name

Editorial responsibility is the end of the line: the person or organisation that would answer for the content. For a business, that is usually the company itself, and inside it a named lead. Write it into your policy.

> **Say this.** "*Published by [company]. Editor responsible: [name].*"

A byline that is a persona, a company mascot or an invented "expert" does not do this job. Do not create fictitious authors for AI text. Beyond the transparency rules, it is a recipe for a misleading-practice complaint.

Build the paper trail once

Create a one-page sign-off form (Appendix D) and a folder structure that keeps the draft, the sources and the sign-off with the final text.

Where the evidence lives: draft, sources, sign-off

A minimum record has seven lines:

1. Title and URL or channel
2. Date published
3. AI tool used and for what
4. Reviewer name and role
5. Sources checked
6. Changes made
7. Editor responsible

When the review fails

Sometimes the reviewer will find the draft is wrong. That is the system working. Options are to fix and re-review, to reject, or to publish with a label. Do not publish an unreviewed draft under a "reviewed" claim, and do not move a failed article to a different channel to avoid the check.

What to do about back catalogue

Content published before 2 August 2026 does not need retroactive labelling, although the Commission encourages it. [2] For older public-interest articles still driving traffic, a quick audit is worth doing: if AI wrote any, add a brief note or have a reviewer check and sign off today.

Checklist: human review

- ☐ A named reviewer with relevant expertise
- ☐ Every factual claim checked against a source
- ☐ Substantive changes made, not just style fixes
- ☐ A real power to reject
- ☐ A named person or entity holding editorial responsibility
- ☐ Sign-off recorded and stored with the published version
- ☐ No fictitious authors or invented experts

CHAPTER 7

Customer-Facing AI Transparency

When a customer talks to your business, they usually assume a person is on the other end. Article 50 turns that assumption into a rule: if they are talking to a machine, they should know.

The chatbot rule, and who carries it

Article 50(1) requires that **AI systems intended to interact directly with natural persons** are designed so that people are informed they are interacting with AI, **unless this is obvious** from the point of view of a reasonably well-informed, observant and circumspect person. [1] The duty sits with the **provider**, the company that builds the system. [2]

The Commission's FAQ sets four cumulative conditions for the rule to apply: an AI system, designed for genuine two-way exchange, interacting directly with a person, who is a natural person. Background systems and machine-to-machine exchanges are out of scope. The "obvious" exception is read restrictively. [2]

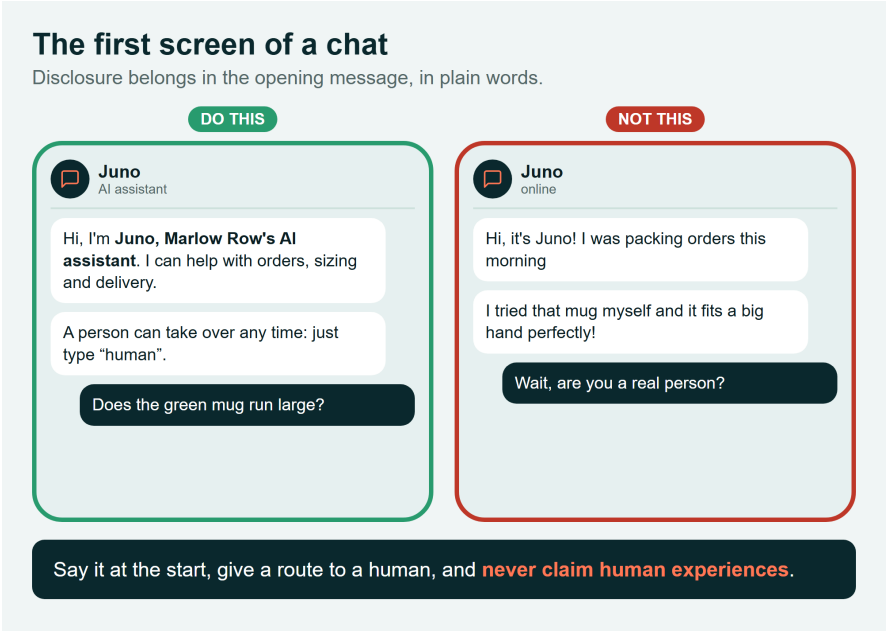

The first screen of a chat: say it in the opening message

As the business that puts the system in front of customers, you are the deployer. The duty is not yours in law, but your configuration can make the vendor's compliance work or break it. In practice:

- **Do not disable** a vendor's "AI assistant" notice.
- **Do not give the bot a human name, photo and sign-off** that suggest a person, without saying what it is.
- **Say it at the start** of the first interaction. The FAQ places the disclosure at the beginning of the first interaction, in line with accessibility requirements. [2]

> **Say this.** *"Hi, I'm Juno, Marlow Row's AI assistant. I can help with orders, sizing and delivery. A person can take over any time. Just type 'human'."*

Three common chatbot mistakes

The pretend human. Bots that say "I just tried it on", "I'll check with my colleague" or "I was in the warehouse this morning". Besides the transparency problem, these are false statements. A tribunal in British

Columbia held an airline responsible for what its own website chatbot told a customer, rejecting the argument that the bot was a separate entity. [15]

The buried notice. A grey line of small print under the input box, visible after the conversation starts. The disclosure must come at the beginning and be perceivable.

The dead end. An AI assistant with no route to a human. Not strictly an Article 50 issue, but it turns a transparency problem into a complaint.

Voice agents and phone lines

Voice systems need an audible disclosure in the first seconds. A silent banner on a website is no use to someone on a phone.

> **Say this.** "*You've reached Harbourside Dental. I'm an automated AI assistant. I can book or change appointments, or put you through to a person.*"

Keep it short, early and unskippable. If the call is recorded or transcribed, say so as well: that is a GDPR notice, and it sits beside the AI disclosure, not instead of it. [10]

Emails, messages and AI-drafted replies

Article 50(1) concerns systems that interact directly with people. An AI that *drafts* a reply, which a human reads and sends, is not interacting directly with the customer. An autonomous system that sends replies without human review is. If your support inbox is fully automated, treat it like a chatbot: tell people at the start of the thread, and give them a route to a person.

Emotion recognition and biometric categorisation

Article 50(3) requires deployers of **emotion-recognition** or **biometric-categorisation** systems to inform the people exposed to them. [1][2] Examples that small businesses meet include "sentiment" tools that analyse a caller's tone for stress or anger, camera analytics that infer shoppers' mood, and software that infers age or gender from faces.

Three cautions:

- **Some uses are prohibited.** The AI Act bans emotion recognition in workplaces and education settings, with narrow exceptions. [14] Do not use it to monitor staff.
- **Biometric data is special-category data** under the GDPR, with strict conditions for processing. [10]
- **Tell people before it starts,** not after a complaint.

Most small businesses should simply not use these tools. If you do, take specialist advice.

Reviews, testimonials and personas

The most serious customer-facing error has nothing to do with chatbots. It is using AI to create **fake reviews, fake testimonials or fake customers**.

The EU's consumer rules treat submitting or commissioning fake consumer reviews, or misrepresenting reviews, as an unfair practice prohibited in all circumstances. [9] A synthetic "customer" photo with an AI-written quote presented as a real customer falls squarely within it. A label does not make it acceptable. It must not be published.

What is acceptable:

- **Real reviews, edited for length only,** with the original kept.
- **Summaries of real reviews,** labelled as automatically generated and linked to the full reviews.
- **Fictional personas** used clearly as examples ("Meet Sam, a typical customer, a fictional character") and not presented as real people.

> **The rule.** AI may help you read, sort and summarise real customer feedback. It must never invent it.

AI in customer data and decisions

Where AI profiles customers, personalises prices or makes decisions that significantly affect them, the GDPR's transparency and automated decision-making rules apply in addition to Article 50. Update your privacy

notice to say that AI tools are used and for what. [10] Chapter 9 includes wording.

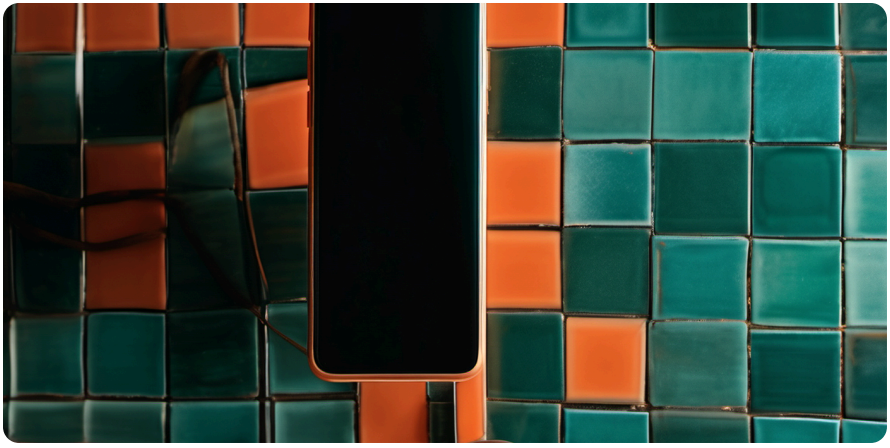

Your customer-facing map: chat, voice, email, reviews

Checklist: customer-facing AI

- ☐ Chatbot opens by saying it is AI
- ☐ Bot has no pretend-human name, photo or claims
- ☐ Voice agent discloses in the first seconds
- ☐ A route to a human is always available
- ☐ Fully automated email replies disclose AI at the start of the thread
- ☐ No fabricated reviews, testimonials or customer personas
- ☐ No emotion recognition in staff monitoring
- ☐ Privacy notice describes AI use

CHAPTER 8

Practical Disclosure Workflows

Knowing the rules is half the job. The other half is making sure that, every day, the right person asks the right question and writes down the answer. This chapter turns the previous chapters into a workflow you can run with a small team.

The five-step workflow

Every piece of content passes through five steps. The steps are the same whether you are a solo creator or a ten-person agency. Only the people change.

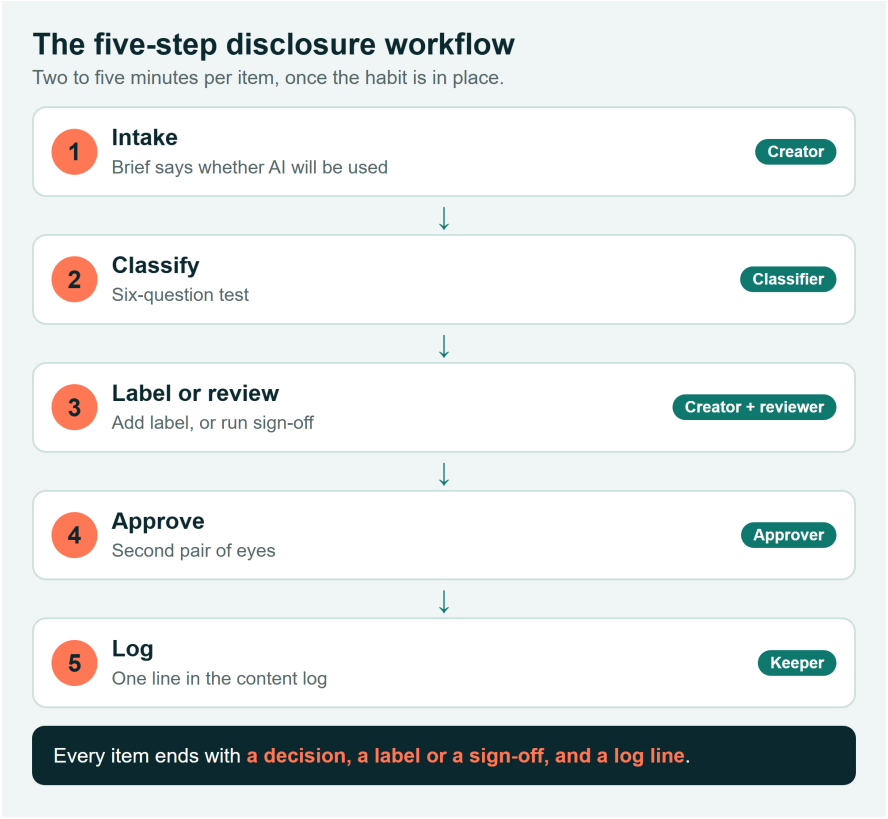

The five-step disclosure workflow

Step 1. Intake. When a brief arrives, record whether AI will be used and what for. Add a "Will AI be used?" field to your brief template. It is the cheapest control you can add.

Step 2. Classify. Run the six-question test from Chapter 2. The output is one of three decisions: label, review or record only.

Step 3. Label or review. Add the right label from the library (Chapter 9) or run the sign-off (Chapter 6).

Step 4. Approve. A named approver checks that the label is present and correct, and that the claims in the content are true. This is not the creator. A second pair of eyes is the point.

Step 5. Log. Write one line in the content log: what it is, where it is published, which AI tools were used, the decision, the label, the approver and the date.

Roles for a small team

You do not need a compliance department. You need four roles, which one person can combine in a very small business.

Role	Does what	Typical owner
Creator	Makes the content, records AI use	Designer, writer, producer
Classifier	Applies the test, picks label or review	Content lead
Approver	Checks label and claims before publication	Editor, account lead or owner
Keeper	Maintains the register, log and policy	Operations or owner

A solo creator holds all four. The discipline is to do them as four separate moments, on separate days if you can, not as one.

Tools that help

You do not need special software. A shared spreadsheet is enough for the register and the log. Three small additions make it sturdier:

- **Content credentials.** The open C2PA standard lets files carry tamper-evident information about how they were made. Some tools add it automatically. Keep it intact. [11]
- **IPTC metadata.** The photo-metadata standard includes a "digital source type" vocabulary for AI-generated media. Set it in your asset manager if you can. [11]
- **A label library** stored where everyone works: a shared document or a snippet tool, so nobody retypes the wording.

Workflow for agencies

Agencies carry the biggest exposure, because they sit between tools, freelancers and clients. Four additions make the workflow fit.

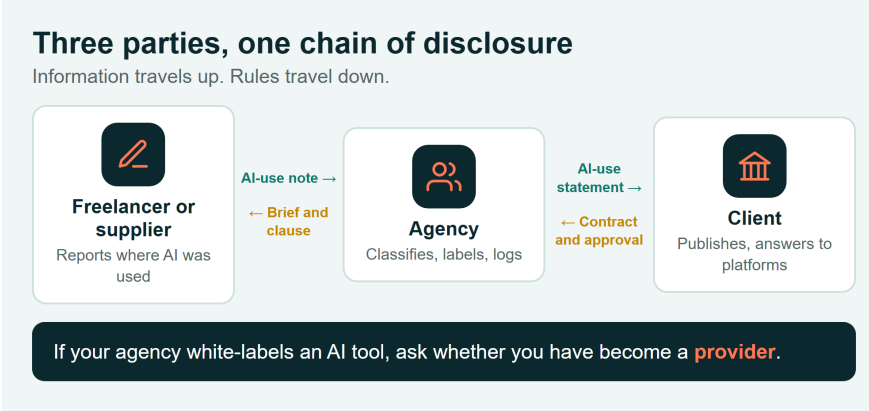

Three parties, one chain of disclosure

1. **Contract the disclosure.** The agency agreement should say which AI tools may be used, who classifies content, who adds labels, who holds editorial responsibility, and who keeps records. Appendix E is a checklist.
2. **Pass it down.** Freelancers and suppliers must tell you where AI was used. Make it a field on the invoice or hand-over form.
3. **Pass it up.** Give clients a short AI-use statement with each deliverable. They are often the deployer and need it for their own records.
4. **Clarify who is the provider.** If you white-label or package an AI tool for clients, take advice on whether you have become a provider.

> **Fictional.** Kestrel & Lane delivers a client campaign with six AI-generated hero images. The hand-over note lists each image, the tool used, the label applied and the log reference. When the client's platform asks whether the campaign contains synthetic media, the client forwards the note. The query closes in a day.

Workflow for online shops

Shops have more content and less time, so the workflow leans on rules rather than judgment.

- **Pre-classify recurring content.** Product descriptions are "record only". AI-generated model imagery is "label". Chat is "disclose at start". Write these defaults once.

- **Batch the log.** Log a campaign or a product range as one entry with a list of assets.
- **Guard the product truth.** Add a check to your launch list: *does every image show the real product?*
- **Reviews stay human.** Put a hard rule in the team handbook: no AI-written reviews.

Workflow for creators and solo brands

Solo creators have one advantage: a small, steady voice.

- **Put an AI statement on your site.** A short page that says where you use AI and where you do not. It is honest, useful to readers, and answers most questions in advance.
- **Keep a personal log.** One spreadsheet tab, five minutes a week, and the same short label each time.

Four failures to design out

Failure	How it shows up	The fix
Nobody owns it	Everyone assumes someone else decided	Name a Classifier and an Approver
Labels drift	Fifteen different phrasings across channels	One label library
Paper trail missing	Cannot show what was decided	The one-line log, every item
Label lost in reshare	Label is in the caption only	Put it in the file or the frame

Checklist: workflow

- ☐ "Will AI be used?" added to the brief template
- ☐ Six-question test pinned where the team works
- ☐ Four roles named
- ☐ Label library shared
- ☐ Content log set up and used for every item
- ☐ Freelancers report AI use in writing; clients get an AI-use statement

CHAPTER 9

The Label and Wording Library

Good wording is short, plain and specific. It tells the audience what they need to know, in the place where they will see it. This chapter is a library of wording you can adapt.

Principles

1. **Say what, not just that.** "Image generated with AI" is better than "AI was used".
2. **Plain words.** Use "AI-generated", "AI-edited" or "AI assistant". Avoid "synthetic", "algorithmically produced" and other jargon. Where an abbreviation is needed, "AI" is the accepted one. [6]
3. **First exposure.** At or before the first moment the viewer encounters the content.
4. **In the content.** Prefer a label inside the image, frame or audio over one in a caption.
5. **Accurate degree.** Say whether it is fully generated or only edited.
6. **No false comfort.** Do not claim "human-reviewed" or "fact-checked" unless it is true.

Label quick reference

Content	Where	Wording
Fully AI-generated image	Corner of the image, plus alt text	**AI-generated image**
AI-edited real photo	Caption or corner	**Photo edited with AI**
AI-generated person	Corner, plus caption	**AI-generated person, not a real customer**
AI video	First frame, or lower third	**AI-generated video**

Content	Where	Wording
AI-generated voice	Spoken at the start, plus notes	**This voice is AI-generated**
Cloned voice	Spoken at start	**This voice is an AI recreation of [name], used with permission**
Public-interest text, unreviewed	Top of the article	**This article was generated with AI and has not been reviewed by an editor**
Public-interest text, reviewed	Byline	**Drafted with AI, checked and edited by [name], [role]**
Chatbot	Opening message	**I'm [name], an AI assistant. A person can take over any time**
Voice agent	First sentence	**I'm an automated AI assistant**
Summary of reviews	Above summary	**AI-generated summary of [n] customer reviews. Read them all below**

Ready-to-use wording

Say this: product pages. "*Some images on this page were created with AI to show styling ideas. The product itself is photographed. See the product photos for the real item.*"

Say this: campaign footer. "*Visuals in this campaign were created with AI tools and edited by our team.*"

Say this: podcast intro. "*Our show notes and summaries are drafted with AI. Everything we say is written and recorded by us.*"

Say this: email footer for automated replies. "*This reply was written by an AI assistant. Reply 'person' to reach someone on our team.*"

Say this: site-wide AI statement. "*We use AI tools to help draft, illustrate and translate. We never use AI to write customer reviews, and we have a named editor for every article on regulation, health or finance.*"

Say this: privacy notice line. "*We use AI tools to help answer queries and personalise content. They process [categories of data] on the basis of [lawful basis]. You can ask for human review of any decision that significantly affects you.*"

Where to put the label

Label placement that survives a reshare

Format	Good placement	Poor placement
Image	Inside the image, in a corner with contrast	Caption only
Social video	Inside the first frame and the caption	Platform toggle only
Long video	Opening frame, repeated if the content changes	Description box only
Audio	Spoken at the start	Show notes only
Article	Top of the page, near the title	Footer
Email	Opening line or header	Small print at the end
Chat	First message	Placeholder text in the input box

What not to do

- **Do not hide the label** in a colour that matches the background.
- **Do not use jargon** nobody understands.
- **Do not label everything** equally. Over-labelling hides the labels that matter.
- **Do not copy the official icons** from memory. Download the artwork from the Commission's page. [6]
- **Do not translate sloppily.** If you publish in several languages, label in each one.

- **Do not claim "AI-free"** unless you can prove it. A false "100% human" claim is a misleading statement.

Accessibility

Labels are for everyone. Add alt text that says the image is AI-generated, give audio labels a spoken form, make sure contrast and size meet your normal accessibility standards, and avoid labels that appear for only a fraction of a second. [6]

Checklist: wording

- ☐ A written label for each content type
- ☐ Label library shared with every publisher
- ☐ Labels in each language you publish
- ☐ Alt text and spoken equivalents defined
- ☐ No claims of review, authorship or "AI-free" status that cannot be proved

CHAPTER 10

Your 30-Day Rollout and Three Case Files

You do not need to fix everything at once. You need a short sequence that gets the highest-risk items under control first and builds a habit.

Four weeks to a working system

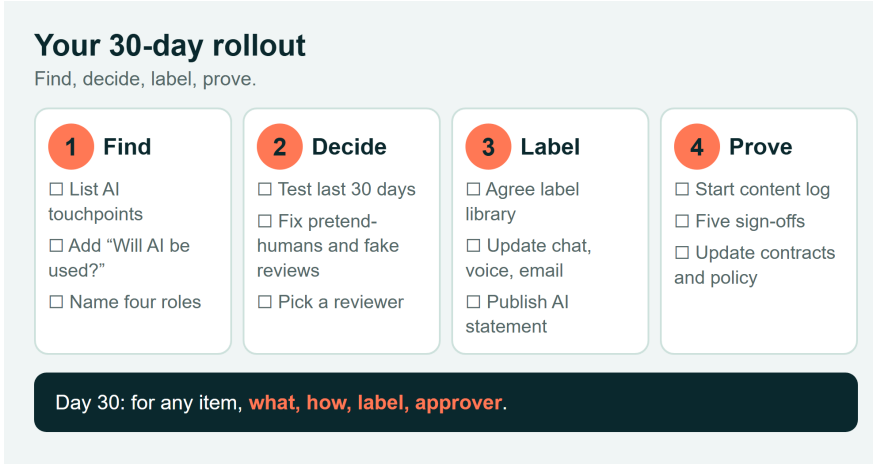

Four weeks: find, decide, label, prove

Week 1: Find

- ☐ List every place AI touches your content or your customers. Use Appendix C.
- ☐ Note, for each, who uses it and what it produces.
- ☐ Add "Will AI be used?" to the brief or content template.
- ☐ Name the four roles: creator, classifier, approver, keeper.

Week 2: Decide

- ☐ Run the six-question test over your last 30 days of published content. Mark each item: label, review or record only.
- ☐ Fix anything that is clearly wrong: pretend-human chatbots, unlabelled realistic synthetic people, fabricated reviews.
- ☐ Pick a human reviewer, with relevant expertise, for any public-interest content.

Week 3: Label

- ☐ Agree the label library (Chapter 9) and publish it where the team works.
- ☐ Update the chatbot opening, the voice greeting and the automated-email footer.
- ☐ Publish a short site-wide AI statement.
- ☐ Check that exports keep embedded marks and metadata.

Week 4: Prove

- ☐ Start the content log and use it for every new item.
- ☐ Fill in the sign-off form for the next five public-interest pieces.
- ☐ Update client and freelancer agreements with the AI-use clauses (Appendix E).
- ☐ Adopt the one-page disclosure policy (Appendix F).
- ☐ Put a recurring thirty-minute review in the calendar, monthly.

> **The rule.** At the end of thirty days you should be able to show, for any item you published, what it was, how it was made, what you told the audience and who approved it.

Case file 1: Marlow Row Ceramics (fictional)

> **Fictional.** Marlow Row Ceramics, its owner and customers are invented.

A seven-person online shop uses AI for lifestyle images, product descriptions, a chatbot called Juno and a monthly newsletter.

What went wrong. Juno chats with a human name and photo. Lifestyle images feature a smiling "customer" holding a mug, captioned with a first name and a quote. The newsletter's "Know Your Glaze" section is AI-drafted and published unchecked.

What the playbook does.

Problem	Test result	Fix
Juno hides that it is AI	Question 1: interaction	Opening message states it is an AI assistant, with a route to a person
Synthetic "customer" with quote	Questions 2 and 4: deepfake and fake review	Remove the quote and name; replace with real reviews or a clearly labelled illustration
Product descriptions	Question 6: record only	Logged as a default
Glaze explainer	Question 3: probably not public interest, but includes safety claims on food contact	Treated as in scope: reviewed by the owner against supplier documentation and signed off

Result. Ten minutes of changes, one default rule, a log line per campaign.

Case file 2: Kestrel & Lane (fictional)

> **Fictional.** Kestrel & Lane is an invented six-person agency.

An agency produces a regional tourism campaign: a video with an AI-generated narrator, AI-assisted still images of a local harbour, and a landing page explainer on new visitor tax rules.

The tests.

- The narrator is a synthetic voice reading an approved script. It resembles no real person. The agency labels it anyway, in the first seconds: "Narration generated with AI."
- The harbour images use a real harbour with generated boats and crowds. That resembles an existing place and appears authentic. The agency labels the images "AI-generated imagery" and tells the tourism client in its hand-over note.

- The visitor-tax explainer is public-interest text. A freelance tax adviser reviews and signs it off. The page byline names the client as responsible publisher and the adviser as reviewer.

The paper trail. A one-page AI-use statement went to the client with the files. When the client's regional partner asked about synthetic media a month later, the answer was one forwarded page.

Case file 3: Ines Falk, newsletter writer (fictional)

> **Fictional.** Ines Falk and her newsletter are invented.

Ines writes a weekly newsletter on personal finance for 12,000 subscribers. She uses AI for first drafts, headline ideas and a podcast version with a synthetic voice.

The risk. Personal finance is a public-interest topic. Drafts are AI-generated. She used to scan the draft and publish. The podcast voice sounds like a real presenter.

The change.

1. **Review standard.** She now lists every number and rule in the draft, checks each against the central bank, regulator or statistics source, and cuts what she cannot verify. She keeps a short sign-off note in a spreadsheet.
2. **Responsibility.** The newsletter footer reads "Written and edited by Ines Falk, who is responsible for its content."
3. **Podcast.** The first sentence of each episode says: "This episode is voiced by an AI. The script is written and checked by Ines."
4. **Site statement.** A short page explains how she uses AI and where she does not.

She did not add an AI label to the newsletter. She can show why it is not needed.

Conclusion: Decide, Label, Own, Record

The AI Act's transparency rule has a reputation for complexity. At the level of your daily work, it comes down to four verbs.

Decide. Before you publish, ask whether the content could mislead someone about who or what they are dealing with.

Label. Where the answer is yes, tell people clearly, early and in the content itself.

Own. Put a named human behind the content that matters, with the expertise and authority to stop it.

Record. Write down what you decided and why, so that you can show it.

Businesses that do these four things will find the rest of the rules, the guidelines, the Code of Practice and the platform policies, far easier to follow, because they are already working in the same direction. Businesses that skip them will spend their time reacting.

You do not have to be perfect. You do have to be deliberate. Pick one thing and do it today: rewrite your chatbot's first message, name your reviewer, check where an AI label goes, or add the AI-use clause to your next contract. Then start with the template in Appendix A and the next piece of content you publish.

APPENDIX A

Content Classification Sheet

Use one row per item, or one row per campaign. Copy into a spreadsheet.

Item	Format	AI tool and use	Q1 chat	Q2 deepfake	Q3 public-interest text	Q4 review	Q5 emotion	Outcome

Outcome key: L = label, R = review, N = record only.

Classifier name: _____ **Date:** _____

APPENDIX B

Label Quick Reference Card

Print this page and pin it where your team works.

If the content is	Label with
Fully AI-generated realistic image	AI-generated image
Real photo, AI-edited so it could mislead	Photo edited with AI
Realistic AI person	AI-generated person, not a real customer
AI video	AI-generated video
AI voice	This voice is AI-generated
Cloned voice	AI recreation of [name], used with permission
Public-interest text, no review	Generated with AI, not reviewed by an editor
Public-interest text, reviewed	Drafted with AI, checked and edited by [name]
Chatbot	I'm an AI assistant. A person can take over
Voice agent	I'm an automated AI assistant

Always: label at first exposure, inside the content where possible, in plain words, with alt text.

Never: invent reviews, invent experts, claim review that did not happen, strip marks and metadata.

APPENDIX C

AI Use Register

List every AI tool you use that touches content or customers.

Tool	Used for	Output type	Who uses it	Customer-facing?	Marks kept?	Owner

Register reviewed on: _____ **By:** _____

APPENDIX D

Human Review Sign-Off Form

One form per public-interest item.

Title and URL or channel:

_____

Date published: _____

AI tool used, and for what:

_____

Reviewer name and role:

_____

Reviewer's relevant expertise:

_____

Sources checked (list):

_____

_____

Claims changed, softened or removed:

_____

Decision: ☐ Approved ☐ Approved with changes ☐ Rejected

Editor or entity holding editorial responsibility: _____

Signature or approval reference: _____ **Date:** _____

APPENDIX E

Client AI-Use Clause Checklist

Use when drafting or reviewing agency, freelancer and client agreements. Have a lawyer adapt it.

- ☐ Which AI tools may be used, and for what
- ☐ A duty to tell the other party where AI was used in each deliverable
- ☐ Who classifies content and applies labels
- ☐ Who holds editorial responsibility for public-interest content
- ☐ Who keeps the log and for how long
- ☐ No removal of embedded marks or metadata
- ☐ No fabricated reviews, testimonials or experts
- ☐ Written consent and scope for any real person's likeness or voice
- ☐ Rules on confidential and personal data in prompts
- ☐ Ownership and intellectual-property position on generated output
- ☐ Who handles a platform query or regulator request
- ☐ Notice and correction process if a label is missing
- ☐ Whether the agency is acting as a provider of any AI system

APPENDIX F

One-Page AI Disclosure Policy Template

Adapt and adopt. Replace the bracketed items.

1. Purpose. [Company] uses AI tools to help create content and serve customers. We tell people when AI is involved where it matters, and we do not mislead them about who or what they are dealing with.

2. Scope. All content published in our name and all customer-facing systems.

3. The test. Before publication, the classifier applies the six-question disclosure test and records one of three outcomes: label, review, record only.

4. Labels. We use the approved label library. Labels appear at first exposure, in the content where possible, in plain words, with alt text or an audible equivalent.

5. Chat and voice. Our chat and voice assistants say they are AI at the start and offer a person on request.

6. Public-interest content. Text on regulation, health, finance or other public-interest matters is either labelled or reviewed by a named person with relevant expertise who holds editorial responsibility.

7. Reviews and people. We do not create fake reviews, fake customers or fictitious experts. We use a real person's likeness or voice only with written consent.

8. Marks and records. We do not remove embedded marks or metadata. We keep a content log and an AI-use register.

9. Roles. Classifier: _____ Approver: _____ Keeper: _____

10. Review. This policy is reviewed every six months, or when the rules or guidance change.

Adopted on: _____ **Signed:** _____

Glossary

AI-generated content. Text, images, audio or video produced or substantially altered by an AI system.

AI literacy. The skills and understanding that let people use AI systems competently and be aware of their risks. An AI Act duty for providers and deployers. [13]

Biometric categorisation. Assigning people to categories based on biometric data.

C2PA / Content Credentials. An open standard for attaching tamper-evident provenance information to media files.

Code of Practice. The voluntary Code on transparency of AI-generated content, in two sections: one for providers and one for deployers.

Deepfake. AI-generated or manipulated image, audio or video that resembles existing persons, objects, places, entities or events and would falsely appear authentic.

Deployer. An organisation that uses an AI system under its own authority in a professional capacity.

Editorial responsibility. Ultimate legal responsibility for what is published.

Emotion recognition. Inferring a person's emotions or intentions from biometric data.

EU icon. An optional visual label offered by the Code of Practice to signal AI involvement.

Human review. Deliberate, substantive examination of content by a person with relevant expertise who can approve, change or reject it.

Labelling. A human-perceivable disclosure that content is AI-generated or manipulated.

Marking. A machine-readable signal, such as a watermark or metadata, that content is AI-generated.

Matter of public interest. A topic of general concern, such as politics, public health, the economy, justice or consumer safety.

Provider. An organisation that develops an AI system, or has one developed, and places it on the market under its own name.

A Note on GDPRGard.eu

Disclosure sits beside a wider question: what does the GDPR expect of a small business? If this book has raised questions about privacy notices, records of processing or the data you put into AI tools, **GDPRGard.eu** offers free tools, guides and articles written in plain English for small organisations.

About the Author

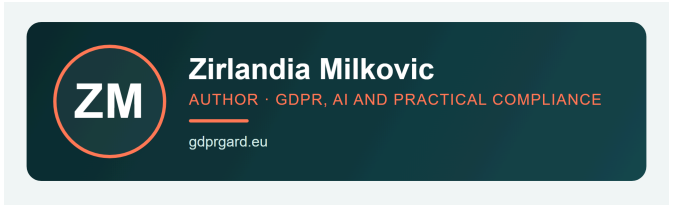

Zirlandia Milkovic

Zirlandia Milkovic writes practical, plain-English guides on data protection, AI and building a business with confidence. She believes privacy is not a barrier to ambition but a foundation of trust, and that business owners should not need a legal degree to understand the rules that apply to them.

She is the author of a series of business guides, including *GDPR Made Simple for Women Founders*, and her work focuses on turning complex European rules into simple steps, checklists and templates that busy people can actually use. She is also behind **GDPRGard**, a growing collection of free GDPR tools, guides and articles for small businesses, available at gdprgard.eu.

Visit **gdprgard.eu** for free tools, articles and updates, including notes on changes to the rules covered in this book. If this guide helped you, an honest review where you bought it helps other readers find it.

Notes and Sources

Research cut-off: 11 November 2026. Items marked (read) were consulted for this edition, and the note says how much of each. Others rest on the author's working knowledge and should be re-checked against the primary source before you rely on them.

1. Regulation (EU) 2024/1689 (AI Act), Article 50 (transparency obligations) and Article 3 (definitions, including provider, deployer and deepfake). eur-lex.europa.eu. Text not re-read for this edition; the substance was confirmed through the Commission pages below.
2. European Commission, FAQ on transparency obligations under Article 50 of the AI Act, last updated 24 July 2026 (read). digital-strategy.ec.europa.eu.
3. European Commission, Guidelines on transparency obligations for providers and deployers of AI systems, published 20 July 2026. The library page was read; the guidelines document itself was not, and its contents are described here as summarised in the Commission's FAQ and icon pages.
4. European Commission, Code of Practice on transparency of AI-generated content, policy page, last updated 19 October 2026 (read). The final Code was published on 10 June 2026. The Code document and the adequacy opinion were not read.
5. European Commission, FAQ on the Code of Practice on transparency of AI-generated content, last updated 10 June 2026 (read).
6. European Commission, EU icons for labelling AI-generated content, last updated 13 October 2026 (read). The per-format placement specifications are in Section 2 of the Code, which was not read; the placement figure in this book is the author's illustration of the general principles.
7. AI Act, Article 99 (penalties): up to EUR 15 million or 3% of worldwide annual turnover for breaches of Article 50 (figures as stated in the Commission FAQ). The treatment of SMEs under Article 99 should be checked in the Official Journal text.
8. Digital Omnibus on AI. The Commission pages read for this edition mention no change to the Article 50 dates beyond the transition to 2 December 2026 for Article 50(2) marking; a legislative amendment may exist, so confirm the current dates in the Official Journal before relying on them.
9. Directive 2005/29/EC (unfair commercial practices) and Directive (EU) 2019/2161, which added fake and unverified consumer reviews to the practices prohibited in all circumstances (working knowledge).
10. Regulation (EU) 2016/679 (GDPR), including Articles 5(1)(a), 9, 12 to 14 and 22.

11. C2PA, Content Credentials technical specification (c2pa.org); IPTC, "digital source type" vocabulary for AI-generated media (iptc.org). Working knowledge.
12. Help and policy pages of major social and video platforms on labelling realistic AI-generated content. These change frequently; check the current version for each platform you use.
13. AI Act, Article 4 (AI literacy), applicable from 2 February 2025.
14. AI Act, Article 5(1)(f) (prohibition of emotion recognition in the workplace and in education, with exceptions), applicable from 2 February 2025.
15. Moffatt v. Air Canada, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal, 14 February 2024). Working knowledge.