EU AI ACT · ARTICLE 50 · PRACTICAL GUIDE

THE AI DISCLOSURE TRAP

When Must Your Business **Tell Customers** That AI Is Involved?

CHATBOTS · AI IMAGES · AI TEXT · REVIEWS

ZIRLANDIA MILKOVIC

EU AI ACT · ARTICLE 50 · PRACTICAL GUIDE

THE AI DISCLOSURE TRAP

When Must Your Business Tell Customers That AI Is Involved?

ZIRLANDIA MILKOVIC

The AI Disclosure Trap
When Must Your Business Tell Customers That AI Is Involved?

© 2026 Zirlandia Milkovic. All rights reserved. No part of this book may be reproduced without written permission, except for brief quotations in reviews.

First edition, 2026. Research cut-off: 5 November 2026. British English. Companion resource: GDPRGard.eu.

This book is practical guidance, not legal, security or data-protection advice. All examples, businesses and people in worked scenarios are fictional. Cover and illustrations created by the author with design and AI image tools.

Contents

START HERE

Educational Disclaimer 4
Introduction: The Small Print Nobody Wrote 4
Who This Book Is For 5
How to Use This Book 6
The One-Page Disclosure Check 7

CHAPTERS

1. The Disclosure Trap 9
2. The Map of Duties 13
3. Chatbots, Voice Bots and Virtual Assistants 18
4. Images, Video and Audio: Deepfakes, Models and Cloned Voices 22
5. AI-Written Text: Articles, Product Pages, Newsletters and Translations 28
6. Marketing, Reviews, Testimonials and AI Influencers 32
7. Agencies, Freelancers and Creators: Clients, Contracts and Credit 37
8. People, Emotions and Recorded Meetings 41
9. The Wording Library and the Labelling Workflow 45
10. Your 30-Day Plan and Five Case Files 51

CLOSING

Conclusion: Say It Plainly 54

APPENDICES

Appendix A: AI Use and Disclosure Register 55
Appendix B: Disclosure Decision Sheet 56
Appendix C: One-Page AI Disclosure Policy Template 57
Appendix D: Content Labelling Log 58
Appendix E: Client AI-Use Clause Checklist 59

BACK MATTER

Glossary 60
A Note on GDPRGard.eu 60
About the Author 62
Notes and Sources 63

Educational Disclaimer

This book is practical, educational guidance for small and medium-sized businesses, agencies, online retailers and content creators. It is not legal advice, and it does not create a professional relationship between the author and the reader.

The EU AI Act, consumer law, advertising rules and data-protection law differ by country, apply differently to different businesses, and are still being clarified by guidelines, codes of practice and regulators. Dates and details changed several times in 2026. Check the current position with a qualified adviser or your regulator before relying on anything here for a decision that matters.

Following the steps in this book may reduce your risk. It cannot guarantee that you will meet every legal duty. Checklists, wording and templates are starting points to adapt, not finished legal documents.

Every business, person and product used in a worked example is **fictional** and is labelled as such. Where the book describes a real case or rule, it is sourced in the notes at the back. Where it gives a recommendation, it says so.

Introduction: The Small Print Nobody Wrote

Somewhere in your business, a machine is probably speaking, drawing or writing on your behalf.

It may be the chat window in the corner of your website. It may be the model in your product photos, the voice that answers your phone out of hours, the caption on your latest post, the first draft of a client's blog, the summary of a customer's review. Each of these is useful. Each of them raises a question that, until recently, nobody had to answer:

Does the person on the other end need to be told?

For years the honest answer was "it depends, and nobody is checking". In 2026 that changed. Article 50 of the EU AI Act, the transparency rule, has applied since **2 August 2026**. It says, in outline, that people must be told when they are talking to a machine, that certain synthetic images, audio and video must be labelled, and that some AI-written text published to inform the public must be marked. Around it sit older rules that were never switched off: consumer law on misleading practices, the ban on fake reviews, advertising standards and the GDPR's transparency duties.

The result is a trap with two jaws. In the first, a business says nothing because "nobody told us we had to", and discovers it was wrong. In the second, a business panics and stamps "AI-generated" on everything, including things that need no label, which confuses customers, undermines its own brand and, in some cases, still misses the one disclosure that mattered.

This book is about walking between the jaws. Its idea is simple:

> **The rule.** Disclose when a reasonable person would otherwise be misled about who, or what, they are dealing with. Do it clearly, early and in plain words. Then keep a record that you did.

That is a working rule, not a legal definition. The chapters that follow turn it into specifics: which tools trigger which duties, what to say, where to say it, who in your business is responsible, and how to prove you did it.

You do not need a lawyer to build a sensible disclosure habit, although you may want one for the grey areas. You need a list of where AI touches your customers, a handful of standard phrases, and a person who owns the job. All three are in this book.

Who This Book Is For

This book is for you if you run, manage or advise a small or medium-sized organisation in Europe or the UK and any of the following is true:

- You run an **online shop** that uses AI for product descriptions, product images, chat support or reviews.

- You work in or run an **agency** and produce copy, images, video or voiceovers for clients, with or without AI tools.
- You are a **content creator, newsletter writer, podcaster or consultant** who uses AI to draft, illustrate, translate or edit.
- You use a **chatbot, voice assistant or AI receptionist** with customers.
- A client, platform, insurer or regulator has started asking, "Did you use AI for this?"

It assumes no technical background. Technical terms are explained when they first appear and collected in the glossary.

It is **not** for organisations that build and sell general-purpose AI models or AI systems to others. They carry heavier duties, including technical marking of outputs. Treat this book as an introduction if that is you, and take specialist advice.

How to Use This Book

You can read the book in an evening, or use it as a manual. Chapters 1 and 2 explain the trap and the map of duties. Chapters 3 to 8 go channel by channel: **chatbots, images and video, text, marketing and reviews, agency and creator work, and people-facing AI**. Chapter 9 is a wording library and a labelling workflow. Chapter 10 is a 30-day plan and five fictional case files.

Look for these markers:

- **The rule.** A principle worth remembering, set in a dark panel.
- **Fictional.** A worked example using invented people and businesses.
- **Say this.** Wording you can adapt.
- Check boxes (☐). A step you can tick off. Where a line is blank, write on it.

The appendices contain templates: a disclosure register, a decision sheet, a content labelling log, a one-page disclosure policy and a client AI-use clause checklist. Copy them into a spreadsheet or print them.

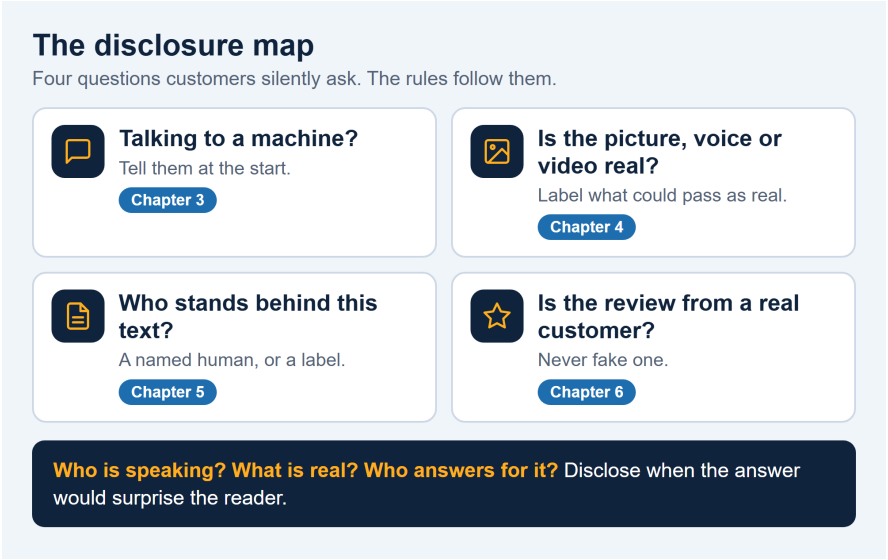

The disclosure map: who must be told what

The One-Page Disclosure Check

Before you read further, answer these ten questions about your business. Every "yes" points to a chapter.

1. ☐ Do customers chat or speak with an AI system in your name? (*Chapter 3*)
2. ☐ Is it always obvious that it is a machine? (*Chapter 3*)
3. ☐ Do you publish AI-generated or AI-edited images, video or audio that look real? (*Chapter 4*)
4. ☐ Do any of them show a person, place or event that did not happen as shown? (*Chapter 4*)
5. ☐ Do you publish AI-written text about news, health, law, politics or other public-interest topics? (*Chapter 5*)
6. ☐ Does a named human review and take responsibility for it before it goes out? (*Chapter 5*)
7. ☐ Do you use AI to write, summarise or generate reviews, testimonials or personas? (*Chapter 6*)
8. ☐ Do you produce AI-assisted work for clients, or use AI-assisted work from suppliers? (*Chapter 7*)

9. ☐ Do you use AI that listens to, records or reads emotion or characteristics of people? (*Chapter 8*)
0. ☐ Could you show a regulator or client, today, what you disclosed and when? (*Chapter 9*)

Questions 2, 6 and 10 are the ones where a "no" is the problem. If you answered "no" to any of them, or "yes" to several of the others, start with the 30-day plan in Chapter 10, then read the chapters your answers point to.

CHAPTER 1

The Disclosure Trap

> **Fictional.** Cobble & Thread, its owner Mara and every person below are invented for teaching. They are not real businesses or incidents.

Cobble & Thread is a nine-person online shop selling linen clothing. Mara, the founder, uses AI for three things. A chatbot called "Poppy" answers sizing and delivery questions. An image tool creates the lifestyle photos, with models who do not exist wearing garments that do. And a writing assistant produces the first draft of every product description and the weekly newsletter.

None of this was hidden. None of it was announced either. Poppy has a friendly photo avatar and signs off with "Speak soon!". The models have names in the captions. The newsletter is signed by Mara.

In October, a customer emails: *"I asked Poppy whether the green dress ran large. She said she had 'tried it on' and it fits true to size. Is Poppy a real person? Were those models real? Did you write that review of your own dress?"*

Mara had assumed three things. That nobody would care. That if someone did, a small shop would be left alone. And that "AI disclosure" was a rule for big tech. Each assumption was about half right, which is exactly why it is dangerous.

A packing desk, a laptop, and a handful of tags that say nothing

Two ways to fall into the trap

The disclosure trap has two jaws, and they snap on opposite kinds of business.

Jaw one: silence. The business assumes that because nobody asked, nothing is required. Poppy says she "tried it on". The product photos present a person who never existed as a satisfied customer. The newsletter looks like personal expertise but was drafted by a tool and published unchecked. A customer who feels deceived does not need Article 50 to complain. Consumer law on misleading practices has been enough for years. Article 50 simply makes the expectation explicit and the regulators' job easier.

Jaw two: noise. The business, alarmed by headlines, adds "Made with AI" to every page, every image, every email footer. The label stops meaning anything. Customers learn to ignore it. Real deepfake-style content ends up with the same tiny badge as a spell-checked caption. And the business still misses the disclosure that mattered, because it was lost in the clutter.

Both jaws have the same cause: **no one decided, on purpose, where disclosure belongs.**

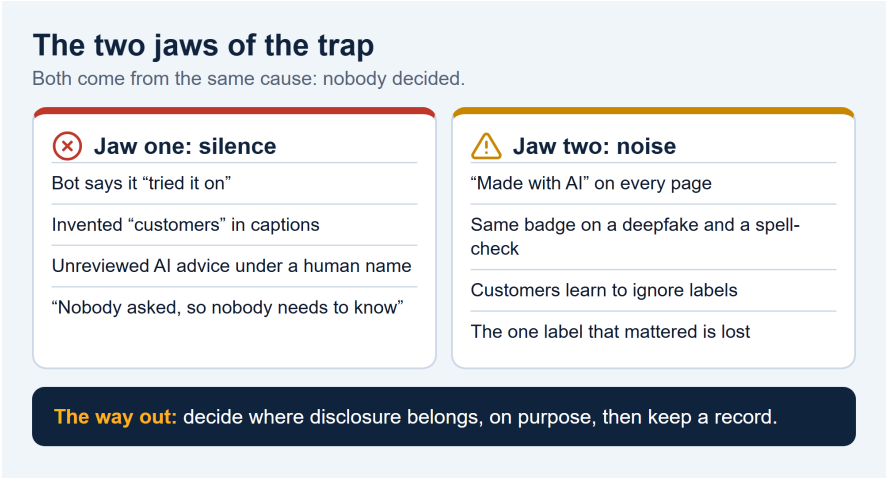

Silence and noise are both failures of design

What a customer actually wants to know

Strip the law away and ask what a reasonable customer is trying to find out:

- **Am I talking to a person or a machine?** Because I will trust, tolerate and share differently.
- **Is this picture, voice or video real?** Because I may rely on it to decide what to buy, believe or share.
- **Who stands behind this text?** Because advice and information are only as good as whoever answers for them.
- **Is this review or testimonial from someone who actually bought the product?** Because it is the thing I am about to spend money on.

These four questions are the real structure of the rules in the next chapter. If you can answer them honestly, in the places where customers look, you are most of the way to compliance.

What this book will not do

It will not tell you to stop using AI. Used well, AI helps a five-person shop publish like a company of fifty. It will not give you a label to paste everywhere. And it will not promise that following it keeps you safe from every regulator or complaint. It will give you a method, decided on purpose, that you can explain.

> **The rule.** Disclosure is not a confession. It is information a customer needs at the moment they need it.

CHAPTER 2

The Map of Duties

Before choosing wording, you need to know which rules point at you. There are four disclosure duties in Article 50 of the AI Act, and the first question for each is who carries it: the **provider** (the organisation that builds or places an AI system on the market) or the **deployer** (the organisation that uses an AI system in a professional capacity).

For most small businesses the answer is simple. You are the **deployer** of tools built by someone else: a chatbot platform, an image generator, a writing assistant. Your duties are the deployer ones. You become a **provider** only if you build your own AI system, or put your own name and branding on someone else's and sell it on as yours. That is the scenario to watch if you are an agency or developer.

The four duties of Article 50

Most small businesses are deployers: the highlighted cards land on you.

ARTICLE 50(1)
Chatbots and voice bots

Built so people are told they are dealing with AI, unless obvious.

Carried by the provider

ARTICLE 50(2)
Machine-readable marking

Synthetic audio, image, video and text detectable as AI-made.

Carried by the provider

ARTICLE 50(3)
Emotion and biometric tools

Tell people exposed that the system is operating.

YOU, THE DEPLOYER

ARTICLE 50(4)
Deepfakes and public-interest text

Label deepfakes; label AI text unless human-reviewed.

YOU, THE DEPLOYER

You become a **provider** if you build your own AI system or sell someone else's under your own name.

The four duties in plain terms

Duty	Who carries it	In plain words
Art. 50(1)	Provider	Build AI that talks to people so they are told they are dealing with a machine, unless that is obvious
Art. 50(2)	Provider	Mark synthetic audio, image, video and text so that machines can detect it as AI-generated
Art. 50(3)	Deployer	Tell people when an emotion-recognition or biometric-categorisation system is used on them
Art. 50(4)	Deployer	Label deepfakes, and label AI-generated text published to inform the public on matters of public interest, unless a human has reviewed it and taken editorial responsibility

A few details matter for a small business.

The chatbot duty (50(1)) sits with the provider, not you. If you buy a chatbot, the vendor should build in a disclosure. But as the business putting it in front of your customers, you still need to configure it honestly and not defeat it. A chatbot that says "I'm Poppy, I tried it on" while a vendor banner says "AI assistant" is a mess you made. Chapter 3 covers this.

The marking duty (50(2)) is technical and sits with the provider. It means watermarks and metadata in the output. As a deployer you do not have to build them, but you should not strip them out. Deleting metadata from AI images before posting is a bad habit to drop.

The deployer duties (50(3) and 50(4)) are the ones that land on you. These are the labels you add to deepfakes and to AI-generated text for the public, and the notice you give people subject to emotion recognition. They are the heart of Chapters 4, 5 and 8.

Exceptions are narrow. The Commission's FAQ[2] says the "obvious to a reasonable person" exception for chatbots is read restrictively. Evidently artistic, creative, satirical or fictional works have a lighter labelling duty that must not spoil the work. Human review and editorial responsibility removes the text-labelling duty. Personal, non-professional use by individuals is outside the rules. Chapters 3 to 5 explain each in turn.

Dates and penalties

Article 50 has applied since **2 August 2026**.[1][2] There is one transitional period: generative AI systems **already on the market before 2 August 2026** have until **2 December 2026** to comply with the machine-readable marking duty in 50(2). That grace period is for providers and covers marking only. It does not postpone the labels you add as a deployer.[2][8]

Content you generated before 2 August 2026 does not need retroactive labelling, although the Commission encourages it.

Breaches of Article 50 can attract fines of up to **€15 million or 3% of worldwide annual turnover**. For SMEs, including start-ups, the AI Act says the lower of the two figures applies, and regulators are told to consider proportionality.[9] Enforcement is mainly by national market surveillance authorities. Expect a first conversation to be a request for information rather than a fine, but do not plan around that.

Soft law you should know

Two documents sit beside the Article. The Commission has published **guidelines** on the transparency obligations that explain the terms and give examples. It has also confirmed a voluntary **Code of Practice on transparency of AI-generated content**, which helps providers and deployers show they comply with Article 50(2), (4) and (5).[10] The Code is not mandatory. Signatories get more legal certainty. If you do not follow it, you need another way to show you comply.

The older rules still apply

Article 50 is not the only source of disclosure duties. Four older layers stay in force and are often the ones that actually bite a small business.

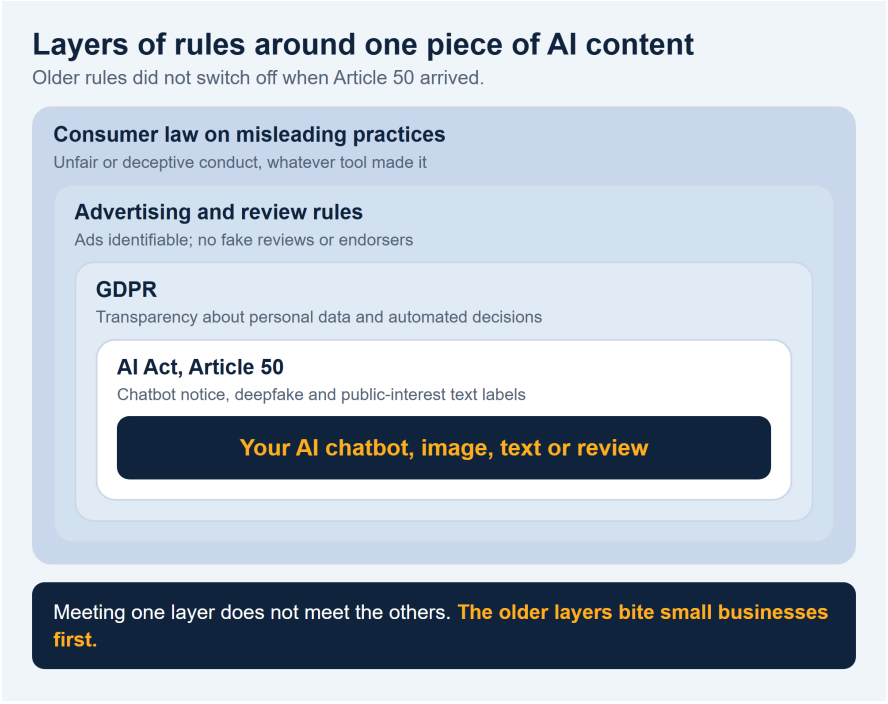

Layers of rules around one piece of AI content

1. **Consumer law on unfair and misleading practices.** In the EU, Directive 2005/29/EC prohibits practices that deceive an average consumer or hide material information.[11] In the UK, the Digital Markets, Competition and Consumers Act 2024 now puts the equivalent regime on a statutory footing with direct enforcement powers.[12] A fake "customer" who does not exist can be misleading whether or not any AI Act label was required.
2. **Fake and undisclosed reviews.** EU law bans submitting or commissioning false consumer reviews, and the UK has a similar ban.[11][12] Chapter 6 covers what that means when a tool can write a thousand reviews in a minute.
3. **Advertising rules.** Advertising codes and influencer rules require ads to be identifiable as ads. AI does not change that, but an AI-generated "testimonial" in an ad has to be true to be fair.
4. **The GDPR.** If AI processes personal data, you owe people transparency about it. If AI makes decisions about people, Article 22 and the duty to explain the logic may apply.[13] Chapter 8 covers this.

The UK, Switzerland and others

The AI Act applies to businesses outside the EU when their AI output is used in the EU. A UK shop selling to Germany can be in scope. The UK itself has no equivalent of Article 50, but it has consumer, advertising and data-protection law that reaches much of the same ground. This book flags where the UK differs. If you sell across borders, assume the stricter rule.

A working summary

If you...	Then think about...
Use a chatbot or voice bot with the public	Telling people it is a machine, at the start (Ch. 3)
Publish realistic AI images, video or audio	Labelling anything that could pass for real events or people (Ch. 4)
Publish AI text on topics of public interest	Human review, or a label (Ch. 5)
Use AI for reviews, personas or ads	Not faking customers (Ch. 6)
Work for clients	Clauses and a labelling hand-over (Ch. 7)
Analyse faces, voices or emotions	Notice, a lawful basis and often a stop sign (Ch. 8)

The rule. Ask who is speaking, what is real, and who answers for it. If the answer would surprise the reader, disclose.

CHAPTER 3

Chatbots, Voice Bots and Virtual Assistants

Of all the places AI meets customers, the chat window is where disclosure matters first. The question is the oldest one: *am I talking to a person?*

What the rule says

Article 50(1) requires AI systems that interact directly with people to be designed so that those people are informed they are dealing with an AI system, unless that is obvious to a reasonably well-informed, observant and circumspect person.[1][2] The Commission's FAQ applies four tests: it is an AI system, built for genuine two-way exchange, communicating directly with a natural person rather than through a human intermediary.[2] Background systems and machine-to-machine exchanges are outside it. The "obvious" exception is read restrictively.

The duty sits with the provider. But two things stay with you as the business using it:

1. **You configure how the bot presents itself.** If the platform lets you give it a human name and photo, those choices are yours.
2. **You decide where it is placed.** A chat widget on your site, a WhatsApp number, a phone line, a social-media DM bot. Each is a different context in which "obvious" may or may not be true.

Same question, two chatbots

Disclosure and honest behaviour work together.

✗ Hides what it is

Hi! I'm Sarah. How can I help?

Does the green dress run large?

I tried it on last week and it fits true to size!

A human name, no disclosure, and a made-up personal experience.

✓ Says what it is

Hi, I'm Poppy, Cobble & Thread's **AI assistant**. I can help with sizes and delivery. Ask for a person any time.

Does the green dress run large?

I can't try clothes on, but customers' reviews say it fits true to size. Want the size chart?

Disclosure in the first message, no false experience, a human route.

A chatbot that tells the truth in its first message

A reception desk where a voice answers before a person does

"Obvious" is a trap word

A grey robot icon and a bot name do not make it obvious to everyone. Customers on a phone with a small screen, older customers, someone who has never used chat support: all are "reasonable persons". A good test is the **worst-case customer**: if the one most likely to be fooled would not realise they were talking to software, then say so.

Voice is harder than text. Modern voice bots sound human. A phone call has no sidebar, no icon, no scrolling back. The disclosure has to be spoken, in the first sentence. *"Hello, this is Harbour Dental's virtual assistant. I'm an AI, and I can take a booking or pass you to a person."*

> **Say this.** "Hi, I'm Poppy, Cobble & Thread's AI assistant. I can help with sizes and delivery. For anything else I'll pass you to a person."

What good looks like

- **In the first message,** not buried in a footer or a privacy policy.
- **Plain words:** "AI assistant", "virtual assistant", "automated". Not "digital team member" or "smart helper".
- **A name that does not pretend.** Poppy is fine. "Sarah from customer care", with a stock photo of a woman, is not.
- **No false claims of human experience.** A bot must not say "I tried it on", "I've been to your area" or "my colleague told me".
- **A human route.** Offer to hand over, and make it work.
- **Accessibility.** The disclosure must be available to screen-reader users and appear in the language the customer is using.

Human in the loop: when a person types the final message

Many small businesses use AI to draft replies that a person reviews and sends. If a human reads, edits and sends each message, the interaction is with a human, assisted by a tool. Article 50(1) is not triggered by AI-assisted drafting where a person genuinely communicates. The line moves the moment the human stops reading and the draft goes straight out. At that point the AI is talking and the customer should be told.

The Air Canada lesson

The best-known example of a business blaming its own chatbot is *Moffatt v. Air Canada* (2024). A customer relied on the airline's website chatbot for bereavement-fare information. The tribunal rejected the airline's suggestion that the chatbot was a separate entity responsible for its own statements, and ordered it to pay.[3] The amounts were small. The lesson is general: a disclosure that says "this is an AI" does not protect you from what the AI promises. Disclosure and accuracy work together.

Where chatbots hide

- **Website chat widgets** and in-app assistants.
- **WhatsApp, Instagram and Messenger auto-replies,** where the bot uses your business name.
- **Phone bots and IVR with natural-language speech.**
- **Email auto-responders** that generate individual answers.
- **Booking and quote tools** that "chat".
- **Voice and avatar assistants** in kiosks or shop displays.

Add each to your register (Appendix A). If you are not sure whether a tool counts as AI, ask the vendor in writing.

Checklist for a chatbot

- ☐ The first message says it is an AI assistant, in plain words.
- ☐ The name, avatar and tone do not suggest a real person.
- ☐ It never claims to be human or to have human experiences.
- ☐ A route to a person works and is offered.
- ☐ Voice bots say so in the first sentence.
- ☐ The prompt and settings have been reviewed by a named owner.
- ☐ Screenshots of the opening message are saved with the date.

> **The rule.** If a customer might wonder, they should not have to ask.

CHAPTER 4

Images, Video and Audio: Deepfakes, Models and Cloned Voices

Synthetic images and video are where small businesses most often stumble, because the tools are cheap, the results look excellent and the line between a "creative visual" and a "deepfake" is not where most people expect it.

What a deepfake is, in the law

The AI Act defines a deepfake as AI-generated or manipulated image, audio or video content that **resembles existing persons, objects, places, entities or events** and would **falsely appear to a person to be authentic or truthful**.[1] The Commission's FAQ treats these as three cumulative tests: resemblance to something that exists, existence of a real counterpart, and a false appearance of authenticity.[2] Context and the audience's expectations are part of the assessment.

That definition is narrower than "any realistic AI picture" and broader than "a famous person saying something they did not say". It includes:

- A realistic video of a real person doing or saying something they never did.
- A photo-real image of a real building, landmark, store or event that never looked like that.
- A cloned voice of a real person endorsing something.
- A fake "event" photo, such as a crowd at an opening that never took place.

It does not automatically include:

- An invented person who resembles no real individual, such as a purely synthetic fashion model. That is not resembling an existing person, although it may still mislead for other reasons (below).
- A clearly fantastical, stylised or illustrated image.
- Standard editing: colour correction, retouching, cropping, background clean-up that does not alter the substance.

The duty: label deepfakes

If you are the deployer, you must disclose that deepfake content has been artificially generated or manipulated.[1] The Commission says the disclosure must be clear, distinguishable, available at the latest on first exposure, and understandable without special tools. A hidden metadata tag the provider added does not satisfy your duty on its own.[2]

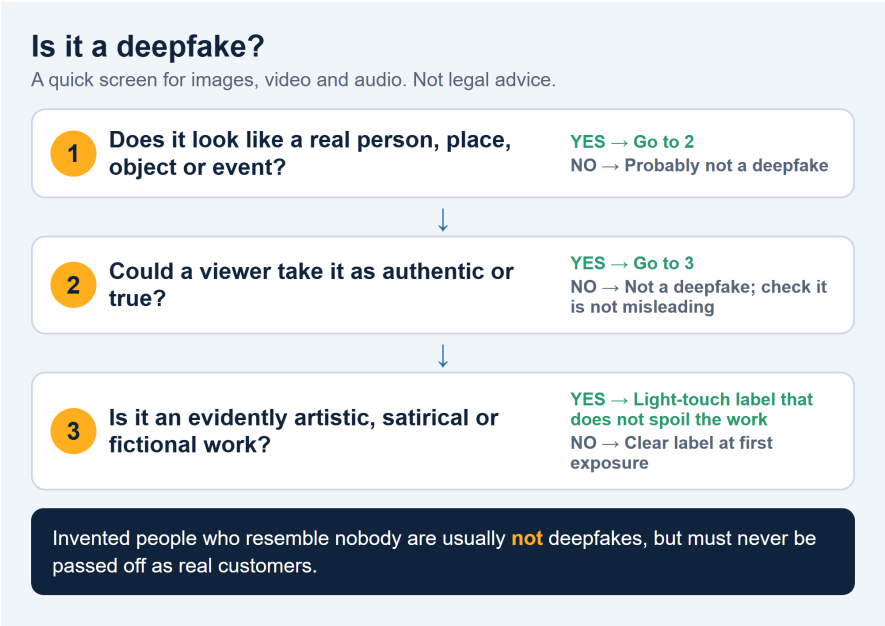

Is it a deepfake? A quick screen for images, video and audio

Artistic, creative, satirical or fictional works have a lighter duty. You must still disclose, but in a way that does not spoil the display or enjoyment of the work, for example a small caption rather than a banner across the picture.

A studio where nothing in the frame is quite what it seems

The invented model problem

Cobble & Thread's AI models raise a different issue. An invented person is probably not a "deepfake" in the legal sense. But there are other ways to get this wrong.

- **Misleading about the product.** If the AI-generated garment looks different from the real one (colour, drape, fit), the image misleads regardless of any label.
- **Misleading about who wore it.** A caption such as "Our customer Amelia in the green dress" when Amelia does not exist is a false testimonial.
- **Diversity-washing.** Showing a range of invented models that you do not serve or stock is a trust problem even if it is not a legal one.
- **Looks like a real person.** If the generator produces a face that resembles a real, identifiable person, you may have a deepfake and a personality-rights problem.

> **Say this.** "Images on this page were created with AI. Colours and fit are shown from our real product photos below."

The cleanest approach: use AI imagery for mood and context, show the real product accurately, and say so.

A real warning: when the picture sold the dream

In February 2024 a Glasgow event marketed as "Willy's Chocolate Experience" drew widespread attention when families who had been shown lavish AI-generated promotional images arrived at a sparsely decorated warehouse. Police were called and refunds were offered.[4] The pictures were not labelled as AI and did not resemble the event. Whatever label rules apply, the case is a reminder that **AI imagery used to sell something you cannot deliver is a misleading-advertising problem**, whether or not Article 50 is involved.

Voice cloning and synthetic audio

A voice that resembles a real, identifiable person and could pass as that person is a deepfake for these purposes.

- Do not clone a person's voice without their documented permission, covering the use you intend.
- Label synthetic voice-overs when they could be taken for a human recording, particularly in testimonials, news-style reads and phone messages.
- For a podcast or video with an AI voice, state it in the description and, where it matters, in the audio.

For an AI voice that talks live with the public, Chapter 3 applies instead.

Video and avatars

Talking-head avatars are common in training videos, explainer clips and social ads.

- If a presenter is a synthetic avatar of nobody, say "AI presenter" in the first seconds or on screen.

- If it is an avatar of a real, consenting person, say "AI-generated version of [name]" and keep their written consent.
- Never use a synthetic presenter to deliver a testimonial, claim or professional advice as if a real person gave it.

Provider marks: do not strip them

Most major image and video tools now embed machine-readable provenance information, such as **C2PA Content Credentials** or IPTC metadata that records the item as AI-generated.[5] Article 50(2) puts the marking duty on providers. As a deployer:

- Do not remove metadata to hide the origin of AI content.
- Prefer export settings that preserve credentials.
- Remember that social platforms often strip metadata on upload, which is why a **visible label** is your own duty, not the provider's.

Platforms have their own rules

The major social and video platforms have introduced their own toggles or labels for realistic AI-generated content, and they may add a label automatically or penalise creators who fail to use the tool.[6] Policies change often. Check the current help pages for each platform you post on. Platform compliance is separate from legal compliance: meeting one does not meet the other.

Checklist for images, video and audio

- ☐ We know which of our visuals are AI-generated or AI-edited.
- ☐ None depicts a real person, place or event in a way that never happened, without a label.
- ☐ Invented people are never presented as real customers, staff or experts.
- ☐ Voice clones have written permission and a label.
- ☐ Product images show the product as it truly is.
- ☐ We keep provider metadata and add a visible label where needed.
- ☐ We use each platform's AI label tool where it exists.

> **The rule.** A label is not decoration. It is the answer to "Did this really happen?"

CHAPTER 5

AI-Written Text: Articles, Product Pages, Newsletters and Translations

Text is where the rule is most often misunderstood. Many businesses assume that anything written with AI needs a label. It does not. And a few kinds of text need one even though no one thinks of them as dangerous.

What the rule actually covers

Article 50(4) requires deployers to disclose that text has been artificially generated or manipulated when it is **published with the purpose of informing the public on matters of public interest**.[1] The Commission's FAQ sets three cumulative tests: the text is published, it is informative to the public, and it concerns a matter of public interest. Listed topics include politics, public administration, justice and law enforcement, fundamental rights, public security, public health, the environment, consumer safety, and economic, scientific or cultural developments.[2]

The important exception: the labelling duty does not apply to text that has undergone **human review or editorial control**, where a person or organisation holds editorial responsibility for the publication. The Commission describes human review as deliberate, substantive examination by people with relevant expertise. A spell-check or formal check does not count.[2]

Does your AI-written text need a label?

Text that is almost certainly outside the label duty

- Product descriptions and catalogue copy.
- Advertising copy and slogans.
- Internal documents, emails and meeting notes.
- Personal or brand blog posts that are not about public-interest topics.
- Text a person has substantively reviewed, edited and approved as their own.
- Translations and grammar fixes.

Text that sits near the line

- A blog post "explaining" a new tax rule or health guideline.
- A newsletter with advice on consumer rights, safety or public policy.
- A local-news-style site that summarises council decisions.
- Social posts about elections, public health or emergencies.
- A business explaining a regulatory change to its customers.

These are the places where an unreviewed AI draft is both a labelling risk and an accuracy risk. The fix is the same for both: **a qualified human reads it, checks the facts and takes responsibility.**

An editor's desk: review is a person with a pencil, not a button

"Human review" has to be real

- A named person (or editorial team) has authority to change, approve or reject the text.
- They have the expertise to judge accuracy on that subject.
- They actually read the text and check key claims against sources.
- There is a record: who reviewed, when, what changed.

If you can meet all four, the label duty falls away, and you gain a better article. If you cannot, label it.

> **Say this.** "This article was drafted with AI and edited and fact-checked by [Name]." Or, if no human review took place: "This article was generated by AI and has not been reviewed by a person."

Beyond Article 50: when text misleads

Even where no label is required, AI text can create other problems.

- **Fake expertise.** A bio, byline or "ask our nutritionist" column that implies a human professional where there is none.
- **Invented facts and sources.** Hallucinated figures, quotations or studies. A widely discussed 2023 case involved a US media brand reported to have published articles under invented author profiles with AI-generated portraits, which it later removed.[7]
- **Misleading claims in marketing.** An AI-drafted claim ("clinically proven", "carbon neutral") must be substantiated by you, not by the tool.
- **Professional rules.** If you are in a regulated profession, your own rules may say more about disclosing AI use to clients than the AI Act does.

Translations, summaries and rewrites

A machine translation of text you wrote is not an "AI-generated text" disclosure issue under Article 50, although a note for important documents is good practice: "Machine-translated; the English original prevails." An AI summary of a third party's document presented to your customers should say what it is.

Checklist for text

- ☐ We know which published text is AI-generated or AI-edited.
- ☐ Public-interest content has a named human reviewer with expertise.
- ☐ Reviewed text records who, when and what was checked.
- ☐ Unreviewed public-interest text carries a clear label.
- ☐ No byline or bio implies a human who did not write or check it.
- ☐ Every factual claim in marketing copy is substantiated.
- ☐ Translations of important documents say they are machine-translated.

> **The rule.** A person who stands behind the text is the best disclosure there is. If nobody will, say so.

CHAPTER 6

Marketing, Reviews, Testimonials and AI Influencers

If Chapters 3 to 5 are about telling people what they are looking at, this chapter is about the one thing you must never fake: **the customer**.

The review rule, in one paragraph

EU consumer law lists, among the practices that are unfair in all circumstances, submitting or commissioning false consumer reviews, and presenting reviews as coming from real consumers when they do not.[11] The UK banned fake reviews as part of its consumer-protection reforms, in force since April 2025.[12] In the US, the Federal Trade Commission's 2024 rule on fake reviews and testimonials expressly covers AI-generated ones.[15] None of these needs Article 50. A review written by a tool and posted as a customer is a fake review, however it was made.

Real customer, real review?

The four-way test for anything that looks like a customer voice.

- **Written by a real customer, published honestly** — OK
- **AI summary of real reviews, labelled, with links to the originals** — OK
- **Customer used AI to polish their own words and approved them** — OK, with care
- **AI-written “customer” review or invented testimonial** — NEVER

You can invent a scene. **You cannot invent a customer.**

Real customer, real review, or not? The four-way test

Small tags, big promises

What you must not do

- Ask a tool to write "customer reviews" or "testimonials" and publish them as real.
- Create fictional customers, with names, photos and quotes, and present them as real.
- Pay for, or commission, reviews from people who did not use the product.
- Use AI to rewrite real reviews so they sound more positive or to hide negative ones.
- Gate reviews so only happy customers can post.
- Present an AI-generated "expert" quoting facts as if a professional said them.

What you may do

- **Summarise real reviews with AI,** as long as the summary is fair, says it is AI-generated and links to the full reviews.
- **Help customers write their own reviews** with prompts or spell-check, provided the content reflects their genuine experience and they approve it.
- **Use fictional personas for internal marketing planning,** such as audience profiles, as long as they never appear to customers as real people.
- **Use AI-generated illustrative content** in clearly labelled examples: "Example customer (fictional)".

> **Say this.** "This summary was generated by AI from 214 verified customer reviews. Read the individual reviews below."

Verified reviews and how you say so

If you display reviews, EU consumer law expects you to tell shoppers whether and how you check that they come from people who actually used the product.[11] If you say reviews are "verified", be able to show how. AI does not change this. It does make it easier to claim verification you have not done.

AI influencers and virtual brand ambassadors

A virtual influencer with a name, a backstory and a feed is a common marketing format. The honesty questions are the same as for a human endorser.

- **Is it clear it is an AI character?** If followers could think it is a real person, say so in the bio and on sponsored posts.
- **Is it clearly an ad?** Sponsored posts need the usual advertising label, whether the endorser is human or not.
- **Is the character saying things that are true?** A virtual ambassador cannot "love" a product, "lose ten kilos" or "feel healthier". Claims of personal experience by a character who has none are misleading.
- **Is it based on a real person?** Then consent, personality rights and the deepfake rules apply.

Ads and AI visuals

Use AI visuals in ads, but follow three rules:

1. **The product must be shown as it is.** AI-generated visuals may set a mood; they must not misrepresent what the buyer receives.
2. **Before-and-after, results and testimonial imagery must be real.** An AI-generated "result" photo is a false claim.
3. **Label realistic synthetic people in sensitive categories** such as health, finance and children's products, where the risk of misleading is higher and regulators pay closer attention.

Personalised and dynamic content

AI makes it easy to send different messages to different people: "dynamic" emails, product pages that change by visitor, chat-based sales funnels. Where this involves personal data, the GDPR's transparency duties apply. Where it exploits vulnerabilities or uses manipulative design, consumer-protection and AI Act prohibited-practice rules may be relevant.[18] The practical rule is not to use AI to target people's weaknesses or to disguise sales pressure as friendly advice.

Transparency is enforced

Italy's data-protection authority fined OpenAI EUR 15 million in December 2024, partly over transparency failures toward users and non-users.[16] A large provider, not a small business, but proof that transparency is a real requirement, not a courtesy.

Checklist for marketing

- ☐ No review, testimonial or customer story is AI-written and presented as real.
- ☐ Any AI-generated summary of reviews says so and links to the originals.
- ☐ We can show how "verified" reviews are verified.
- ☐ Virtual influencers are identified as AI characters and ads as ads.
- ☐ Characters make no first-hand experience claims they cannot have.
- ☐ Product visuals show the real product.
- ☐ Before-and-after and results images are real.

The rule. You can invent a scene. You cannot invent a customer.

CHAPTER 7

Agencies, Freelancers and Creators: Clients, Contracts and Credit

When you produce work for others, disclosure becomes a three-way matter: you, your client and the public. The AI Act's rules do not change because you were hired to do the work. But who carries them does, and that is a contract question.

Who is the deployer?

If an agency uses an AI tool to create an image that the client publishes, who must label it?

- The party that **publishes** the content to the public, in a professional capacity, is the deployer for the labelling duty. That is usually the client.
- The agency, as the **producer**, carries a practical duty: it must tell the client what it used, so the client can meet the duty.
- If the agency publishes on the client's behalf (managing their social accounts, for example), it is likely to be acting as, or for, the deployer, and should treat the duty as its own.

The Commission notes that a legal person remains the deployer even if contractors operate the system for it, and that employees acting under the organisation's control are not separate deployers.[2] The safest working assumption: **if your name or system presses publish, you share the duty.**

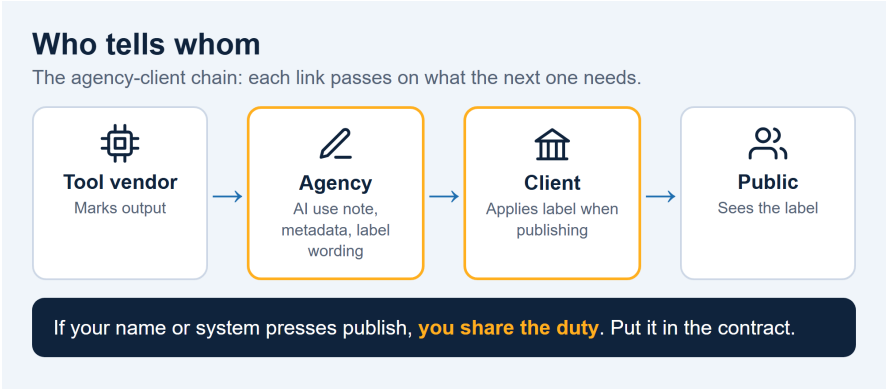

Who tells whom: the agency-client disclosure chain

What your clients need from you

1. **A statement of where AI was used.** Which deliverables were drafted, generated or edited with AI, and with which tools.
2. **A statement of what a human did.** Who reviewed, edited, checked facts and approved.
3. **The metadata and credentials,** left intact where possible.
4. **A suggested label.** Wording and placement that fit the channel.
5. **Rights and risk information.** The tool's terms on commercial use, any reported training-data issues and your own warranties.

Do this as a one-page "AI use note" delivered with each project. Appendix E gives you a clause checklist.

What your clients may ask of you

Clients increasingly want assurances. Reasonable requests include:

- A **no-AI** or **AI-with-approval** clause for sensitive work.
- A requirement to **disclose** tools and processes.
- A promise not to put **client confidential information** into public AI tools.
- **Indemnities** if AI output infringes someone's rights.
- A **human review** of anything customer-facing.

Agree each of these in writing. Do not promise "100% human-made" unless you can prove it.

Credit and honesty with clients

The question "Did you use AI?" should have a truthful answer. Quietly using AI for a fixed-fee job is not illegal. Misrepresenting what you did, or billing as if hours were worked that were not, is a contract and honesty issue.

- If you charge by the hour, bill the time actually spent.
- If you price by value, say that AI is part of your process.
- If a client expressly forbade AI, do not use it.
- If you subcontract to someone who uses AI, flow the terms down.

Creators: your audience is the client

For newsletter writers, YouTubers, podcasters and course creators, the audience is the person being told. A few habits help.

- **A short, standing statement** on your site: how you use AI in your work and where you do not.
- **Per-item labels** where AI plays a major role: an AI voice, a synthetic image, a drafted section.
- **Personal brand honesty.** If you are known for a voice, use AI in ways that do not erode it, and say so when it matters.
- **Sponsored content** still needs the advertising label.

> **Say this.** "I use AI to help draft and edit. Everything published under my name is reviewed and approved by me. Images marked 'AI' are generated."

Using other people's work in AI tools

Putting a client's or a third party's material into a public AI tool can breach confidentiality, copyright or data-protection duties. Before you do:

- Check the tool's terms for training on inputs and retention.
- Use business or enterprise plans with no-training commitments where you handle client data.
- Do not paste personal data into tools without a lawful basis and a data-processing agreement.

- Keep a list of which tools see which client data.

AI literacy applies to you

The AI Act's AI-literacy duty, in force since February 2025, expects providers and deployers to make sure their staff and others operating AI on their behalf have a sufficient level of AI literacy.[17] For a small agency, that means a short written guide, a briefing for anyone who uses the tools, and a record that it happened.

Checklist for agencies and creators

- ☐ We keep a list of AI tools used on each project.
- ☐ Each delivery includes an AI use note.
- ☐ Contracts say what we may use and what we may not.
- ☐ Metadata and credentials are kept intact.
- ☐ Staff are briefed on confidentiality and on labelling.
- ☐ We answer "Did you use AI?" truthfully every time.
- ☐ Our website has a short, honest AI-use statement.

> **The rule.** Whoever presses publish should know what the machine did. Whoever made it should tell them.

CHAPTER 8

People, Emotions and Recorded Meetings

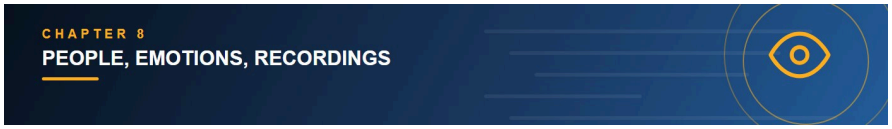

Some AI systems do not talk to people or produce content for them. They watch, listen and infer. These uses have a different disclosure logic, because the people affected are often not customers who chose to be involved.

Emotion recognition and biometric categorisation

Article 50(3) requires deployers of emotion-recognition or biometric-categorisation systems to inform the people exposed to them that the system is operating, and to process personal data in line with the GDPR.[1] The Commission's FAQ says deployers need not explain the purpose in the notice, and that the duty applies whether the system works in real time or after the fact.[2]

Before the notice, however, comes a more basic question: *should you be using this at all?*

- The AI Act **prohibits** emotion recognition in the **workplace and in education**, except for medical or safety reasons, and has done so since 2 February 2025.[14] Sentiment scoring of employees' faces or voices in a call centre falls near this line.
- Other biometric categorisation that infers sensitive characteristics, such as ethnicity, political opinions or sexual orientation, is also restricted.
- Biometric data used to identify a person is "special category" data under the GDPR and requires a stricter legal basis.

For a small business, the practical advice is blunt: **avoid facial or voice emotion analysis on customers and staff.** If a vendor offers "sentiment", "engagement" or "mood" detection, ask exactly what it measures, from

what input, and what you must tell people.

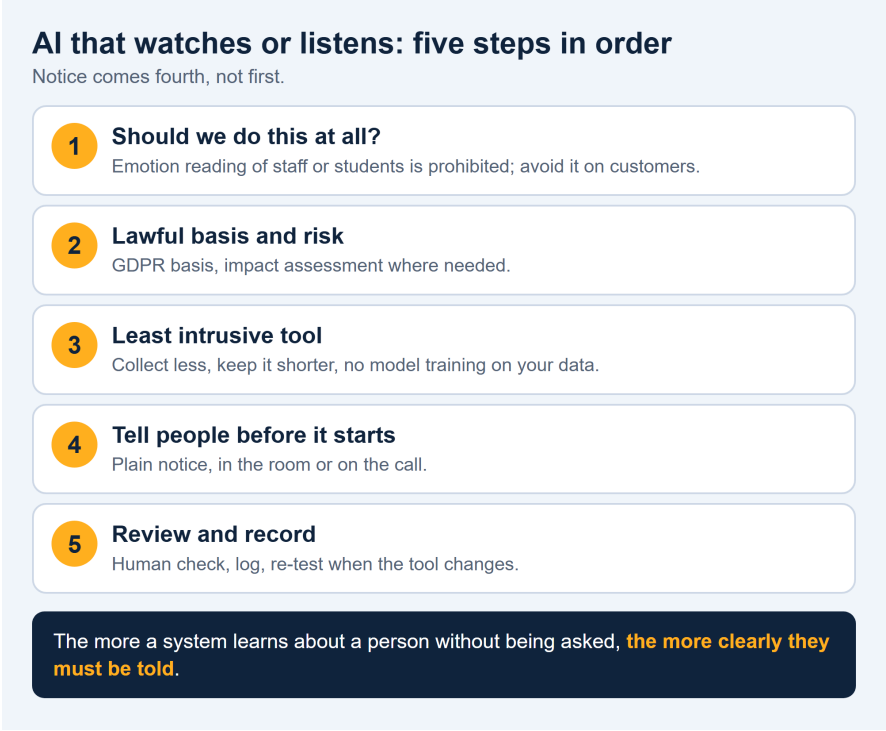

Notice is the last step, not the first

AI meeting recorders and note-takers

Transcription and summary bots are among the most widely adopted AI tools in small businesses, and the least thought through.

- **Tell participants before recording.** Joining a call as a silent bot is poor practice and may breach recording and consent laws that vary by country.
- **Use clear wording:** "This meeting is being recorded and transcribed by an AI note-taker. The summary will be shared with attendees."
- **Know where it goes.** Many tools store transcripts on third-party servers and may use them to improve their models, unless you turn that off.
- **Think about who is in the room.** A meeting with a client's staff, a job candidate or a patient needs a stricter approach.

- **Do not use AI summaries as the only record** of anything important without a human check.

> **Say this.** "Before we start: I use an AI tool to record and transcribe this call. It is shared only with me and you. Tell me if you would rather I did not."

Hiring and AI

AI tools that screen CVs, rank candidates or analyse video interviews are among the uses the AI Act treats as high-risk, with stronger duties on providers and deployers, which are phased in under the AI Omnibus.[8] The GDPR's rules on automated decisions also apply.[13] A disclosure that says "we use AI in recruitment" is a start, not a defence. Tell candidates what is used, what it does and how to ask for human review.

Personal data in prompts

When staff paste customer or employee information into a public AI tool, that is processing personal data. Your privacy notice should say that you use AI tools where it applies, with whom data is shared, and why. You need a lawful basis, and a data-processing agreement with the provider where it processes data for you. For a deeper treatment, see the companion titles on AI privacy and data leaks.

AI in customer profiling and pricing

Segmenting customers, scoring leads or setting personalised prices with AI involves profiling. The GDPR requires you to tell people about profiling and automated decision-making, including meaningful information about the logic involved and its consequences, where Article 22 decisions are made.[13] Dynamic pricing based on personal data needs particular care and a clear notice.

Children and vulnerable people

Children should not be your test audience for new AI features. If a chatbot or content feature may be used by minors, apply stricter design: plain language, no manipulative personas, no collection of unnecessary data,

parental controls where relevant and a rapid human escalation.

Checklist for people-facing AI

- ☐ We do not use emotion recognition on staff or on students.
- ☐ We have asked vendors in writing what "sentiment" tools actually analyse.
- ☐ Recorded meetings are announced before the recording starts.
- ☐ Notetaker settings keep data private and stop model training.
- ☐ Recruitment tools are documented, and candidates are told.
- ☐ Our privacy notice names AI tools that process personal data.
- ☐ Anything affecting children has a human escalation route.

> **The rule.** The more a system learns about a person without being asked, the more clearly they must be told.

CHAPTER 9

The Wording Library and the Labelling Workflow

You now know what needs disclosing. This chapter gives you the words and the routine to do it consistently.

Principles for good wording

1. **Plain.** "AI-generated", "AI assistant", "created with AI". Not "synthetically augmented" or "algorithmically enhanced".
2. **Specific.** Say what the AI did: "image generated by AI", "text drafted by AI and edited by our team".
3. **Early.** At first exposure: the first chat message, the top of the page, the first seconds of the video.
4. **Visible.** Readable on a phone, in the language of the audience, not hidden in a footer or a hover.
5. **Short.** If a label needs a paragraph, link to a longer page and keep the label itself brief.
6. **Consistent.** Use the same words everywhere, so customers learn what they mean.

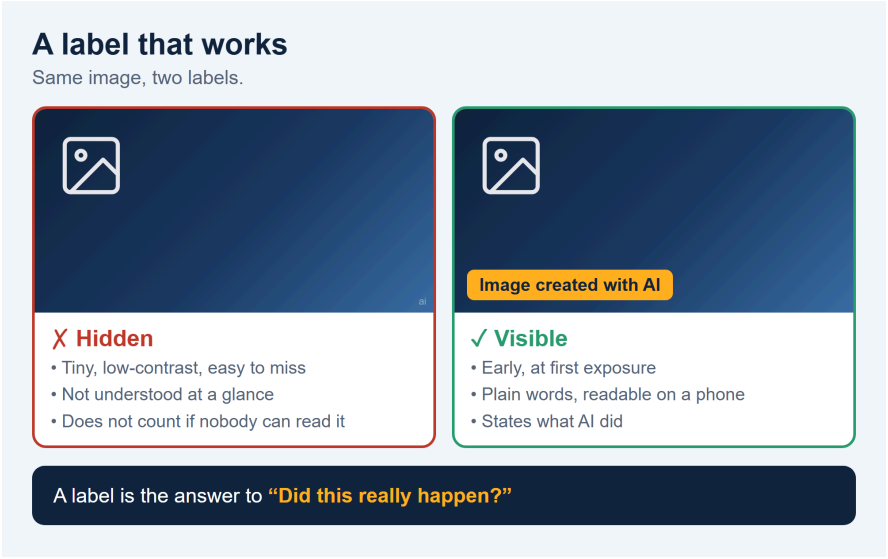

A label that works: early, visible, specific

The wording library

Copy, adapt, and keep the final versions in your policy.

Channel	Say this
Website chat	"Hi, I'm [Name], [Business]'s AI assistant. I can help with [topics]. Ask for a person any time."
Voice bot	"Hello, this is [Business]'s virtual assistant. I'm an AI. I can [task] or connect you to a person."
Email auto-reply	"This reply was generated by an AI assistant. A member of our team will follow up if needed."
AI image (caption)	"Image created with AI."
AI image with real product	"Illustration created with AI. Product shown as it is in the photos below."
Synthetic person (not real)	"AI-generated person. Not a real customer."
Deepfake of a real person (with consent)	"AI-generated video of [Name], used with their permission."
Synthetic voice	"Voice generated by AI."
AI avatar presenter	"Presented by an AI avatar."
AI-assisted article (reviewed)	"Drafted with AI; edited and fact-checked by [Name]."

Channel	Say this
AI-generated article (unreviewed)	"Generated by AI. Not reviewed by a person."
AI review summary	"Summary generated by AI from [number] verified reviews."
Machine translation	"Machine-translated. The [language] original prevails."
Meeting recording	"This call is recorded and summarised by an AI tool. Tell us if you would rather not."
Fictional example	"Fictional example created for illustration."

Where to put the label

- **Chat:** inside the first message, not only in the widget header.
- **Image:** a caption beside or under the image, or a visible mark on it. For short video, include the label in the first frames and in the description.
- **Audio:** a spoken line at the start, plus a note in the show notes.
- **Article:** under the headline or at the top, in the byline area.
- **Social post:** in the text of the post and, where one exists, the platform's own AI toggle.
- **Email:** in the body near the top.
- **PDF or ebook:** on the copyright page and at the point of use.

What not to do

- **Do not hide labels in light grey on white.** If a reasonable person cannot read it, it does not count.
- **Do not bury them in the terms and conditions.**
- **Do not over-label.** A label on every spell-checked sentence teaches customers to ignore labels.
- **Do not use ambiguous jargon.** "Enhanced by technology" is not a disclosure.
- **Do not claim "human-made" falsely.** It is as misleading as a missing "AI" label.

The checkpoint every piece of content passes through

The workflow: five steps for every piece of content

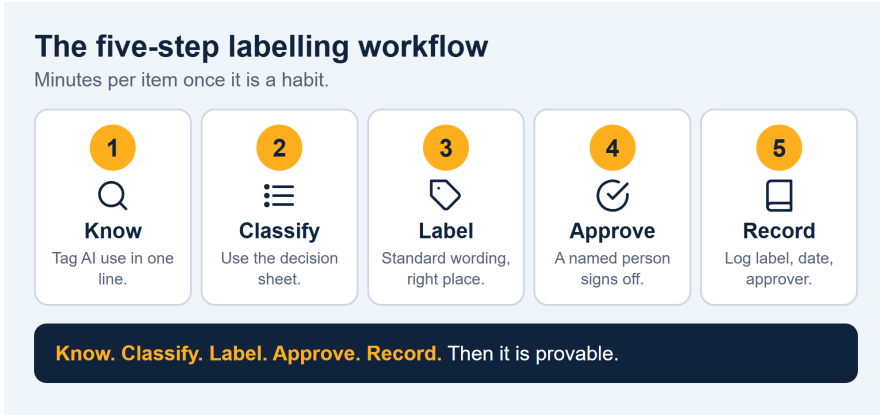

A five-step labelling workflow

1. **Know.** Record that AI was used and what for. A one-line tag in your content tool is enough: AI: *image*; *human: edited*.

2. **Classify.** Use the decision sheet (Appendix B): chatbot, deepfake, public-interest text, review, or other.
3. **Label.** Apply the standard wording and place it where the audience will see it.
4. **Approve.** A named person signs off any content that needs a label, or that depends on a human review exemption.
5. **Record.** Save the label, the date and the approver in the labelling log (Appendix D).

For a small team, this takes minutes per item and almost nothing once it is a habit.

Provenance and metadata

Where tools support it, keep **content credentials** (C2PA) or equivalent metadata through your editing pipeline. They carry the record of how an image or video was made. They are not a substitute for a visible label, but they back it up, and platforms and press increasingly read them.[5] If you must compress or convert files, use tools that preserve the data.

Training your team

The AI Act's AI-literacy duty is easy to satisfy in a small business with three things: a short written guide, a 30-minute briefing, and a record of who attended.[17] The briefing should cover: what counts as AI, which tools are approved, what to disclose and how, and who to ask. Repeat it when tools or rules change.

A note on "AI-free" claims

Some businesses will market themselves as "100% human-made". It can be a strong positioning. It must also be true. If you claim it, define it ("no AI-generated images or text in our published content"), set rules to uphold it, and be honest about what falls outside, such as spell-check and translation tools.

Checklist for the workflow

☐ A standard set of labels is approved and saved.

- ☐ Every channel has a defined place for the label.
- ☐ Each content item records its AI use in one line.
- ☐ A named person approves labelled content.
- ☐ Labels and approvals are logged.
- ☐ Metadata and credentials are preserved.
- ☐ The team has been briefed in the last 12 months.

The rule. Decide once, label the same way every time, and write it down.

CHAPTER 10

Your 30-Day Plan and Five Case Files

Four weeks to a disclosure habit

Four weeks to a disclosure habit

A few hours a week.

Week 1	Week 2	Week 3	Week 4
Find it	**Classify it**	**Fix it**	**Prove it**
☐ List every AI touchpoint	☐ Run the decision sheet	☐ Rewrite chat openings	☐ Adopt the one-page policy
☐ Ask each person what they use	☐ Name an owner for each	☐ Add labels; remove fake customers	☐ Start the labelling log
☐ Screenshot what customers see	☐ Mark: label, review, stop, none	☐ Keep content credentials	☐ Brief team; add client clause

By day 30: a register, owners, standard labels, a policy and a log.

A 30-day plan for the small business

Week 1: Find it. List every place AI touches a customer, in your register (Appendix A). Include chatbots, images, text, voice, reviews, recorders and the AI built into software you already use. Ask each team member what they use. Capture screenshots of what customers currently see.

Week 2: Classify it. Run each entry through the decision sheet (Appendix B). Mark each one: needs a disclosure, needs a human review, needs to be stopped, or needs nothing. Name an owner for each.

Week 3: Fix it. Update chat openings. Add labels to the images and text that need them. Remove or replace any fake customer content. Switch on content credentials. Rewrite any "human-like" claims in your bots.

Week 4: Prove it. Adopt the one-page policy (Appendix C). Start the labelling log. Brief the team. Add a clause to client contracts (Appendix E). Put a short statement on your website. Diary a monthly review.

Case file 1: Cobble & Thread (fictional)

The problem: a chatbot claimed to have tried on clothes; AI "customers" appeared in product captions.

The fix: Poppy's opening line now says she is an AI assistant; her prompt forbids personal experience claims; the model images carry "AI-generated visual; product shown in photos below"; the captions that named invented customers were removed. Mara logged the changes with dated screenshots.

Lesson: the worst problem was not the missing label but the false experience claim. Fix the behaviour, then the wording.

Case file 2: Harbour Lane Studio (fictional)

The problem: a four-person agency used AI to produce a client's campaign video with a synthetic spokesperson who looked like a real local doctor. The client published without a label.

The fix: the agency stopped the campaign, replaced the avatar with a clearly invented character and added "AI presenter" on screen. It issued an AI use note with every delivery and added contract clauses (Appendix E) making the client's publication duty explicit.

Lesson: if a synthetic person resembles a real one, it is a deepfake problem. Check who the face resembles before you publish.

Case file 3: Mira Voss, newsletter writer (fictional)

The problem: Mira's weekly newsletter on consumer rights was drafted by AI and sent with light editing. A reader pointed out a wrong refund-period claim.

The fix: she introduced a rule that every claim of law or fact is checked against a primary source before sending, recorded who checked it in her log, and added a line: "Drafted with AI; checked and edited by Mira."

Lesson: human review is both a legal exemption and a quality control. It only works if it is real.

Case file 4: Oakmere Physio (fictional)

The problem: a clinic's AI phone assistant greeted callers in a warm human voice and asked about symptoms. One caller asked if they were speaking to a nurse.

The fix: the greeting now begins with "This is Oakmere Physio's virtual assistant, an AI". The assistant no longer asks clinical questions. It books and passes calls to a person. A privacy notice update covers recorded calls.

Lesson: health context raises the stakes. Disclose first, collect less, hand over faster.

Case file 5: Fernwood Home (fictional)

The problem: an intern asked an AI tool for "ten five-star reviews" to fill a new product page. Six went live before a colleague noticed.

The fix: all six removed within the hour. The shop wrote a short policy against fake reviews, restricted review-posting rights, and introduced a rule that review text is only ever customer-written. A real-review request email now goes to buyers after delivery.

Lesson: a tool will do exactly what it is told. The control belongs in your process and policy.

What to do next

Start with the three exposures that cause the most harm: **a bot that pretends to be human, a synthetic person presented as real, and a fake review.** Fix those first. Then work through the rest in order. You will not reach perfect. You will reach defensible, which is what the rules and your customers expect.

Conclusion: Say It Plainly

The AI disclosure trap catches businesses that decide by accident: some say nothing, some say everything, and most have never written a word down. The way out is to decide on purpose.

Ask who is speaking, what is real, and who answers for it. Tell people at the moment it matters, in words they understand. Keep proof. Repeat the check when your tools change.

If you do that, you will meet most of what Article 50 asks, and you will have earned something no regulation can give you: customers who know exactly what they are dealing with, and keep coming back anyway.

APPENDIX A

AI Use and Disclosure Register

Copy one row per AI use. Review monthly.

#	Tool and supplier	What AI does	Who sees it	Disclosure needed?	Wording and place	Owner
1						
2						
3						
4						
5						
6						

APPENDIX B

Disclosure Decision Sheet

Run each AI use through these questions in order. Stop at the first "yes".

Question	If yes
1. Does a person chat or speak with an AI system directly?	Tell them at the start that it is an AI (Ch. 3)
2. Is it a realistic image, video or audio that resembles a real person, place or event and could pass as real?	Label as AI-generated (Ch. 4)
3. Is it text published to inform the public on a matter of public interest, with no real human review?	Label it, or add real review (Ch. 5)
4. Does it present a customer, reviewer or expert who is not real?	Remove it (Ch. 6)
5. Does it analyse emotions, faces or voices of people?	Stop, or get advice, notice and a lawful basis (Ch. 8)
6. Does it make or shape decisions about people?	Check GDPR Article 22 and tell them (Ch. 8)
7. None of the above	No label required; keep the record

Decision: _____ **Owner:** _____ **Date:** _____ **Next review:** _____

APPENDIX C

One-Page AI Disclosure Policy Template

Copy, fill in and adapt. Have it reviewed if you operate in a regulated field.

1. Purpose. This policy explains how [Business] tells people when AI is involved in what we say, show and publish.

2. Owner. [Name, role] is responsible for this policy and reviews it every [3 / 6 / 12] months.

3. Chatbots and voice. Every customer-facing AI assistant says in its first message that it is an AI. It does not claim to be human or to have human experience. It offers a route to a person.

4. Images, video, audio. We label realistic AI-generated content that could be mistaken for real people, places or events. We do not strip provenance metadata. We do not clone a real voice or likeness without written consent.

5. Text. We label AI-generated text on matters of public interest unless a named person with the right expertise has reviewed it and takes editorial responsibility. We check every factual claim before publishing.

6. Reviews and testimonials. We never publish AI-written or invented reviews, customers or experts as real.

7. Clients. We tell clients where we use AI and ask them to tell us where they do.

8. People. We do not use emotion recognition on staff or students. We announce AI recording of meetings before it starts.

9. Records. We keep a register of AI uses and a labelling log.

10. Training. Everyone who uses AI at work is briefed once a year.

Approved by: _____ Date: _____

APPENDIX D

Content Labelling Log

Record each labelled item. Keep for at least [12] months.

Date	Content and channel	What AI did	Label used	Placement	Approved by

APPENDIX E

Client AI-Use Clause Checklist

Use with your own lawyer or template. These are topics to cover, not finished clauses.

- ☐ **Definitions.** What "AI tool" and "AI-generated content" mean in this contract.
- ☐ **Permitted use.** Whether and for which deliverables the supplier may use AI.
- ☐ **Disclosure to client.** The supplier tells the client which tools and where.
- ☐ **Client's publication duty.** The client is responsible for labels on content it publishes, and the supplier supplies what is needed to meet them.
- ☐ **Human review.** Which deliverables a named human must review before delivery.
- ☐ **Confidential and personal data.** What may and may not go into AI tools.
- ☐ **Metadata and credentials.** Kept intact unless agreed otherwise.
- ☐ **Rights.** Ownership of outputs; warranties on training data and third-party rights; indemnity limits.
- ☐ **Likeness and voice.** Written consent for any real person's face or voice.
- ☐ **No fake reviews.** The supplier will not create or post fake reviews or testimonials.
- ☐ **Changes.** How either side proposes a change to AI use during the project.
- ☐ **Records.** How long each side keeps the AI use note and labelling log.

Glossary

AI-generated content. Text, images, audio or video produced or substantially altered by an AI system.

AI literacy. The skills and understanding that let people use AI systems competently and be aware of their risks. An AI Act duty for providers and deployers.

Biometric categorisation. Assigning people to categories, such as ethnicity or political opinion, based on their biometric data.

C2PA / Content Credentials. An open standard for attaching tamper-evident provenance information to media files.

Deepfake. AI-generated or manipulated image, audio or video that resembles existing persons, objects, places, entities or events and would falsely appear authentic.

Deployer. An organisation that uses an AI system under its own authority in a professional capacity.

Emotion recognition. Inferring a person's emotions or intentions from biometric data.

Human review. Deliberate, substantive examination of content by a person with relevant expertise who can approve, change or reject it.

Matter of public interest. A topic of general concern, such as politics, public health, justice, the environment or consumer safety.

Provider. An organisation that develops an AI system, or has one developed, and places it on the market under its own name.

Watermark. A machine-readable mark embedded in generated content so it can be detected as synthetic.

A Note on GDPRGard.eu

Disclosure sits beside a wider question: what does the GDPR expect of a small business? If this book has raised questions about privacy notices, records of processing, impact assessments or data-subject requests, **GDPRGard.eu** offers free tools, guides and articles written in plain English for small organisations.

About the Author

Zirlandia Milkovic writes practical, plain-English guides on data protection, AI and building a business with confidence. She believes privacy is not a barrier to ambition but a foundation of trust, and that business owners should not need a legal degree to understand the rules that apply to them.

She is the author of a series of business guides, including *GDPR Made Simple for Women Founders*, and her work focuses on turning complex European rules into simple steps, checklists and templates that busy people can actually use. She is also behind **GDPRGard**, a growing collection of free GDPR tools, guides and articles for small businesses, available at gdprgard.eu.

Visit **gdprgard.eu** for free tools, articles and updates, including notes on changes to the rules covered in this book. If this guide helped you, an honest review where you bought it helps other readers find it.

Notes and Sources

Research cut-off: 5 November 2026. Items marked (read) were consulted for this edition; the others rest on the author's working knowledge of well-known public records and should be re-checked against the primary source before you rely on them.

1. Regulation (EU) 2024/1689 (AI Act), Article 50 (transparency obligations) and Article 3 (definitions of provider, deployer and deepfake). eur-lex.europa.eu.
2. European Commission, FAQ on transparency obligations under Article 50 of the AI Act, last updated 24 July 2026 (read). digital-strategy.ec.europa.eu.
3. Moffatt v. Air Canada, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal, 14 February 2024).
4. Public reporting on the "Willy's Chocolate Experience" event in Glasgow, February 2024, including police attendance and refunds.
5. C2PA, Content Credentials technical specification (c2pa.org); IPTC, "digital source type" vocabulary for AI-generated media (iptc.org).
6. Help and policy pages of major social and video platforms on labelling realistic AI-generated content. These change frequently; check the current version for each platform you use.
7. Public reporting (Futurism, November 2023) on AI-generated author profiles at a US media brand, and the brand's subsequent removal of the content.
8. Regulation (EU) 2026/1744 amending the AI Act ("AI Omnibus"), published in the Official Journal on 24 July 2026 and in force from 27 July 2026; high-risk dates of 2 December 2027 and 2 August 2028; a transitional period to 2 December 2026 for the Article 50(2) marking duty for generative systems placed on the market before 2 August 2026 (read, through the Commission FAQ and law-firm commentary; check the Official Journal text).
9. AI Act, Article 99 (penalties): fines of up to EUR 15 million or 3% of worldwide annual turnover for breaches of Article 50, and the rule in Article 99(6) that for SMEs the lower amount applies (the FAQ states the EUR 15 million / 3% figures; paragraph 6 rests on working knowledge).
10. Code of Practice on transparency of AI-generated content (voluntary; assessed as adequate by the Commission and the AI Board) and the Commission's Guidelines on transparency obligations (read, through the Commission FAQ and a July 2026 law-firm note).

11. Directive 2005/29/EC (unfair commercial practices) and Directive (EU) 2019/2161 (the "Omnibus" consumer directive), which added fake and unverified reviews to the practices prohibited in all circumstances.
12. Digital Markets, Competition and Consumers Act 2024 (UK), Part 4, Chapter 1 (unfair commercial practices, including fake reviews), in force from 6 April 2025.
13. GDPR (Regulation (EU) 2016/679), Articles 5(1)(a), 12 to 14, 22 and 35.
14. AI Act, Article 5(1)(f) (prohibition of emotion recognition in the workplace and education institutions), applicable from 2 February 2025.
15. US Federal Trade Commission, Trade Regulation Rule on the Use of Consumer Reviews and Testimonials (16 CFR Part 465), final rule announced August 2024.
16. Garante per la protezione dei dati personali (Italy), decision of December 2024 imposing a EUR 15 million fine on OpenAI.
17. AI Act, Article 4 (AI literacy), applicable from 2 February 2025.
18. AI Act, Article 5(1)(a) and (b) (prohibited manipulative and exploitative AI practices).