AI COMPLIANCE · BUSINESS

THE AI LITERACY FILE

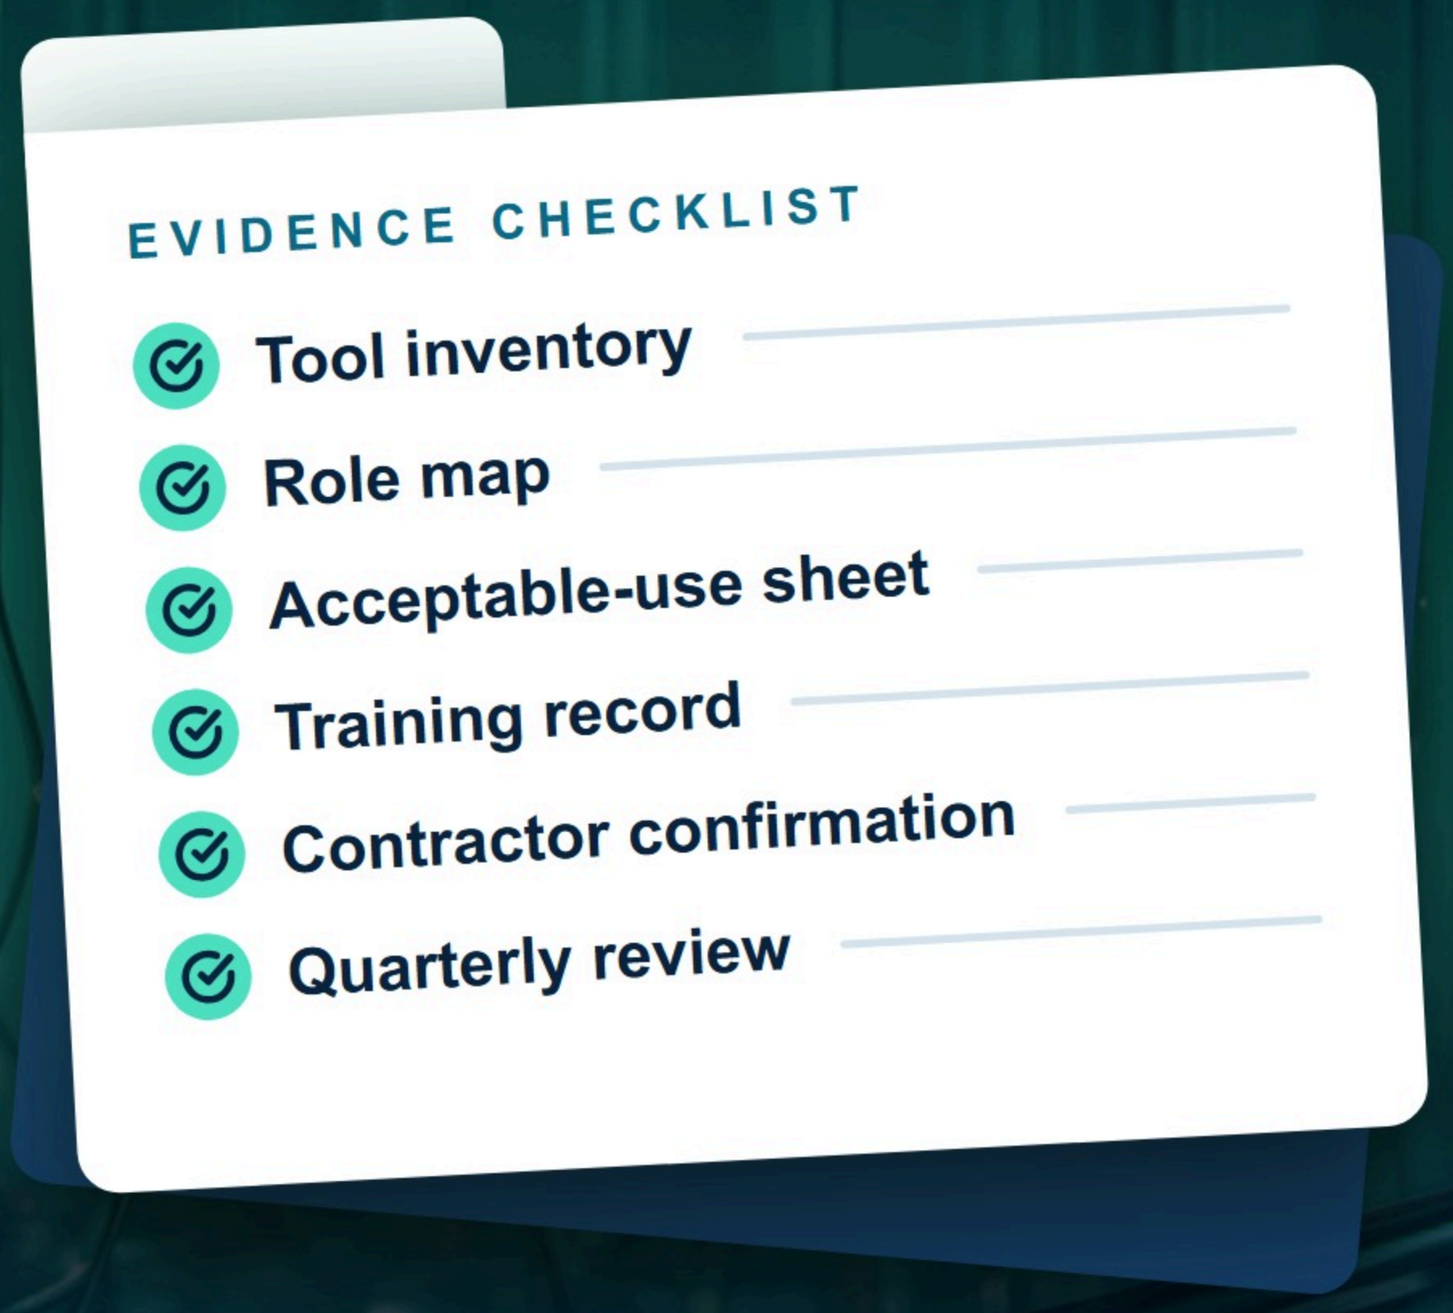

How Small Businesses Can Build **Practical AI Training** and **Evidence** Under the EU AI Act

TOOL INVENTORY · TRAINING SCENARIOS · EVIDENCE CHECKLIST · ARTICLE 4

ZIRLANDIA MILKOVIC

AI COMPLIANCE · BUSINESS

THE AI LITERACY FILE

How Small Businesses Can Build Practical AI Training and Evidence Under the EU AI Act

ZIRLANDIA MILKOVIC

The AI Literacy File

How Small Businesses Can Build Practical AI Training and Evidence Under the EU AI Act

© 2026 Zirlandia Milkovic. All rights reserved. No part of this book may be reproduced without written permission, except for brief quotations in reviews.

First edition, 2026. Research cut-off: 11 November 2026. British English. Companion resource: GDPRGard.eu.

This book is practical guidance, not legal advice. All examples, businesses and people in worked scenarios are fictional.

Contents

START HERE

Educational Disclaimer 4
Introduction: The Question You Cannot Yet Answer 4
Who This Book Is For 6
How to Use This Book 6
The Ten-Minute AI Literacy Check 7

CHAPTERS

1. What Article 4 Actually Says 9
2. What "AI Literacy" Means Inside a Small Business 14
3. Find Your AI: The Tool Inventory 17
4. Who Needs What: Roles, Staff and Contractors 21
5. Employee and Contractor Responsibilities 24
6. Designing the Programme: A 90-Day Plan 27
7. Training Scenarios, Part One: Habits Under Pressure 30
8. Training Scenarios, Part Two: New Risks 34
9. Delivering Training on a Small Budget 37
10. The Evidence File 40
11. Review, Update and Improve 43
12. When Someone Asks for Proof 46

CLOSING

Conclusion: A File That Tells the Truth 49

APPENDICES

Appendix A: AI Tool Inventory 50
Appendix B: Role Map 51
Appendix C: One-Page Acceptable-Use Sheet 52
Appendix D: Training Record 53
Appendix E: Contractor and Supplier Confirmation 54
Appendix F: Scenario Facilitator Cards 55
Appendix G: Quarterly Review Checklist 56

BACK MATTER

Glossary 57
A Note on GDPRGard.eu 58
About the Author 59
Notes and Sources 60

Educational Disclaimer

This book is practical, educational guidance for small-business owners, managers, HR professionals and the consultants who support them. It is not legal advice, it does not create a professional relationship between the author and the reader, and it cannot tell you whether your particular business is compliant.

The EU AI Act, its amendments, the Commission's guidance and national enforcement arrangements are all still developing. The description of Article 4 in this book reflects the Commission's published guidance and secondary legal commentary available at the research cut-off in the notes at the back. **The exact wording of the amended Article 4 should be read in the Official Journal text of Regulation (EU) 2026/1744 and Regulation (EU) 2024/1689 before you rely on it.** Whether a business is a "provider" or a "deployer", and whether any of its tools is a high-risk AI system, depends on facts this book cannot see.

Following the steps here will help you build a sensible, documented AI literacy programme. It cannot guarantee that a regulator, client, insurer or court will accept it.

Every business, person and incident in a worked example is **fictional** and labelled as such. Real cases, statistics and official positions are sourced in the notes at the back. Templates are starting points: adapt them to your business and have them reviewed where the stakes justify it.

Introduction: The Question You Cannot Yet Answer

Priya runs Orchard Lane Studio, a fourteen-person design and print business with five regular freelancers. We will use her as an example throughout. She and the business are fictional.

On a Tuesday in September, a new client's procurement team sends Priya a supplier questionnaire. It is two pages long, and most of it is routine. Question 17 is not:

"*Please describe the measures you have taken to support AI literacy among staff and other persons using AI systems on your behalf, and the records you keep of them.*"

Priya knows her team uses AI. The designers use image tools. The account manager pastes client emails into a chatbot to tidy the wording. A freelancer uses an AI note-taker on client calls. Nobody ever decided any of this. Nobody wrote it down. There is no policy, no list of tools, and no record of anyone being shown how to use them safely.

She can answer the question *honestly* in one sentence: "We have not done anything." What she cannot do is answer it *well*, because she does not know what "AI literacy" means, whether it is a legal duty for a business her size, or what she would need to show.

This book is for Priya, and for every owner, manager and adviser in her position. It does three things.

1. **It explains the duty accurately.** Article 4 of the EU AI Act is an obligation, but a much more modest one than many headlines and vendor brochures suggest. It does not require you to guarantee any particular level of literacy, and it does not require certificates.
2. **It gives you a programme you can run.** Not a course catalogue, but a short sequence of practical steps: find your AI, decide who needs what, set clear responsibilities, train with real scenarios, and keep proof.
3. **It gives you the paperwork.** An inventory, a role map, an acceptable-use sheet, training records, a contractor confirmation and a review checklist, ready to adapt.

The aim is not a perfect file. It is a **credible, proportionate, up-to-date file**: one that a client, an insurer, an auditor or a regulator would recognise as the work of a business that took the duty seriously.

> **The rule.** AI literacy is not a certificate you buy. It is a habit you can show: you know which AI you use, you have told people how to use it well, and you can prove it.

Who This Book Is For

This book is written for people who have to make AI use in a small organisation safer, and who need something they can finish rather than something they must study:

- **Owners and directors** of businesses with roughly 5 to 100 people who use AI tools, whether or not they have chosen to.
- **Managers** who have been handed "the AI question" with no brief and no budget.
- **HR and people-operations professionals** who run induction, training and policy refreshes.
- **Compliance consultants, DPOs and accountants** who advise several small clients and want one repeatable programme.
- **Freelancers and agencies** who work for larger clients and are now being asked for evidence.

You do not need a technical background. You do not need to be a lawyer. If you can run a team meeting and keep a spreadsheet, you can run this programme.

A word on tone. Nothing here assumes that your staff are careless or that you have failed. AI tools arrived in workplaces through browsers, phones and software updates, faster than any policy could follow. The goal is to catch up in an orderly way.

How to Use This Book

Read it in an evening, then work through it over about three months.

- **In a hurry?** Read Chapter 1 (what the law says), Chapter 3 (the inventory) and Chapter 10 (the evidence file). Then fill in Appendices A and D.
- **Starting from nothing?** Follow Chapter 6, the 90-day plan. It sequences the other chapters.
- **Training your team?** Go to Chapters 7 and 8 for the scenarios, and Appendix F for facilitator cards.
- **Working with freelancers or agencies?** Read Chapter 5 and use Appendix E.
- **Advising clients?** Chapters 2, 4 and 11 give you the structure; the appendices are the deliverables.

Each chapter ends with a short action list. The appendices are designed to be printed, photocopied or pasted into a spreadsheet. A good first output is a single folder, which we call **the AI Literacy File**, holding everything in this book you have completed.

The Ten-Minute AI Literacy Check

Answer yes or no. Be honest; nobody is marking this.

1. Could you list, within ten minutes, every AI tool or AI feature that people in your business use for work?
2. Do you know which of those tools are used on personal accounts rather than business accounts?
3. Does anyone, anywhere in the business, have written permission to use AI tools for work?
4. Have you told staff what information must never be put into an AI tool?
5. Does every person who uses AI at work know that its output can be confidently wrong?
6. Does someone check AI output before it goes to a customer, a regulator or the public?
7. Do your freelancers, agencies and contractors know your AI rules?
8. Could you show a client a dated record of what training or guidance your team received?

9. Does any AI tool make or influence decisions about people (hiring, pricing, credit, discipline)?
0. Does a named person own AI literacy, with a date set for the next review?

Scoring. Count your "yes" answers.

Score	What it means
0 to 3	You are at the start. Follow the 90-day plan in Chapter 6 in order.
4 to 7	You have habits but not proof. Focus on the inventory (Chapter 3) and the evidence file (Chapter 10).
8 to 10	You are well placed. Use Chapters 10 and 11 to tighten the records and set a review rhythm.

A "no" to question 9 is worth checking twice. Many businesses discover, when they look, that a recruitment platform, an advertising tool or a customer-scoring feature uses AI in a way they had not noticed.

CHAPTER 1

What Article 4 Actually Says

CHAPTER 1
WHAT ARTICLE 4 ACTUALLY SAYS

The EU Artificial Intelligence Act, Regulation (EU) 2024/1689, entered into force on 1 August 2024. Its early provisions, including the general provisions in Chapter I and the list of prohibited practices, began to apply on **2 February 2025**.[1] Article 4 sits in that first group. Its subject is **AI literacy**.

The AI literacy plan for a small business starts with understanding the duty

The duty in plain English

Article 4 applies to **providers and deployers of AI systems**. A provider develops an AI system, or has one developed, and places it on the market under its own name. A **deployer** is a person or organisation that *uses* an AI system under its own authority in a professional setting.[2] A café that uses an AI scheduling tool, an accountancy firm that drafts with a chatbot and a design studio that generates images are all, in the language of the Act, deployers.

The duty is to take measures to support the AI literacy of two groups:

- **your staff**, and
- **other persons dealing with the operation and use of AI systems on your behalf**, which the Commission explains includes contractors, service providers and clients working under your remit.[3]

The Act defines AI literacy in Article 3 as the skills, knowledge and understanding that allow people to make an *informed* use of AI systems and to be aware of the opportunities, risks and possible harm AI can cause.[2]

What changed in 2026

The original wording of Article 4 asked providers and deployers to take measures to ensure, "to their best extent", a *sufficient level* of AI literacy. Many businesses read "sufficient" as a target they could fail.

The **Digital Omnibus on AI**, Regulation (EU) 2026/1744, amended the provision. It was published in the Official Journal on 24 July 2026 and entered into force on 27 July 2026, according to law-firm commentary.[4] The Commission's FAQ describes the result like this:[3]

- AI literacy **remains an obligation** for providers and deployers.
- **No specific level**, and no "sufficient" level, is mandated.
- The measures are to **support the development** of AI literacy.
- The Commission and Member States continue to support businesses with tools and examples, but they do not take the duty over.

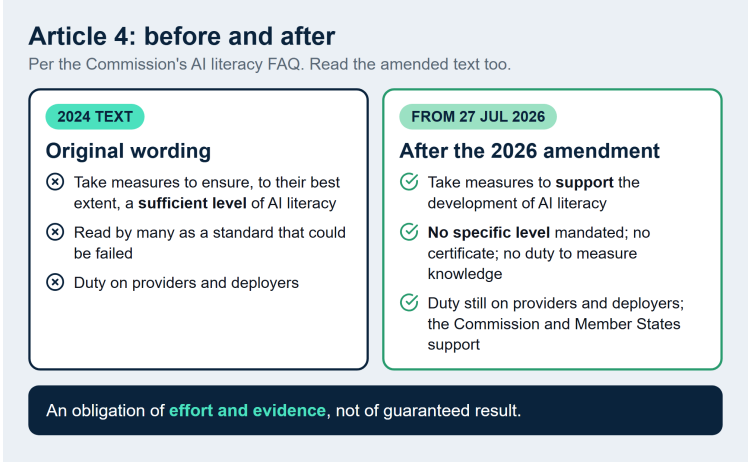

Article 4 before and after the 2026 amendment

So the important distinction is this. **Article 4 is still an obligation, but it is an obligation to take measures and support, not to guarantee an outcome.** You are not required to make every employee reach a

particular standard, and you are not required to measure what they know.

> **The rule.** The question a regulator or client will ask is not "is everyone literate?" but "what did you do, and can you show it?"

What the Commission says you do *not* need

The Commission's FAQ, last updated in October 2026, is unusually direct:[3]

- **No certificate** is needed. Organisations can keep an *internal record* of trainings and other guiding initiatives.
- There is **no obligation to measure** employees' knowledge of AI.
- There is **no mandated format.** Training, guidance, coaching and informal learning can all count.
- **No specific governance structure** is required, so you do not need to appoint an "AI officer".
- Simply handing staff the vendor's instructions for use is described as **ineffective** on its own.

What it does expect

The same FAQ points to factors you should weigh when you decide what to do:[3]

- your **role** (provider or deployer),
- the **risk** of the AI systems you use,
- your staff's existing **knowledge and experience**, and
- the **context** in which the tools are used, including who is affected.

For example, a team that uses a chatbot for drafting should at least be informed about its specific risks, such as made-up facts. Deployers of **high-risk** systems may need more, and Article 26 of the Act separately requires human oversight to be assigned to people with the necessary competence, training and authority.[2][3]

Timing and enforcement

Article 4 has applied since 2 February 2025. The Act's supervision and enforcement rules apply from **2 August 2026**. Enforcement is by **national market surveillance authorities**, not by the EU AI Office, and penalties are set by national law, so the practical position differs between Member States.[3] The FAQ also notes that the AI Act itself creates no criminal offences and no right to compensation, and that any damages claim would run through national law.

Key dates for Article 4

Why small businesses should still care

The FAQ does not exempt small businesses. Even if no inspector ever calls, the evidence question will arrive from other directions:

- **Client questionnaires** and tender documents, like Priya's Question 17.
- **Insurers**, who are beginning to ask how AI is governed.

- **Larger customers** passing their own compliance duties down the supply chain.
- **An incident**, where a leaked client file or a false statement sent to a customer makes "what had you told your people?" the first question.

None of this requires a large programme. It requires a proportionate one, which is what the rest of the book builds.

Your actions

- ☐ Read the Commission's AI literacy FAQ and the current text of Article 4 in the Official Journal.
- ☐ Write one sentence stating whether your business is a deployer of AI systems (it almost certainly is).
- ☐ Appoint one person to own AI literacy, with a first review date.

CHAPTER 2

What "AI Literacy" Means Inside a Small Business

The legal definition is deliberately broad. A small business needs something it can teach in an afternoon. This chapter turns the definition into **five habits** that you can train, observe and record.

Literacy is built through habits, not heavy study

The five habits

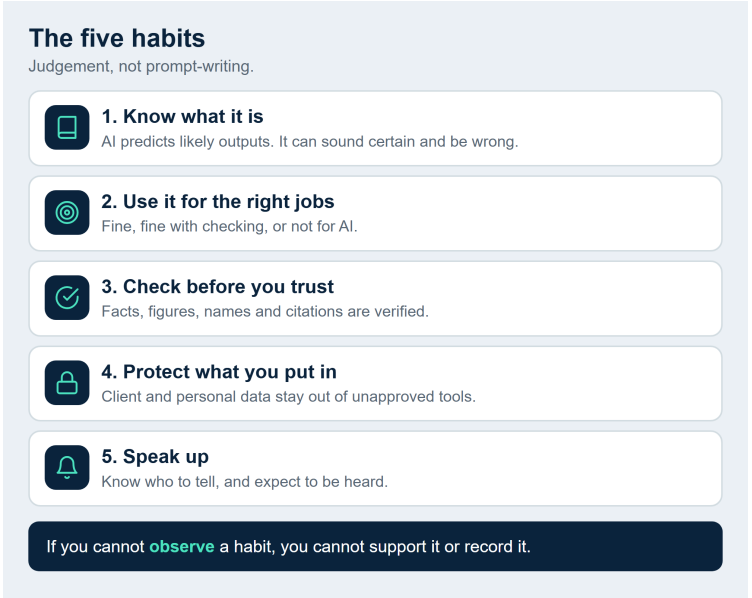

The five habits of an AI-literate team

1. Know what it is. AI tools predict likely outputs from patterns in data. They do not look things up the way a person does, and they do not "know" when they are wrong. A person with this habit can explain, in one sentence, why a chatbot can sound certain and be false.

2. Use it for the right jobs. Drafting, summarising, brainstorming and first-pass analysis are generally lower-risk. Decisions about people, legal or medical judgements and anything that cannot be checked are not. A person with this habit sorts tasks into *fine*, *fine with checking* and *not for AI*.

3. Check before you trust. AI output is a draft. Facts, figures, quotations, names and citations must be verified against a source. A person with this habit never forwards AI output they have not read and tested.

4. Protect what you put in. Client data, personal data, passwords, contracts and unpublished work do not go into tools the business has not approved. A person with this habit knows the difference between a business account and a personal one.

5. Speak up. When an output looks wrong, a tool behaves oddly, or someone suggests using AI in a way that affects customers or colleagues, the person knows who to tell and expects to be listened to.

Literacy is not the same as skill

These habits are about *judgement*, not about writing clever prompts. A confident power user who pastes a client contract into a free chatbot is less literate, for the purposes of Article 4, than a cautious colleague who stops to ask.

Matching depth to the person

Not everyone needs the same depth. The Commission's own approach is context-based, and Chapter 4 builds a role map. For now, notice three levels:

Level	Who	What they need
Awareness	Everyone, including people who never open an AI tool	The five habits in outline, the data rule, who to tell
Working knowledge	Regular users of AI tools	The habits in practice, tool-specific limits, checking routines
Decision-level	Owners, managers, HR, anyone choosing or approving AI	Risk questions, vendor checks, effects on people, record-keeping

Make it observable

If you cannot see a habit, you cannot support it. For each habit, decide what "good" looks like in your business. For example, for habit 3: "*Any client-facing document that used AI has a named person who checked its facts.*" That sentence is something you can put in a policy, teach, and later show a client.

Your actions

- ☐ Rewrite each of the five habits in one sentence that suits your business.
- ☐ Choose which level (awareness, working, decision) each person needs. Chapter 4 helps.

CHAPTER 3

Find Your AI: The Tool Inventory

You cannot support literacy about tools you have not found. The inventory is the foundation of everything else, and it is the document most likely to be requested first.

The tools are already on the desk: the inventory starts by looking

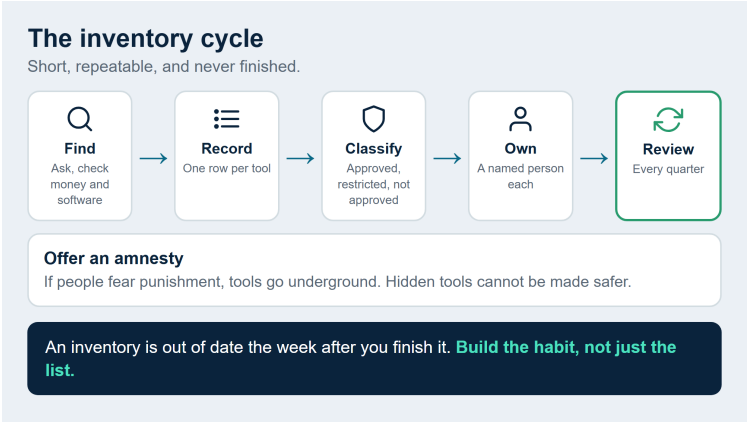

The inventory cycle

Where AI hides

Most businesses undercount. Besides the tools people name when asked ("ChatGPT"), look for:

- **Features inside software you already pay for**: email, office suites, design tools, CRM, accounting and recruitment platforms.
- **Browser extensions** and writing assistants.
- **Note-takers and transcription** in video-call software.
- **Phone features**: dictation, image editing, summaries.
- **Tools bought by one person** on a personal card or free account.
- **Tools used by contractors** on your work, on their own accounts.

The last group is called **shadow AI**: use that happens outside anything the business approved.

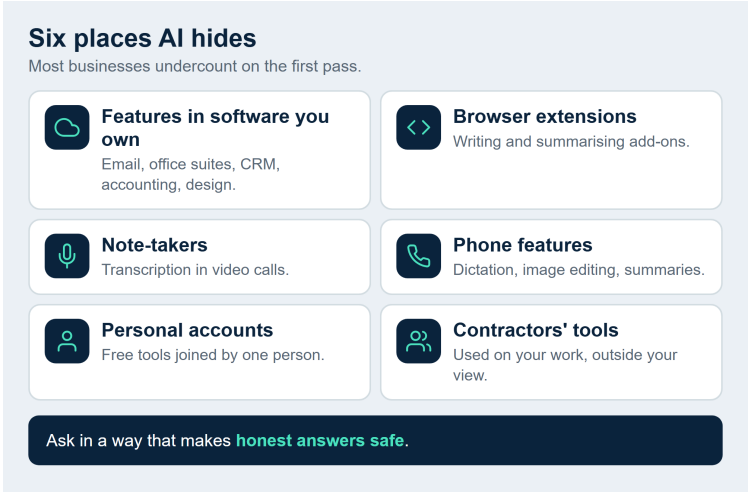

Six places AI use hides in a small business

How to run the discovery in one week

1. **Ask, without blame.** Send a short message: "*We are making a list of every AI tool and AI feature people use for work, so we can support you properly. There is no penalty for anything you list.*"
2. **Check the money.** Scan card statements, expense claims and subscription invoices for AI products.
3. **Check the software.** Look through the admin pages of your main business software for AI features and settings.
4. **Check the browsers.** Ask people to list extensions they use for writing or summarising.
5. **Ask contractors** (Chapter 5).

Offer an amnesty. If people fear punishment they will hide tools, and hidden tools cannot be improved.

What to record

Use Appendix A. For each tool, capture at least:

Field	Why it matters
Tool and supplier	You cannot assess what you cannot name
Who uses it, for what	Links literacy needs to real tasks
Account type (business, personal, free)	Personal and free accounts carry the highest data risk
Data it touches	Personal data, client confidential, public only
Affects people?	Hiring, pricing, discipline or customer decisions raise the stakes
Approved, restricted or banned	Gives staff a clear answer
Owner and review date	Someone is responsible and the record stays current

Classify, simply

Sort each tool into three groups:

- **Approved:** reviewed, business account, clear rules.
- **Restricted:** allowed only for stated tasks or with certain data excluded.
- **Not approved:** do not use for work until reviewed.

If a tool might be a **high-risk** system under the Act, for example one used in recruitment, worker management or access to essential services, flag it for a closer look with your adviser. Most office drafting and design tools are not high-risk, but you should not assume it.

Keep it alive

An inventory is out of date the week after you finish it. Add a standing question to onboarding ("which tools do you use?") and review the list every quarter (Chapter 11).

Your actions

- ☐ Send the amnesty message and collect replies for one week.
- ☐ Complete Appendix A, including contractors' tools.
- ☐ Mark every tool Approved, Restricted or Not approved.

CHAPTER 4

Who Needs What: Roles, Staff and Contractors

Article 4 speaks of staff and **other persons dealing with the operation and use of AI systems on your behalf**. For a small business, the second group is often as important as the first.

Map people to levels

Start with a list of everyone who touches AI for your business, then place each person into a tier. This is the role map. A template is in Appendix B.

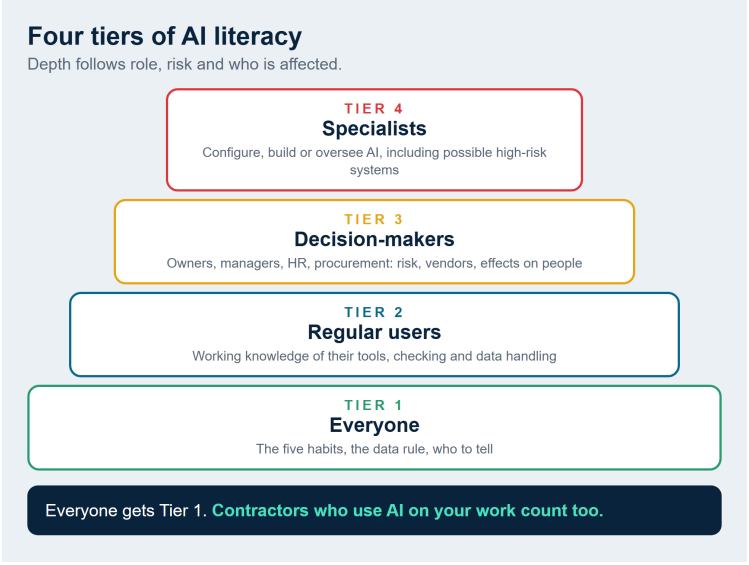

Four tiers of AI literacy by role

Tier 1: Everyone. Every employee, volunteer and long-term contractor receives the basics: what AI is, the data rule, the checking rule, how to raise a concern. Even people who say they never use AI need it, because AI features appear in tools without being announced.

Tier 2: Regular users. People who use AI tools weekly or more need working knowledge: the strengths and limits of their specific tools, how to verify output in their kind of work, and the handling of client and personal data.

Tier 3: Decision-makers. Owners, managers, HR and procurement need the additional questions: when a tool affects a person, what to ask a vendor, how to record a decision, when to escalate.

Tier 4: Specialists. If someone configures, builds or integrates AI, or operates a system that could be high-risk, they need deeper and more specific training, and the rules on human oversight may apply directly. [2]

Include people who are not employees

Under the Commission's explanation, "other persons" include contractors, service providers and clients acting under your remit, and the people working for a contractor need the same AI skills as employees.[3] Practically:

- **Freelancers** using AI on your work should receive the Tier 1 basics and confirm your data rule.
- **Agencies and suppliers** processing your data should confirm their own literacy measures (Appendix E).
- **Clients** you invite into your tools are a special case. Consider a short welcome notice.

You cannot train other businesses' staff, but you can set expectations, ask for confirmation and keep the reply.

Think about who is affected

The Act's phrasing asks you to consider "the persons or groups of persons on whom the AI systems are to be used".[2] A tool that writes your newsletter affects readers lightly. A tool that screens CVs affects

applicants seriously. The more a tool affects people, the more depth the people running it need.

A worked example

At Orchard Lane Studio, Priya maps the people:

Person	Uses AI for	Tier
Priya (owner)	Drafting proposals, approving tools	3
Tomas (HR/office manager)	Job adverts, screening a recruitment platform	3, plus a closer look at the platform
Three designers	Image generation, mock-ups	2
Account manager	Client email drafting	2
Eight other staff	Little or none	1
Five freelancers	Note-taking, translation, imagery	1, plus contractor confirmation

It takes twenty minutes and already reveals Priya's first priorities: Tomas's recruitment platform and the freelancers' note-taker.

Your actions

- ☐ Complete the role map in Appendix B for every person and contractor.
- ☐ Mark anyone who touches decisions about people.
- ☐ List every contractor or agency that uses AI on your work.

CHAPTER 5

Employee and Contractor Responsibilities

Training tells people what good looks like. Responsibilities tell them what is expected. Without the second, the first is a suggestion.

Responsibilities are agreed together, not imposed

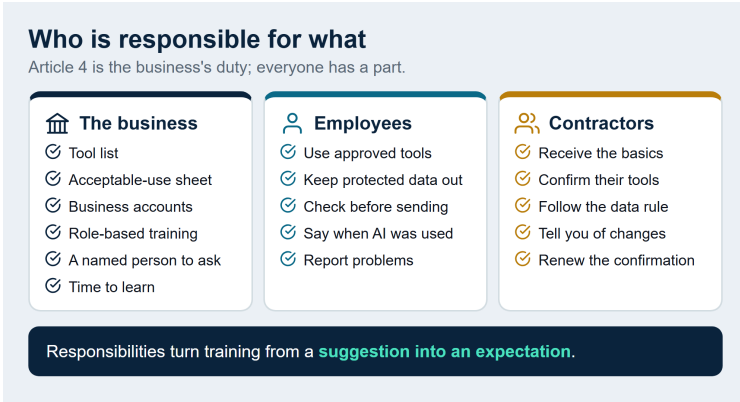

Who is responsible for what

What the business owes its people

Article 4 is the business's duty, so start with what *you* provide:

- A **clear list** of approved, restricted and not-approved tools (Chapter 3).
- A short **acceptable-use sheet** (Appendix C) written in plain language.
- **Business accounts** for approved tools, so people are not forced onto personal ones.
- **Training and guidance** matched to role (Chapters 7 to 9).
- A **named person** to ask, and a promise that asking will not be penalised.
- **Time** to learn. A policy that demands literacy but allows no time for it is hard to defend.

What employees owe the business

Keep the list short enough to remember. Five commitments cover most of the ground:

1. **Use only approved tools** for work, and ask before trying a new one.
2. **Keep protected information out** of tools that are not approved for it.
3. **Check output** before it is sent, published or relied on, and be able to say who checked.
4. **Say when AI was used** where a client, colleague or rule needs to know.
5. **Report problems**: wrong output, leaked data, odd behaviour, or pressure to use AI in a way that feels wrong.

These sit comfortably beside existing employment terms, confidentiality duties and data-protection obligations. If your contracts or handbook need updating, take advice on the wording: an acceptable-use sheet that is not part of the employment framework may be harder to rely on.

What contractors owe the business

Contractors are not employees, so the route is contractual and practical rather than disciplinary.

Step	What to do
Tell them	Send the Tier 1 basics and the acceptable-use sheet before work starts
Ask	Use the confirmation in Appendix E: which AI tools do they use on your work, on what accounts, with what data?
Agree	Record that they will follow your data rule and will not put your or your clients' confidential information into unapproved tools
Keep	File their reply with their contract
Revisit	Ask again at renewal, or when their tools change

For agencies and suppliers that handle personal data on your behalf, your existing data-processing agreement is the right place to add AI expectations, alongside the usual security and sub-processor terms.[5]

A note on clients

If clients give you confidential material, they may have their own AI rules and may forbid certain tools. Read the contract and any security schedule before you put client material into any AI tool. A client's rule that "no AI may be used on our data" overrides your internal approval.

Make it easy to comply

People break rules that are inconvenient. If your approved tool is slow, locked down or unavailable, staff will drift to the one that works. Where possible, give people something at least as good as the tools they are tempted to use, then enforce the line.

Your actions

- ☐ Draft the one-page acceptable-use sheet (Appendix C).
- ☐ Send the sheet and the Appendix E confirmation to every contractor.
- ☐ Ask your adviser whether your contracts or handbook should reference it.

CHAPTER 6

Designing the Programme: A 90-Day Plan

A programme is a sequence, not a pile. This chapter turns the earlier chapters into a plan that a busy owner can finish in about three months, using roughly two hours a week.

A twelve-week plan, visible on one wall

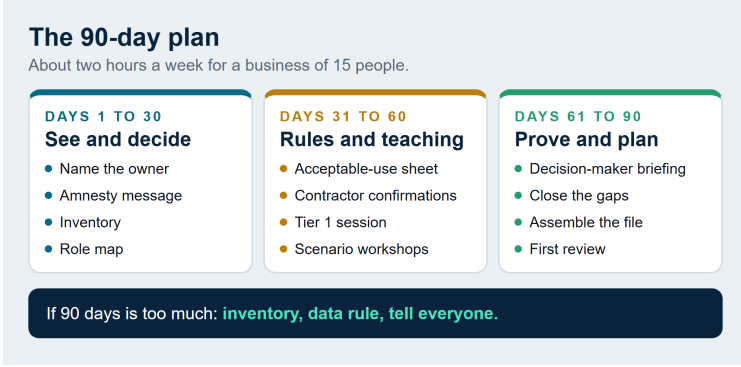

The 90-day plan

Days 1 to 30: See and decide

Week	Task	Output
1	Name the owner. Send the amnesty message.	Owner named, tool replies
2	Check subscriptions, software settings and browsers	Draft inventory
3	Complete the inventory and classify tools	Appendix A
4	Build the role map and decide the tiers	Appendix B

Milestone: you can list your AI tools and say who needs what.

Days 31 to 60: Set the rules and start teaching

Week	Task	Output
5	Draft and agree the acceptable-use sheet	Appendix C
6	Send contractor confirmations	Appendix E
7	Run the Tier 1 session for everyone (30 minutes)	Attendance record
8	Run Tier 2 scenario sessions (Chapters 7 and 8)	Scenario notes

Milestone: everyone has been told the rules, and regular users have practised on real scenarios.

Days 61 to 90: Prove it and plan the next quarter

Week	Task	Output
9	Tier 3 session for owners, managers, HR	Decision notes
10	Close gaps: missed attendees, late contractors	Updated records
11	Assemble the evidence file (Chapter 10)	The AI Literacy File
12	First review meeting; set the next date	Review record

Milestone: you can answer Priya's Question 17 with documents.

Budget and time

Costs are mostly time. A rough guide for a business of 15 people:

Item	Time
Owner: discovery, inventory, policy, evidence file	12 to 15 hours
Staff: Tier 1 session (30 minutes)	8 hours in total
Regular users: two scenario sessions (45 minutes each)	10 to 12 hours in total
Decision-makers: one session	3 hours in total

These are estimates, not benchmarks. Free resources and your suppliers' own guidance can replace paid courses, though the Commission notes that relying solely on vendor instructions is not enough.[3]

Start smaller if you must

If 90 days is too much, do these three things in the first fortnight: **write the inventory, write the data rule, and tell everyone.** That alone is worth more than a polished plan on a shelf.

Your actions

- ☐ Put the twelve weeks in the calendar with owners.
- ☐ Book the Tier 1 session now, before the work is done.
- ☐ Decide what "done" looks like for week 12.

CHAPTER 7

Training Scenarios, Part One: Habits Under Pressure

People learn habits by practising decisions, not by listening to definitions. The next two chapters give eight scenarios. Each takes ten to fifteen minutes to discuss and links to one or more of the five habits. Run two or three per session. Appendix F has facilitator cards.

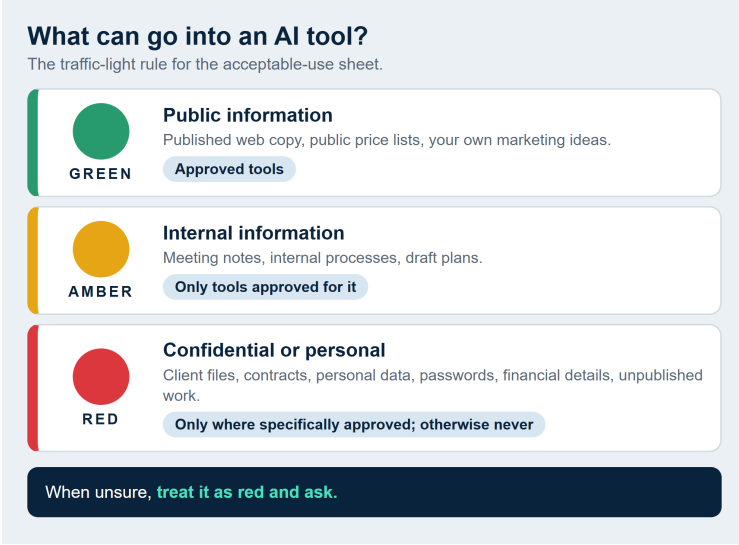

What can go into an AI tool?

How to run a scenario

1. Read the situation aloud.
2. Ask the questions. Let people answer before you do.
3. Read the "good practice" notes and compare.
4. Ask: *what would have to be true in our business for this to go wrong?*

5. Write down one change you will make, and record the session (Chapter 10).

All scenarios are fictional unless a source is named.

Scenario 1: The confident wrong answer

Habit: check before you trust.

Situation. An account manager at a fictional firm asks a chatbot for case law to support a point in a client letter. The reply cites three decisions with names, courts and years. She pastes them in. A colleague later cannot find any of them.

Discuss. - What should she have done before pasting? - Who is responsible for the letter: the tool or the sender? - What kinds of statement in your own work most need a source?

Good practice. AI tools can produce plausible but false citations, figures and quotations. In the widely reported 2023 US case *Mata v. Avianca*, lawyers filed a brief containing made-up cases generated by a chatbot, and the court sanctioned them.[6] Treat any fact, name, number or citation as unverified until you have opened the primary source yourself.

Scenario 2: The client contract in the free chatbot

Habit: protect what you put in.

Situation. A freelancer is short of time and pastes a client's draft contract, including names and payment terms, into a free chatbot account to "simplify the language".

Discuss. - Which parts of that text are confidential or personal data? - What do we not know about where it went and who can see it? - What would a safe alternative look like for this task?

Good practice. Free and personal accounts often have different terms from business ones, and you may not be able to say how inputs are stored or used. In a widely reported 2023 incident, engineers at Samsung pasted confidential source code into a public chatbot and the company restricted its use afterwards.[7] Use the traffic-light rule:

public information may go in; internal information only into approved business accounts; client-confidential and personal data only where the tool is specifically approved for it.

Scenario 3: The chatbot on your own website

Habit: use it for the right jobs; speak up.

Situation. A business adds an AI chat assistant to its site to answer questions about refunds. A customer is told they can claim a refund after travel, which is not the policy. They rely on it and complain.

Discuss. - Who is responsible for what the chatbot says? - How would we find out it was wrong? - Who checks the assistant, and how often?

Good practice. In the 2024 Canadian tribunal decision *Moffatt v. Air Canada*, the airline was held responsible for incorrect information given by its website chatbot, and the tribunal rejected the argument that the chatbot was a separate entity.[8] A customer-facing tool needs a named owner, tested answers, an easy route to a human, and regular review. If it talks to people, Article 50 transparency duties may also apply, so tell users they are dealing with AI.[9]

Scenario 4: The shortlist nobody chose

Habit: use it for the right jobs; speak up.

Situation. An HR manager uses a recruitment platform's "smart ranking" to sort 200 applications. It returns a top ten. Nobody can say how the ranking works. One strong candidate, who took a career break, is not in the list.

Discuss. - Which of the five habits does this test? - What do we need to ask the supplier? - At what point does a person look at the applications the tool ranked low?

Good practice. Tools used in recruitment can fall into the Act's high-risk category, which brings extra obligations, including human oversight by people with the competence, training and authority to do the job.[2] Even before that analysis, the practice is clear: AI may assist, but a trained person decides, records why, and can explain it. Ask the

supplier what the system does, what data it uses, and what it provides for oversight. GDPR rules on automated decision-making may also apply.[5]

Your actions

- ☐ Choose two scenarios and run them with your regular users.
- ☐ Record who attended, the date and one change agreed (Appendix D).
- ☐ Replace one fictional detail with a situation from your own business.

CHAPTER 8

Training Scenarios, Part Two: New Risks

These four scenarios cover newer risks that many small businesses have not yet met. They suit Tier 2 and Tier 3 sessions.

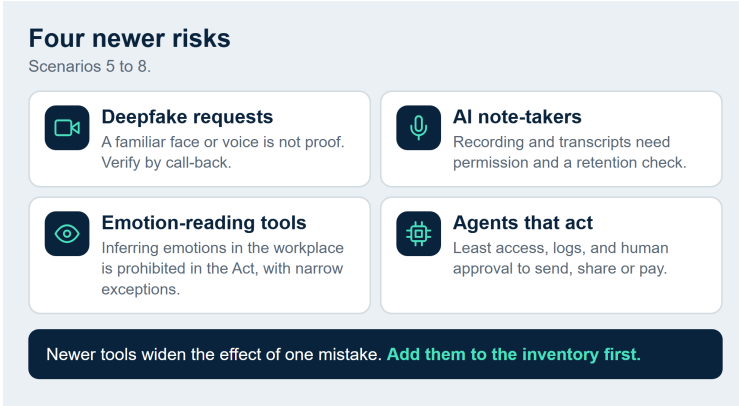

Four newer risks for small teams

Scenario 5: The call that sounded like the boss

Habit: know what it is; speak up.

Situation. A finance assistant receives a video call from someone who looks and sounds like the managing director, asking for an urgent, confidential transfer to a new supplier. The request breaks normal procedure.

Discuss. - What makes this convincing? - What control stops this regardless of how real the call looks? - Who can the assistant check with, and how quickly?

Good practice. In 2024, engineering firm Arup confirmed that an employee in Hong Kong had been deceived by a video call featuring deepfake versions of senior colleagues, resulting in transfers reported at about HK$200 million.[10] Voices and faces can no longer be treated as proof. The defence is process: call-back to a known number, two-person approval for new payees, and a culture where querying an "urgent" request is expected.

Scenario 6: The note-taker in the meeting

Habit: protect what you put in; speak up.

Situation. A contractor joins a client call and an AI note-taker joins with them, recording, transcribing and storing the conversation. The client was not told.

Discuss. - Who needs to know a call is being recorded or transcribed? - Where does the transcript go and who can read it? - When is it appropriate to turn the tool off?

Good practice. Recording and transcription involve personal data, and people usually need to be told, and in some cases to agree. Treat meeting tools as part of your inventory, ask permission at the start of the call, and agree with clients in advance when AI tools are acceptable. Check the tool's retention and sharing settings.[5]

Scenario 7: The mood-reading software

Habit: use it for the right jobs; know what it is.

Situation. A vendor offers a tool that analyses staff faces on video calls to measure "engagement" and "stress", and suggests it for performance reviews.

Discuss. - Would we be comfortable if this were used on us? - What might be wrong with a system that infers emotions from faces? - Who in the business should be told before any trial starts?

Good practice. The AI Act prohibits AI systems that infer emotions in the workplace, other than for medical or safety reasons. That prohibition has applied since 2 February 2025.[11] Part of literacy is

recognising when a proposal falls outside what is allowed. Decision-makers should be able to say "stop, take advice" before a pilot begins.

Scenario 8: The agent that acts on its own

Habit: protect what you put in; check before you trust; speak up.

Situation. A manager connects an AI assistant to the company inbox and calendar so it can "handle scheduling". It begins accepting invitations, replying to suppliers, and sharing a document link with an outside address.

Discuss. - What did we allow the tool to do, and what did we not realise we had allowed? - What is the smallest set of permissions that would do the job? - How would we notice something had gone wrong, and who reviews its actions?

Good practice. Tools that act rather than merely answer widen the effect of a mistake. Give such tools the least access they need, keep a log of what they do, require human approval for sending, sharing or paying, and add them to the inventory with a named owner.

Your actions

- ☐ Run Scenarios 5 and 8 with your finance and admin staff.
- ☐ Add "call-back and two-person approval" to your payment procedure.
- ☐ Ask every vendor about any emotion-recognition or monitoring features.

CHAPTER 9

Delivering Training on a Small Budget

The Commission does not mandate a format, and a small business does not need a learning platform. It needs short, repeatable sessions that people actually attend.

Short, practical sessions beat long lectures

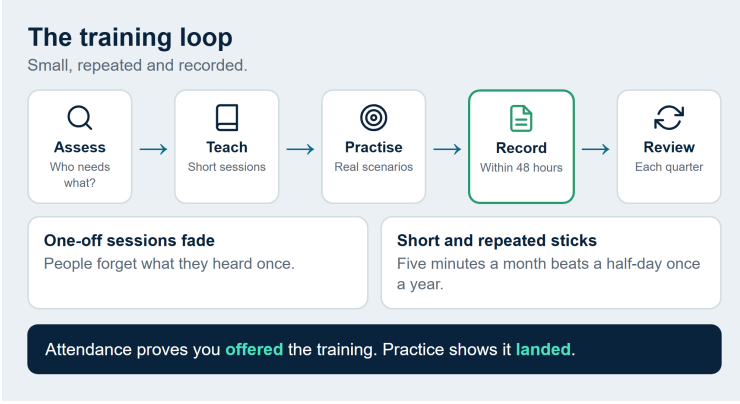

The training loop

Six formats that work

Format	Time	Best for
Kick-off session	30 minutes	Tier 1: the five habits, data rule, who to ask
Scenario workshop	45 minutes	Tier 2: two or three scenarios on real tasks
Tool walk-through	20 minutes	Tier 2: one approved tool, its settings and limits
Decision briefing	60 minutes	Tier 3: vendor questions, effects on people, records
Onboarding module	20 minutes	New starters and new contractors
Monthly "AI minute"	5 minutes	Everyone: one example, one lesson in a team meeting

Mix them. A one-off session is forgotten; short repeated ones stick.

A 30-minute kick-off agenda

Minutes	Content
0 to 5	Why we are doing this: the duty in two sentences, and what has not changed in how we work
5 to 12	The five habits with one example each
12 to 18	The data rule: the traffic-light sheet
18 to 25	One scenario discussed together
25 to 30	Who to ask, how to raise a concern, what happens next

Close with a promise you can keep: no one will be penalised for reporting honestly.

Teach with your own work

Generic examples teach generic lessons. Wherever possible, replace the fictional detail with tasks from your own business: "*Last month, someone used a tool to draft a customer reply. What would checking have looked like?*" People remember what resembles their own week.

Use free and supplier material, sensibly

Good sources of material include the Commission's own AI literacy pages and examples, national authorities, your software suppliers' guidance and reputable free courses. The Commission's FAQ points to practical examples under Article 4, EU Digital Innovation Hubs and the AI Skills Academy as support for businesses.[3] Using outside material is fine, but **adapt it** to your tools and rules and **record** what you used. Vendor instructions alone are not enough.

Make it inclusive

- Offer sessions at more than one time, and a short recording or written version.
- Avoid jargon. If a term needs explaining, explain it once, in the glossary.
- Include part-time staff and contractors, not only the people already in the office.
- Check understanding with questions, not just attendance.

Practise rather than test

You are not required to measure knowledge, and a score on a quiz proves less than people think. A better signal is a short practical task: *"Take this AI draft and show me what you would check."* If you do use quizzes, keep them low-stakes, and treat the results as information about the training, not about the person.

Your actions

- ☐ Choose one format per tier and book the dates.
- ☐ Prepare a 30-minute kick-off from the agenda above.
- ☐ Add a five-minute "AI minute" to one recurring team meeting.

CHAPTER 10

The Evidence File

The Commission says you can keep an internal record of your training and guidance.[3] This chapter turns that into a file you can hand over in a few minutes.

The AI Literacy File: everything in one folder

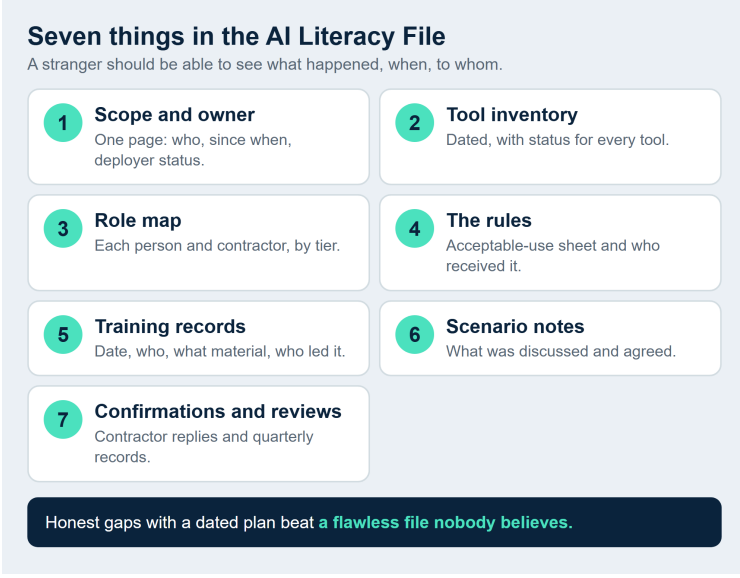

Seven things in the AI Literacy File

What the file contains

1. Scope and owner. One page: who owns AI literacy, the date started, and what the business is (deployer, and whether any tool may be high-risk).

2. The tool inventory (Appendix A), dated, with approved, restricted and not-approved status.

3. The role map (Appendix B), showing each person's tier.

4. The rules. The acceptable-use sheet (Appendix C) and the version history, with the date each person received it.

5. Training records (Appendix D): date, format, who attended, content or slides used, and who delivered it.

6. Scenario notes. For workshops: which scenarios, what was agreed, and what changed afterwards.

7. Contractor confirmations (Appendix E) and review records (Appendix G).

What good records look like

A record is useful if a stranger could read it and see **what happened, when, to whom and with what content.** For each session, capture:

- date and duration,
- format and facilitator,
- the people invited and the people who attended,
- the material used, saved in the file,
- any follow-up actions.

Keep a copy of the slides or notes exactly as presented. "We did AI training in March" is weak; "Training session on 14 March, 11 attendees, slides attached, three actions" is strong.

Keep it proportionate and private

The file should hold evidence of the *programme*, not a file on every employee's thoughts. Do not record test scores or personal judgements unless you need them, and keep training records under your normal retention policy and access rules. They are personal data, so record only what you need, store it securely and delete it when it is no longer useful.[5]

Show the gaps honestly

Records that show a plan with dated steps are better than records that pretend perfection. If a freelancer has not yet replied, write that down with the follow-up date. A candid, current file is more credible than a flawless one that nobody believes.

Answering Question 17

With the file complete, Priya's reply to the client can be short and specific:

> *"We maintain an AI Literacy File, owned by our managing director. It includes an inventory of the AI tools we use, a role map, an acceptable-use policy issued to all staff and contractors, dated training records and a quarterly review. A summary is attached; supporting documents are available on request."*

That sentence is only honest if every part of it is true. That is the point of building the file first.

Your actions

- ☐ Create the AI Literacy File folder with the seven sections.
- ☐ Set a rule: every session produces a record within 48 hours.
- ☐ Decide where the file lives, who can read it and how long you keep it.

CHAPTER 11

Review, Update and Improve

AI tools change monthly. A programme that was current in January may have gaps by summer. This chapter sets a rhythm that keeps the file true without eating your time.

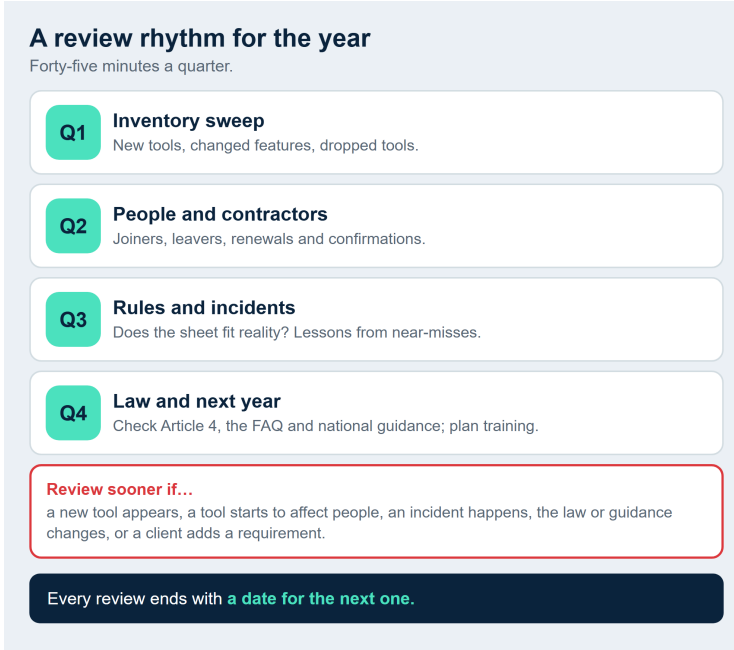

A review rhythm for the year

The quarterly review

Once a quarter, the owner of AI literacy meets whoever manages the teams for 45 minutes. Use Appendix G. The agenda has five items:

1. **Inventory.** What tools have appeared, changed or been dropped?

2. **People.** Who has joined, left or changed role? Does anyone need a different tier?
3. **Incidents and near-misses.** What went wrong or nearly did? What does it teach?
4. **Rules.** Does the acceptable-use sheet still fit what people actually do?
5. **Training.** What is the next session, and what is missing?

Record the date, attendees, decisions and the next review date.

Triggers for an unscheduled review

Do not wait for the calendar when something significant changes:

- **A new AI tool or feature** is adopted, or a supplier switches on AI in existing software.
- **A tool begins to affect people**: hiring, pricing, performance or customer decisions.
- **An incident**: leaked data, a false statement sent to a client, a deepfake attempt.
- **A change in the law or guidance**, including new Commission or national authority material.
- **A new client requirement** about AI use.
- **A new contractor or agency** with access to your work.

Keep up with the law without drowning

Amendments, guidance and national rules are moving. A light routine is enough: subscribe to the Commission's AI Act pages, set a quarterly reminder to check the text and FAQ, and ask your adviser to flag material changes. The 2026 changes to Article 4 are a reminder that wording can shift: other parts of the Omnibus, such as the timing of the high-risk rules, were also moved, so check dates before you rely on them.[4][12]

Learning from incidents

Treat a mistake as training material. A short, blameless write-up (what happened, which habit would have helped, what changes) turns a bad day into a scenario for next quarter. People who see mistakes handled calmly report them sooner.

Measure what matters

You are not required to measure knowledge, but you can track simple signs that the programme is working:

- the share of people covered by their tier's training,
- the number of tools on the inventory with a named owner,
- the number of concerns raised (a rising number often means more trust, not more problems),
- the time between a new tool appearing and it being reviewed.

Your actions

- ☐ Put four quarterly reviews in the calendar for the next twelve months.
- ☐ Write the unscheduled-review triggers into the acceptable-use sheet.
- ☐ Agree who watches for legal and guidance updates.

CHAPTER 12

When Someone Asks for Proof

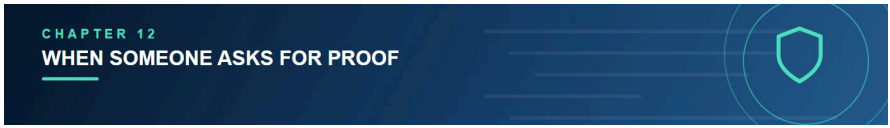

Sooner or later a client, insurer, auditor or authority will ask what you have done. This chapter prepares you for each, and for the awkward case where the file has gaps.

Be ready to hand the file over

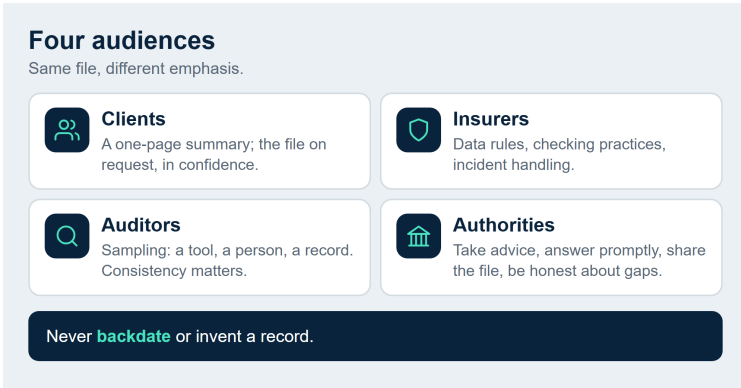

Four audiences, four kinds of answer

Clients and procurement teams

They want a short, specific answer and, increasingly, documents. Prepare a **one-page summary** of the programme: owner, tools approved, who is covered, last review date. Offer the supporting file on

request, under confidentiality, and do not send more than you are asked for.

Insurers and brokers

They are interested in risk controls: data rules, checking practices, incident handling. Use the acceptable-use sheet and the incident log. Be accurate about what you do, because statements to insurers may later matter.

Auditors and consultants

They will sample. They may pick a tool from the inventory and ask who uses it, which rule applies and what training covered it. They may pick a person and ask what they were told. Consistency across the documents matters more than volume.

Regulators

Enforcement of Article 4 sits with national market surveillance authorities from 2 August 2026, and the details differ by country.[3] If an authority contacts you, take advice, answer promptly and accurately, and provide the file. A credible programme with honest gaps and a dated plan is a far better position than no programme at all.

If the file is thin

Do not backdate. Do not invent records. If training took place before you had records, describe it honestly as "informal, undocumented" and show what you did from now on. Backdated documents are a risk in themselves and can turn a modest gap into a serious problem.

A short script for awkward questions

Question	A sound answer
"Is everyone AI-literate?"	"We have measures to support it, matched to roles. The law does not require a particular level, and we do not claim one. Here is what we do."

Question	A sound answer
"Do you have certificates?"	"No, none is required. We keep dated training records and the materials used."
"Which tools do you use?"	"Here is our inventory, last reviewed on [date]."
"What if a freelancer uses AI on our work?"	"They receive our rules, confirm their tools in writing, and we review at renewal."
"What happened when something went wrong?"	"We follow the incident steps, record it and update training. Here is the last example."

Your actions

- ☐ Write the one-page summary and keep it current.
- ☐ Keep the file ready to share in under an hour.
- ☐ Decide who answers client and authority questions.

Conclusion: A File That Tells the Truth

When Priya returned to Question 17, she did not write about a perfect programme. She wrote about a real one, ten weeks old, with a dated inventory, a role map, an acceptable-use sheet, three training sessions, five freelancer confirmations, and a review booked for the autumn. Two items were open and marked as such.

The client's procurement officer replied in one line: *"Thank you. That is more than most."*

That is the realistic prize. Not a badge or a certificate, but a business that knows what AI it uses, has told its people how to use it well, and can prove it.

Article 4 asks you to take measures to support AI literacy. It does not ask for perfection, and since the 2026 amendment it does not even ask for a particular level. It does ask for effort that is real, proportionate and visible. The five habits, the inventory, the tiers, the scenarios and the evidence file turn that effort into something you can finish.

Three things to do this week:

1. **Name the owner** and send the amnesty message.
2. **Write the data rule** on one page and tell everyone.
3. **Open the folder.** Create the AI Literacy File and put today's date on the first page.

The rest follows one step at a time.

APPENDIX A

AI Tool Inventory

Build this as a spreadsheet. One row per tool or AI feature, including those used by contractors and on personal accounts.

Ref	Tool / supplier	Used for	Who uses it	Account type	Data it touches	Affects people?	Status	Owner	Reviewed
A1									
A2									
A3									
A4									
A5									
A6									
A7									
A8									

Account type: Business, Personal, Free, Unknown. **Data it touches:** Public only, Internal, Client confidential, Personal data. **Affects people?** No, Customers, Staff, Applicants, Other. **Status:** Approved, Restricted (state the limits), Not approved.

Rule: a tool with no owner or no review date is not approved for use with client or personal data.

APPENDIX B

Role Map

Name	Role	Employee / contractor	Uses AI for	Tier (1 to 4)	Training due	Done

Tiers. 1 Everyone (awareness). 2 Regular users (working knowledge). 3 Decision-makers (owners, managers, HR, procurement). 4 Specialists (configure, build or oversee AI, including possible high-risk systems).

Check: has every contractor who uses AI on your work been placed on the map?

APPENDIX C

One-Page Acceptable-Use Sheet

Adapt this to your business and have it reviewed. Give it to every employee and contractor and record the date.

Our AI rules

We want you to use AI well. These rules keep our clients, colleagues and business safe.

1. **Use approved tools.** The current list is kept at: _____. If a tool is not on the list, ask before you use it for work.
2. **Use business accounts.** Do not use personal or free accounts for work unless the list says you may.
3. **Protect information.** - Green: public information may go into approved tools. - Amber: internal information only into tools approved for it. - Red: client-confidential material, personal data, passwords, financial details and unpublished work never go into a tool that is not approved for that data.
4. **Check before you send.** AI can be confidently wrong. Check facts, figures, names and citations against a source. You are responsible for what you send.
5. **Say so when it matters.** Tell clients and colleagues when AI contributed, where they need to know or our contract requires it.
6. **People decisions need people.** AI may support, but never decides, hiring, pay, discipline or customer treatment.
7. **Be careful with calls and payments.** Voices and faces can be faked. Verify unusual requests by calling back on a known number.
8. **Speak up.** Tell _____ if an output is wrong, data may have leaked, or you are asked to use AI in a way that feels wrong. You will not be penalised for reporting honestly.

Questions to: _____ **Version:** _____ **Date:** _____

APPENDIX D

Training Record

Complete one for every session and file it within 48 hours.

Session title: _____ **Date:** _____ **Duration:** _____

Format: Kick-off / Scenario workshop / Tool walk-through / Decision briefing / Onboarding / AI minute

Facilitator: _____ **Tier(s) covered:** _____

Materials used (attach copies):

Scenarios or topics covered:

Name	Role	Present	Notes

Actions agreed (owner and date):

Follow-up for absentees:

APPENDIX E

Contractor and Supplier Confirmation

Send with the acceptable-use sheet. File the reply with the contract.

Contractor / supplier: _____ **Contact:** _____ **Date:** _____

Please answer for the work you do for us.

1. Do you or your staff use AI tools or AI features in this work? **Yes / No**
2. If yes, which tools, and on what kind of account (business, personal, free)?

---

3. Will our information, our clients' information or any personal data be entered into those tools? **Yes / No / It depends (explain)**
4. What steps do you take to support AI literacy among the people working on our account (training, guidance, rules)?

---

5. Will you follow our acceptable-use sheet for work done for us? **Yes / Yes, except (explain)**
6. Will you tell us before you change the AI tools used on our work? **Yes / No**
7. How do you handle a suspected incident involving AI (leak, wrong output, impersonation)?

Confirmed by: _____ **Position:** _____ **Date:** _____

For us: reviewed by _____ **Follow-up needed:** _____ **Next check:** _____

APPENDIX F

Scenario Facilitator Cards

Copy each card for the person leading the session. Total time: about 12 minutes per scenario.

Scenario	Habit	Key question	Look for
1. The confident wrong answer	Check before you trust	Who verified the citations?	People opening a primary source rather than asking the tool
2. The client contract in the free chatbot	Protect what you put in	Where did that text go?	The traffic-light rule applied unprompted
3. The chatbot on our website	Right jobs; speak up	Who owns what it says?	A named owner, tested answers, a human route
4. The shortlist nobody chose	Right jobs; speak up	Who decides and who can explain?	A person reviewing, recording reasons and asking the supplier
5. The call that sounded like the boss	Know what it is; speak up	What control works whatever the call looks like?	Call-back, two-person approval, comfort in challenging
6. The note-taker in the meeting	Protect; speak up	Who needs to know?	Permission asked at the start; retention checked
7. The mood-reading software	Right jobs; know what it is	Should we do this at all?	Recognising the prohibition and the escalation route
8. The agent that acts on its own	Protect; check; speak up	What can it do without us?	Least permissions, logs, human approval for sending and paying

Closing question for every scenario: *What is one thing we will do differently on Monday*? Write it down and record it on Appendix D.

APPENDIX G

Quarterly Review Checklist

Review date: _____ **Attendees:**

Next review: _____

- ☐ Inventory updated: new, changed and removed tools added to Appendix A.
- ☐ Every tool has an owner, a status and a review date.
- ☐ Any tool affecting people has been looked at by a decision-maker.
- ☐ Role map updated for joiners, leavers and changes of role.
- ☐ New starters and contractors received the Tier 1 basics.
- ☐ Contractor confirmations current, with late replies chased.
- ☐ Acceptable-use sheet still matches what people do; version and date updated.
- ☐ Incidents and near-misses reviewed and written up.
- ☐ Training since the last review recorded in Appendix D.
- ☐ Next training sessions scheduled.
- ☐ Law and guidance checked: Article 4 text, the Commission FAQ, national authority.
- ☐ The one-page client summary updated.
- ☐ Evidence file complete, correctly stored and access-controlled.

Decisions made:

Actions (owner and date):

Glossary

AI literacy. The skills, knowledge and understanding that allow people to make an informed use of AI systems and be aware of their opportunities, risks and possible harm (AI Act, Article 3).

AI system. A machine-based system that infers from inputs how to generate outputs such as predictions, content, recommendations or decisions.

Deployer. A person or organisation that uses an AI system under its own authority in a professional setting.

Provider. A person or organisation that develops an AI system, or has one developed, and places it on the market or into service under its own name.

High-risk AI system. A system in a category the Act treats as higher risk, such as certain uses in recruitment, worker management or access to essential services, with additional obligations.

Human oversight. Assigning people with the competence, training and authority to monitor and, where needed, override a system's outputs.

Hallucination. An AI output that sounds plausible but is false or invented.

Shadow AI. Use of AI tools outside anything the business has approved or recorded.

Market surveillance authority. The national body responsible for enforcing the AI Act's supervision rules.

Deepfake. Synthetic audio, image or video that imitates a real person.

Agent. An AI tool that takes actions, such as sending messages or changing records, rather than only answering questions.

Digital Omnibus on AI. Regulation (EU) 2026/1744, which amended parts of the AI Act, including Article 4.

A Note on GDPRGard.eu

GDPRGard publishes free tools, guides and articles to help small businesses understand data-protection and AI duties in plain English, including how to map data sharing, handle a data breach and respond to rights requests. Visit **gdprgard.eu** for current guidance, templates and updates to the topics in this book.

About the Author

Zirlandia Milkovic writes practical, plain-English guides on data protection, cyber-security, AI and building a business with confidence. She believes compliance should not be a barrier to ambition but a foundation of trust, and that nobody should need a law degree to run a responsible small business.

She is the author of a series of business guides, including *GDPR Made Simple for Women Founders*, and her work focuses on turning complex rules and fast-moving technology into simple steps, checklists and templates that busy people can actually use. She is also behind **GDPRGard**, a growing collection of free GDPR tools, guides and articles for small businesses, available at gdprgard.eu.

The AI Literacy File grew out of a simple observation: many small businesses use AI every day, almost none can show what they have told their people about it, and the duty to support AI literacy is easier to meet than most vendors suggest. If this guide helped you, an honest review where you bought it helps other readers find it.

Visit **gdprgard.eu** for free tools, articles and updates.

Notes and Sources

Research cut-off: 11 November 2026. Items marked (read) were opened directly; others were seen through search summaries or secondary coverage and should be checked against the original before you rely on them.

1. Regulation (EU) 2024/1689 (Artificial Intelligence Act): entry into force 1 August 2024; Chapters I and II, including Articles 4 and 5, apply from 2 February 2025 (working knowledge; read the text on EUR-Lex).
2. Regulation (EU) 2024/1689, Article 3 (definitions of deployer, provider and AI literacy), Article 4 (original wording: measures to ensure, to their best extent, a sufficient level of AI literacy of staff and other persons dealing with the operation and use of AI systems on their behalf, taking into account their knowledge, experience, education, training, the context of use and the persons affected), Annex III (high-risk uses including recruitment and worker management) and Article 26 (deployer obligations, including human oversight by competent persons) (working knowledge; read the text).
3. European Commission, "AI Literacy: Questions and Answers", digital-strategy.ec.europa.eu, version last updated 20 October 2026 (read via a subagent summary). Points used: obligation remains for providers and deployers; no specific or "sufficient" level mandated; no certificate; internal record of trainings may be kept; no obligation to measure knowledge; no mandated governance structure; "other persons" include contractors, service providers and clients; context, role and risk to be considered; reliance on instructions for use alone is ineffective; enforcement by national market surveillance authorities from 2 August 2026; penalties set by national law; no criminal offences or right to compensation under the AI Act; support through practical examples, EU Digital Innovation Hubs and the AI Skills Academy. The page paraphrases the amended Article 4 and does not quote it; the exact text was not read.
4. Lewis Silkin, "The Digital Omnibus on AI enters into force today", 27 July 2026 (read via a subagent summary): publication in the Official Journal on 24 July 2026 and entry into force on 27 July 2026; Article 4 now requires providers and deployers to take measures, with the duty relaxed. The Commission's FAQ describes entry into force as "mid-July 2026". Confirm dates against EUR-Lex.
5. Regulation (EU) 2016/679 (GDPR), including Articles 5, 22, 28 and 32 (principles, automated decision-making, processors, security). Training records and recordings are personal data (working knowledge; read the text).

6. *Mata v. Avianca, Inc.* (S.D.N.Y., 2023): lawyers sanctioned after filing a brief containing non-existent cases generated by a chatbot (widely reported; working knowledge; read the court's order).
7. Reports in May 2023 (including Bloomberg) that Samsung restricted staff use of generative AI after employees uploaded confidential source code to a public chatbot (widely reported; working knowledge; read the original coverage).
8. *Moffatt v. Air Canada*, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal, February 2024): the airline held responsible for misleading information given by its website chatbot (working knowledge; read the decision).
9. Regulation (EU) 2024/1689, Article 50 (transparency obligations, including informing people that they are interacting with an AI system), applicable from 2 August 2026, with a grace period to 2 December 2026 for certain marking obligations for generative systems already on the market (working knowledge and earlier research; verify).
10. Arup statement, May 2024, and press coverage (including CNN and the Financial Times): an employee in Hong Kong deceived by a video call with deepfake colleagues, with transfers reported at about HK$200 million (about US$25 million) (widely reported; working knowledge; read the coverage).
11. Regulation (EU) 2024/1689, Article 5(1)(f): prohibition on AI systems inferring emotions of individuals in the workplace and education settings, except for medical or safety reasons; applies from 2 February 2025 (working knowledge; read the text and the Commission's guidelines on prohibited practices).
12. Secondary commentary on Regulation (EU) 2026/1744 (including Cuatrecasas, Bristows and Winston & Taylor summaries from search results): high-risk obligations postponed (Annex III systems to 2 December 2027 and Annex I systems to 2 August 2028 as understood by the author at the research cut-off) (secondary sources only; verify against the Official Journal).

Not verified by the author: the exact wording of the amended Article 4 in Regulation (EU) 2026/1744; the national enforcement arrangements in any particular Member State; whether any tool you use is a high-risk system; and the current features of any AI product mentioned generically. The case examples are real where sourced above; Orchard Lane Studio, Priya, Tomas and all other worked-example businesses and people are fictional.